


CWA- mesh generation for MPAS model

YING-JHANG WU, TING-AN CHEN, YU-HUNG LO,
Central Weather Administration
(*Technology Development & Numerical Information Divisions*)
Mentors: Leo Chen and Jay Chen

Mesh generation for MPAS model



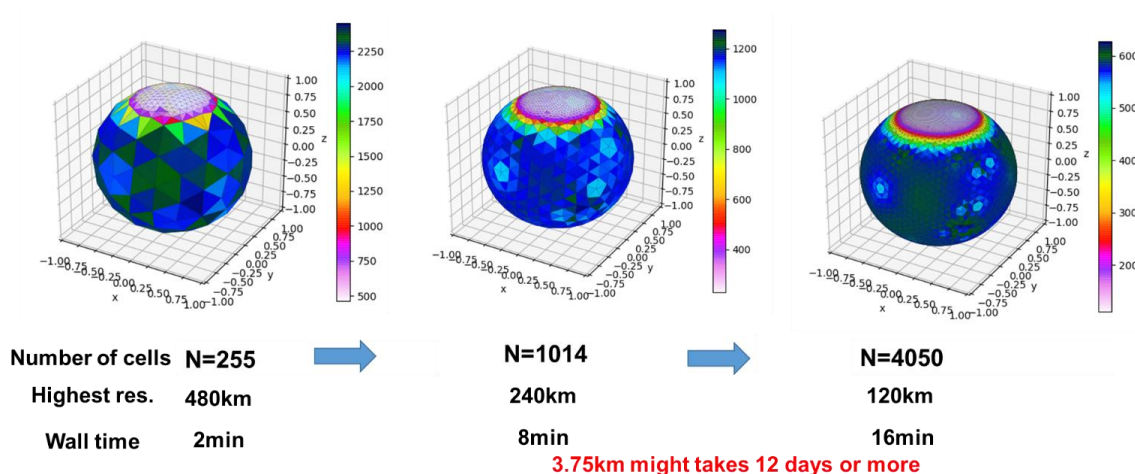
<https://mpas-dev.github.io/>

$$\rho(\mathbf{x}) = \frac{1-\gamma}{2} \left[\tanh \left(\frac{\beta - \|\mathbf{x}_c - \mathbf{x}\|}{\alpha} \right) + 1 \right] + \gamma$$

$$\frac{h_i}{h_j} \approx \left(\frac{\rho(\mathbf{z}_j)}{\rho(\mathbf{z}_i)} \right)^{1/4}$$

Ringler, 2011

- MPAS mesh generation establishes the MPAS model's horizontal resolution
- Fortran Language
- LLOYD'S METHOD




- Require several months to complete a single task

Voronoi Tessellation


各綠點彼此連線為中垂線

1. 首先需要給定任意格點(綠點所示)
2. 找出相鄰三格點的外心(白色點)




Centroidal Voronoi Tessellation

(Ju et al. 2002)



$$z_i \neq z_i^*, i = 1, \dots, k$$



$$z_i = z_i^*, i = 1, \dots, k$$

1. 若無特別處理，質心(z_i^*)與參考點 z_i 並不會重合
2. 我們希望網格各中心連線是垂直平分,且交點恰為質心

Evolution and Strategy

- What was your goal coming here?
 - Porting on GPU
 - Learn to use openACC
- What was your initial strategy?
 - Modify the source code from openMP to openACC
- How did this strategy change?
 - Under planned (stick to the plan.)

Progress

- Two parts of computation:
 1. Voronoi computation: on CPU
 2. Centroid computation: openMP -> openACC (on GPU)
- Porting part 2 (Centroid computation) on A100
- Compiler: gfortran VS pgfortran VS nvfortran

```
Time for Voronoi computation: 0.15 s , 1000000 Hz
Time for centroid computation: 0.28 s , 1000000 Hz
401410 2000 8.8188361151890690E-010 5.1832255029052828E-010
Time for Voronoi computation: 0.15 s , 1000000 Hz
Time for centroid computation: 0.23 s , 1000000 Hz
401410 2100 8.8181093699007062E-010 5.1827546100969774E-010
Time for Voronoi computation: 0.14 s , 1000000 Hz
Time for centroid computation: 0.25 s , 1000000 Hz
401410 2200 8.8172499566554482E-010 5.1822837889265424E-010
```

gfortran

```
Time for Voronoi computation: 0.15 s , 1000000 Hz
Time for centroid computation: 0.24 s , 1000000 Hz
401410 1100 8.8260357064804557E-010 5.1874652629835801E-010
Time for Voronoi computation: 0.15 s , 1000000 Hz
Time for centroid computation: 0.27 s , 1000000 Hz
401410 1200 8.8253126936933654E-010 5.1869940007437832E-010
Time for Voronoi computation: 0.16 s , 1000000 Hz
Time for centroid computation: 0.21 s , 1000000 Hz
401410 1300 8.8244646233434736E-010 5.1865227956434174E-010
```

pgfortran

```
Time for Voronoi computation: 0.16 s , 1000000 Hz
Time for centroid computation: 0.03 s , 1000000 Hz
401410 400 8.8316747102826590E-010 5.1907652823898715E-010
Time for Voronoi computation: 0.16 s , 1000000 Hz
Time for centroid computation: 0.03 s , 1000000 Hz
401410 500 8.8308888330564438E-010 5.1902937352481088E-010
Time for Voronoi computation: 0.15 s , 1000000 Hz
Time for centroid computation: 0.03 s , 1000000 Hz
401410 600 8.8300602649377994E-010 5.1898221990084138E-010
```

nvfortran

Use NVTX

- use nvtx
- call nvtxStartRange("...")

.....

nvtxEndRange

- call nvtxStartRange("...")

.....

nvtxEndRange

- FFLAGS = -acc ... -l nvhpcwrapnvtw

...

LDFLAGS = -acc ... -l nvhpcwrapnvtw

```
+ 1 +-- 76 lines: module voronoi_utils-----
77 |-----
78 | SUBROUTINE COMPUTE_VC
79 |
80 | Compute the Voronoi corners of a set of lat/lon locations.
81 |-----
82 | subroutine compute_vc(rlat, rlon, n, nrow, ntmx, list, lptr, lend, listc, vclat,
83 | use nvtx
84 |
85 | implicit none
86 |
87 | integer, intent(in) :: n, nrow, ntmx, nvc
88 | integer, dimension(nvc), intent(inout) :: list, lptr, listc
89 | real, dimension(nvc), intent(inout) :: vclat, vclon
+ 90 +-- 9 lines: integer, dimension(n), intent(inout) :: lend-----
99
100 | if (nvc < 6*n-12) then
101 |   write(0,*) 'Error: Argument nvc to COMPUTE_VC must be at least 6*n+12'
102 |   return
103 | end if
104 |
105 | call nvtxStartRange("Trans")
106 | call trans(n, rlat, rlon, x, y, z)
107 |
108 | ! It is time again to check whether our triangulation is still valid
109 | !
110 | call nvtxEndRange
111 | if (ttl == 0) then
112 |   !
113 |   ! Compute triangulation based on current generating points
114 |   !
115 |   call nvtxStartRange("trmesh")
116 |   call trmesh(n, x, y, z, list, lptr, lend, lnew, near, next, dist, ierr)
117 |   call nvtxEndRange
118 |   if (ierr /= 0) then
119 |     write(0,*) 'Error: TRMESH returned error code ', ierr
120 |   end if
121 |
122 |
123 |
```

```
18 LDFLAGS = -O3 -fopenmp
19 endif
20
21 ifeq ($(CORE), nv)
22 FC = nvfortran
23 CC = gcc
24 FFLAGS = -acc -Minfo -fast -xvect=levels:8 -l nvhpcwrapnvtw
25 #FFLAGS = -acc -Minfo -fast
26 F77FLAGS =
27 CFLAGS = -O3
28 CPPFLAGS = -DRKIND=8
29 PROMOTION = -r8
30 LDFLAGS = -acc -Minfo -fast -l nvhpcwrapnvtw
31 endif
32 #FC = xlf2003_r
33 #CC = xlc_r
34 #FFLAGS = -O4 -qsmmp=omp
35 #F77FLAGS = -qfixed -O4 -qsmmp=omp
36 #CFLAGS = -O4
+ 37 +-- 11 lines: CPPFLAGS = -DRKIND=8 -Uvector-----
48 #LDFLAGS = -Ofast -mp
49
50 ifeq ($(CORE), pg)
51 FC = pgfortran
52 CC = pgcc
53 FFLAGS = -Mfree -O3 -mp
54 #FFLAGS = -O3
55 CFLAGS = -O3
56 CPPFLAGS = -DRKIND=8
57 PROMOTION = -r8
```

Results and Final Profile

- What were you able to accomplish?
 - 2x faster (*256 CPU vs 1GPU numbers)
- What did you learn?
 - Syntax difference & Nsight profiler & nvtx label
 - Achieved new scientific goals?
 - Speed up about 8x (only modify 25 lines)

Process Overview

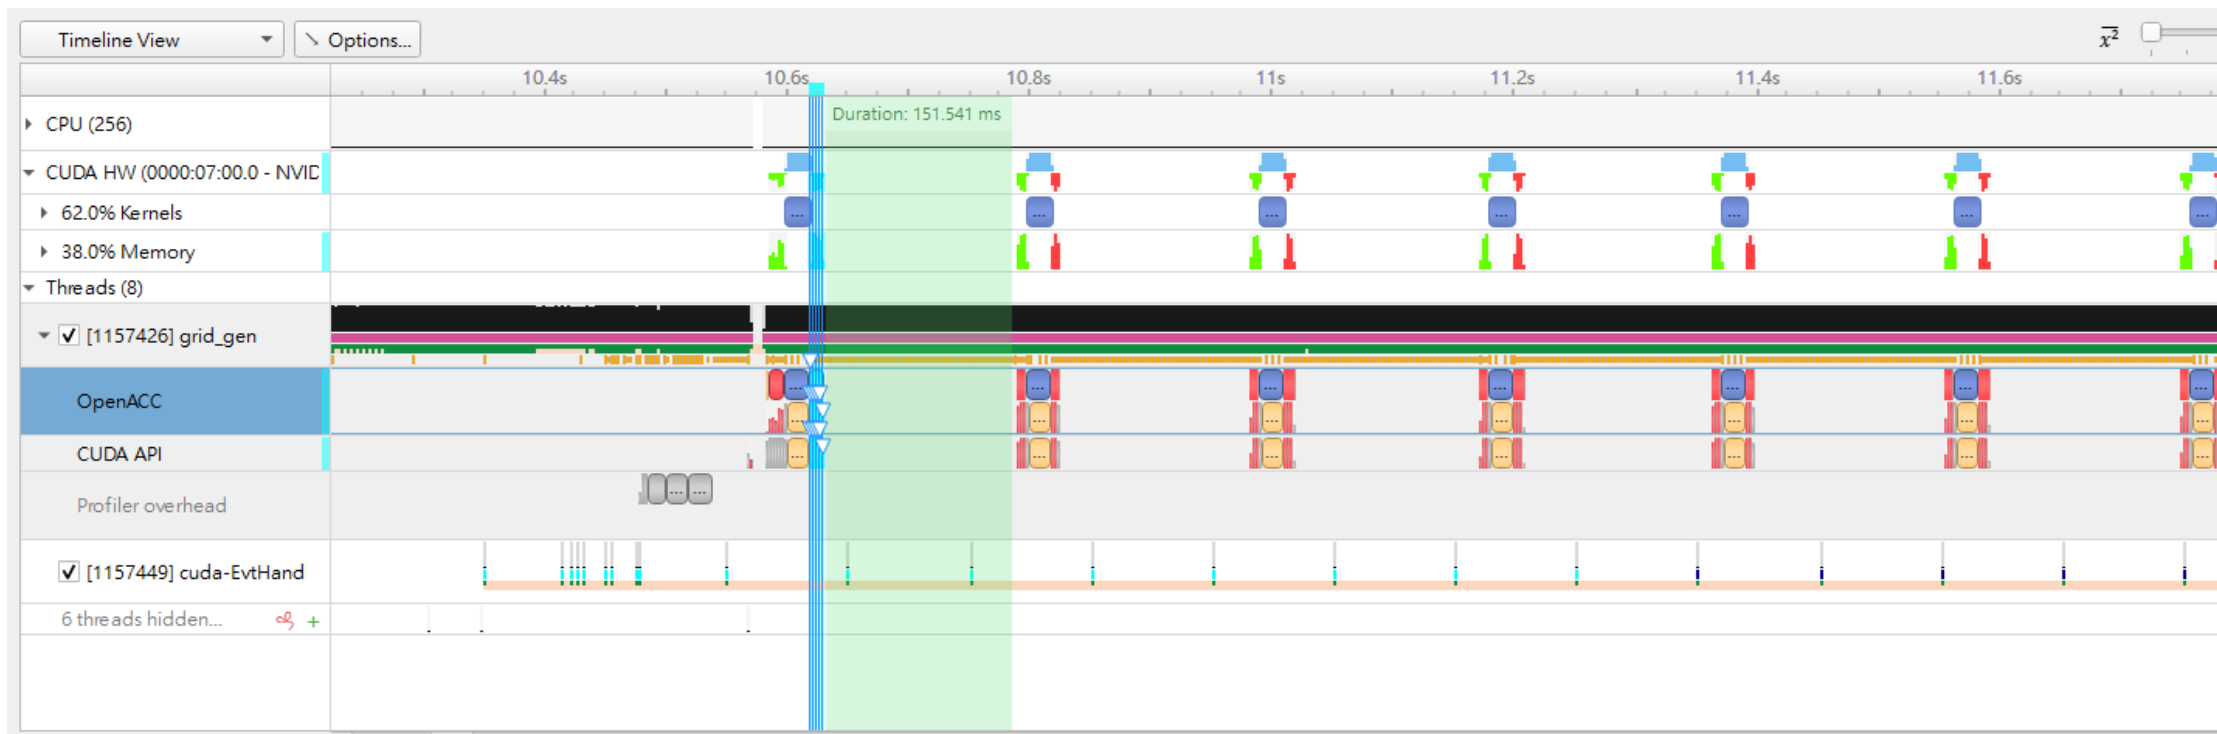
Voronoi Computation: No differences among 3 scenarios

Centroid Computation: nvfortran (on GPU) speeds up, only need 1/8

| | gfortran | pgfortran | nvfortran |
|----------------------|----------|-----------|-----------|
| Voronoi computation | 0.15 | 0.15 | 0.15 |
| Centroid computation | 0.25 | 0.24 | 0.03 |
| Total | 0.4 | 0.39 | 0.18 |
| Ratio | 1 | 0.975 | 0.45 |

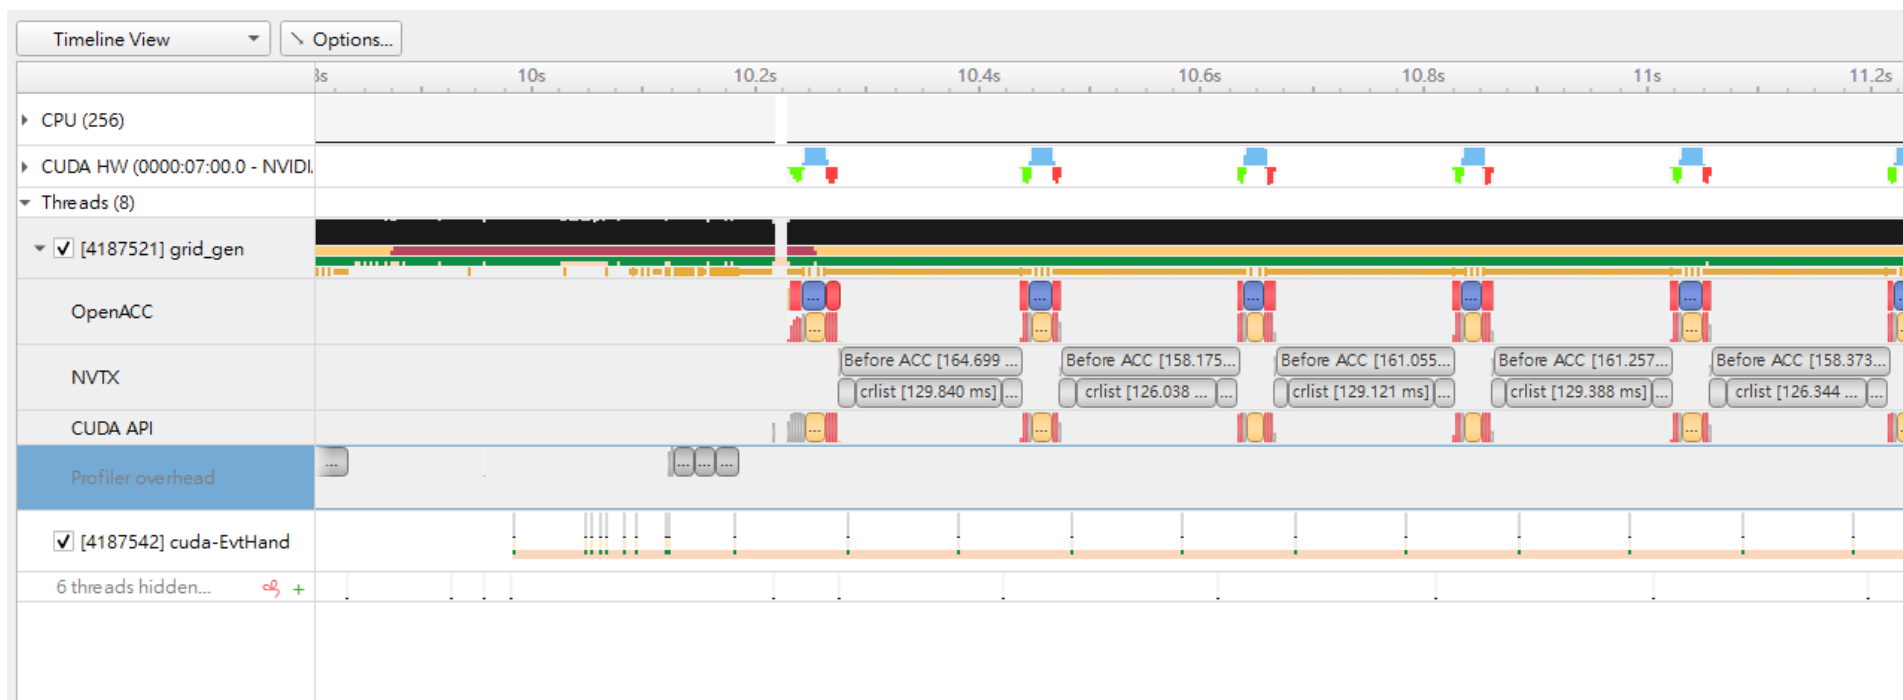
Profiler Output

- Nsight

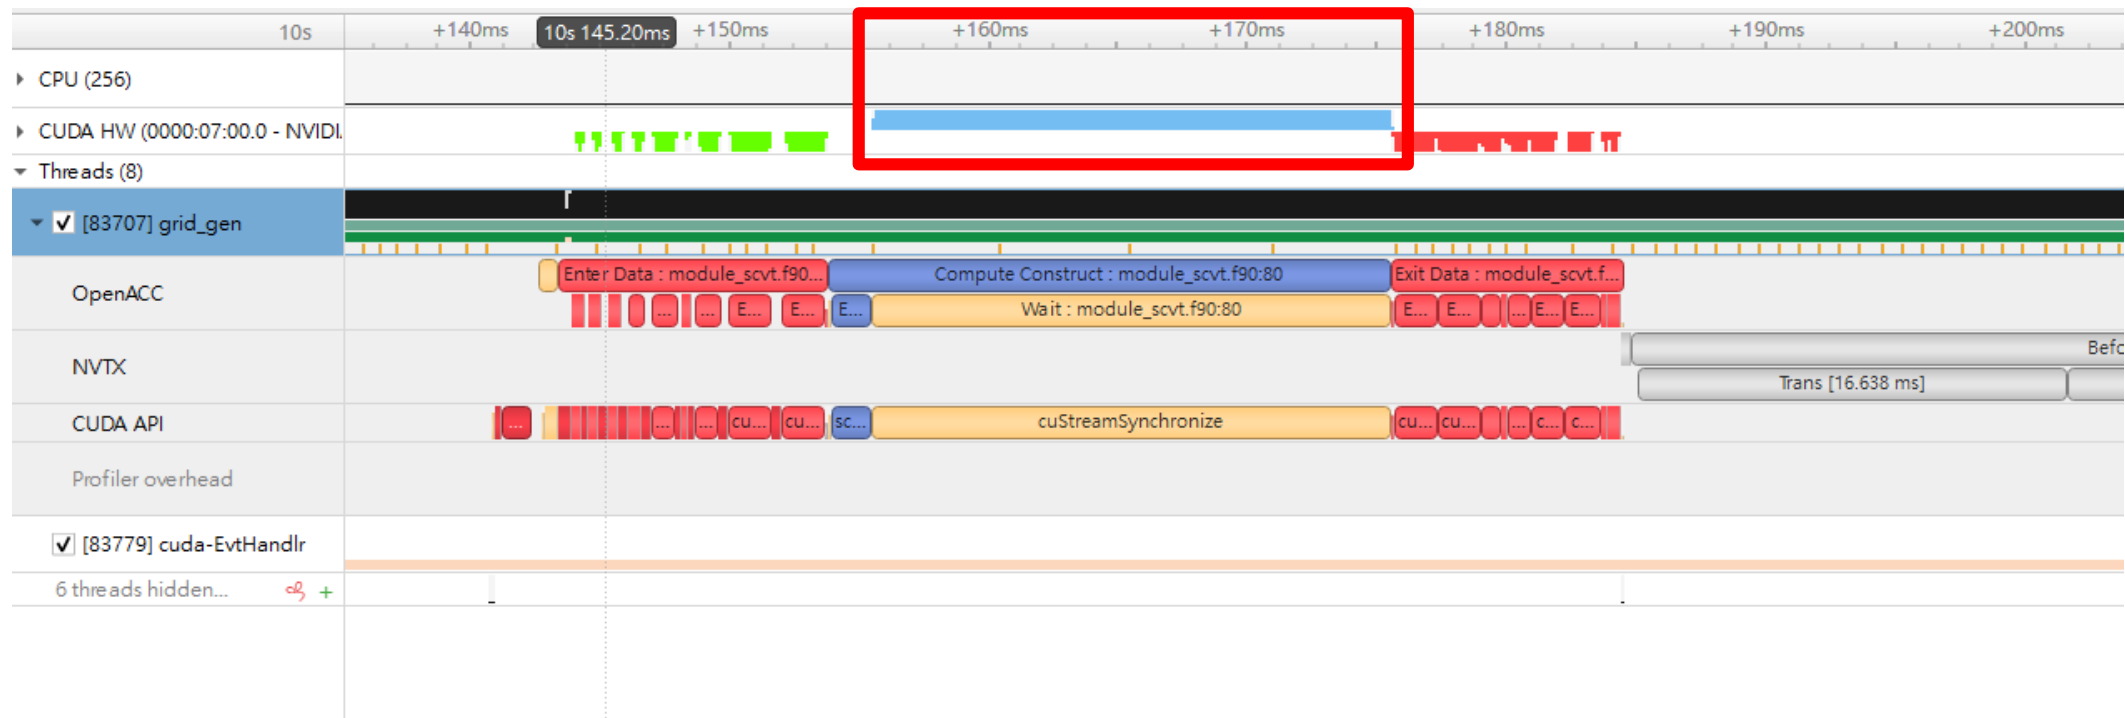


Profiler Output

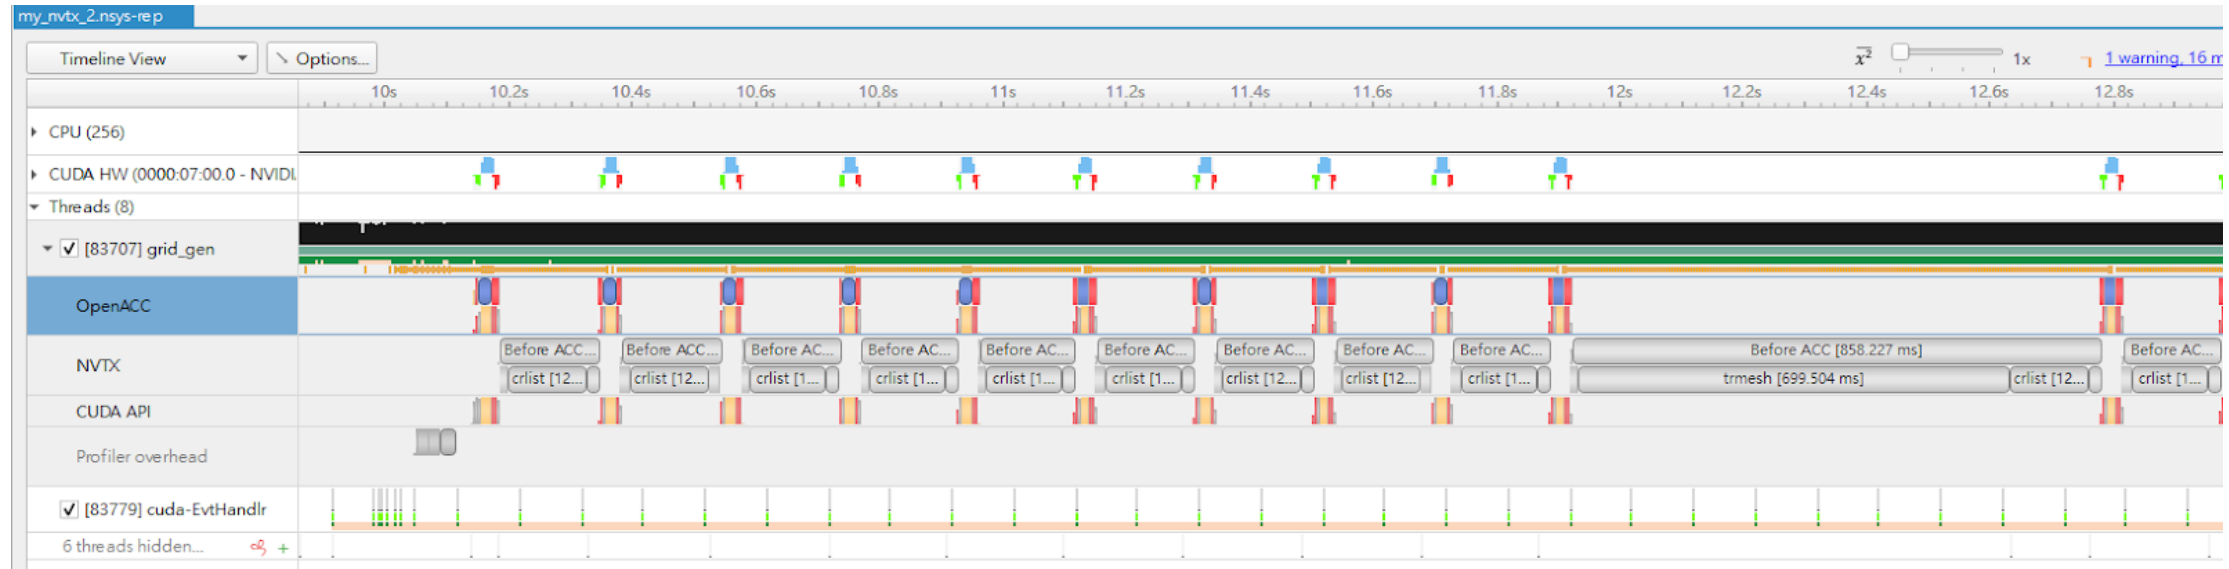
- add nvtx for CPU part



Profiler Output



Process Overview



Problems and Solutions

- Allocation and deallocation of array of type is not supported on device

```
543 type (geo_point), allocatable, dimension(:)
544 type (geo_point), allocatable, dimension(:)
545
546 glevel = nint(log(real(n)) / log(4.0)) ! E
547                                     !
548
549 glevel = (2 ** glevel) + 1
550 allocate(line(glevel, glevel))
551 allocate(p1p2(glevel))
552 allocate(p1p3(glevel))
```

```
549 type (geo_point), dimension(5) :: p1p2, p1
550 type (geo_point), dimension(5,5) :: line
-----
-----
-----
551
552 ! glevel = nint(log(real(n)) / log(4.0)) !
553 !
-----
-----
```

Problems and Solutions

- Syntax difference between gfortran(left) and nvfortran(right)

| | | | |
|-----|--------------------------|-----|--------------------------|
| 127 | | 141 | |
| 128 | do while (k /= lend(i)) | 142 | do while (k /= lend(i)) |
| 129 | k = lptr(k) | 143 | k = lptr(k) |
| 130 | p2 = p3 | 144 | p2%lat = p3%lat |
| | ----- | 145 | p2%lon = p3%lon |
| 131 | p3%lat = vclat(listc(k)) | 146 | p3%lat = vclat(listc(k)) |
| 132 | p3%lon = vclon(listc(k)) | 147 | p3%lon = vclon(listc(k)) |

- The 'write' formatting in Fortran is not supported by the nvfortran compiler

Energy Efficiency

| INPUTS | |
|---------------------|------|
| # CPU Cores | 256 |
| # GPUs (A100) | 1 |
| Application Speedup | 8.0x |

| | |
|------------------|--------|
| Node Replacement | 128.0x |
|------------------|--------|

| GPU NODE POWER SAVINGS | | | |
|------------------------|--------------------|-------------------|----------------|
| | AMD Dual Rome 7742 | 8x A100 80GB SXM4 | Power Savings |
| Compute Power (W) | 140,800 | 6,500 | 134,300 |
| Networking Power (W) | 5,944 | 93 | 5,851 |
| Total Power (W) | 146,744 | 6,593 | 140,151 |

| | |
|-----------------------|-------|
| Node Power efficiency | 22.3x |
|-----------------------|-------|

| ANNUAL ENERGY SAVINGS PER GPU NODE | | | |
|------------------------------------|--------------------|-------------------|------------------|
| | AMD Dual Rome 7742 | 8x A100 80GB SXM4 | Power Savings |
| Compute Power (kWh/year) | 1,233,408 | 56,940 | 1,176,468 |
| Networking Power (kWh/year) | 52,069 | 814 | 51,255 |
| Total Power (kWh/year) | 1,285,477 | 57,754 | 1,227,723 |

| | |
|---------------------|-----------------|
| \$/kWh | \$ 0.34 |
| Annual Cost Savings | \$ 417,425.82 |
| 3-year Cost Savings | \$ 1,252,277.46 |

| | |
|------------------------------------|--------|
| Metric Tons of CO2 | 870 |
| Gasoline Cars Driven for 1 year | 188 |
| Seedlings Trees grown for 10 years | 14,389 |

[\(source: Link\)](#)

| POWER ASSUMPTIONS | | |
|------------------------------|-------------------------------------|----------------------------------|
| Node Configurations | Baseline Node AMD Dual Rome 7742 | Alternative 8x A100 80GB SXM4 |
| CPU SKU | 7742 | 7742 |
| # CPU | 2 | 2 |
| # CPU Cores | 128 | 128 |
| CPU Power (W) | 450 | 450 |
| GPU SKU | 0 | A100 80GB SXM4 |
| # GPU | 0 | 8 |
| GPU Power (W) | 0 | 3200 |
| Network Type | IB EDR | IB EDR |
| # Network Ports | 1 | 2 |
| Network Card Power (W) | 30 | 60 |
| RBoM Power (W) | 300 | 450 |
| Total Compute Node Power (W) | 1100 | 6500 |
| Core Network Power / Node | 46 | 93 |
| Total Power / Node | 1146 | 6593 |

ASSUMPTIONS

- (1) The workload being input will run 24/7/365 on the node in question
- (2) When the workload runs on a fraction of a CPU or GPU server, no other bottlenecks occur to stop it from scaling up to occupy the full server
- (3) The calculations use TDP for both CPU and GPU. In reality, neither server will run full time at TDP. The comparison here is "worst case CPU" vs. "worst case GPU"
- (4) Annual cost savings are operational for electricity only. Capital, personnel, etc are not included

Wishlist

- What do you wish existed to make your life easier?
 - Automatic Tools (openMP to openACC)
 - Lessons
 - Expand the research probability

Please use 100 words to summarize your team's achievements during this Hackathon

1. Realize the skills on GPUs, including segmentation and parallelization.
2. Adapt existing Fortran code from openMP to OpenACC , and achieve 8x speed-up on Centroid computation.
3. Realize the property of the mesh generation code, and find the bottleneck of acceleration.
4. Acquire proficiency in profiling tools (nsight) for code analysis.

Future work:

Endeavor to apply in research and operational environments in CWA.