財團法人
國家實驗研究院
National Applied Research Laboratories

Day 3 –
LLM inference with TensorRT-LLM on NCHC servers

Team member:
Fang-An Kuo, Kuo-Teng Ding, Meng-Chi Huang,
NCHC Speedrunning team

NVIDIA Mentor:
Anthony
Cliff

Inference Engine Optimization

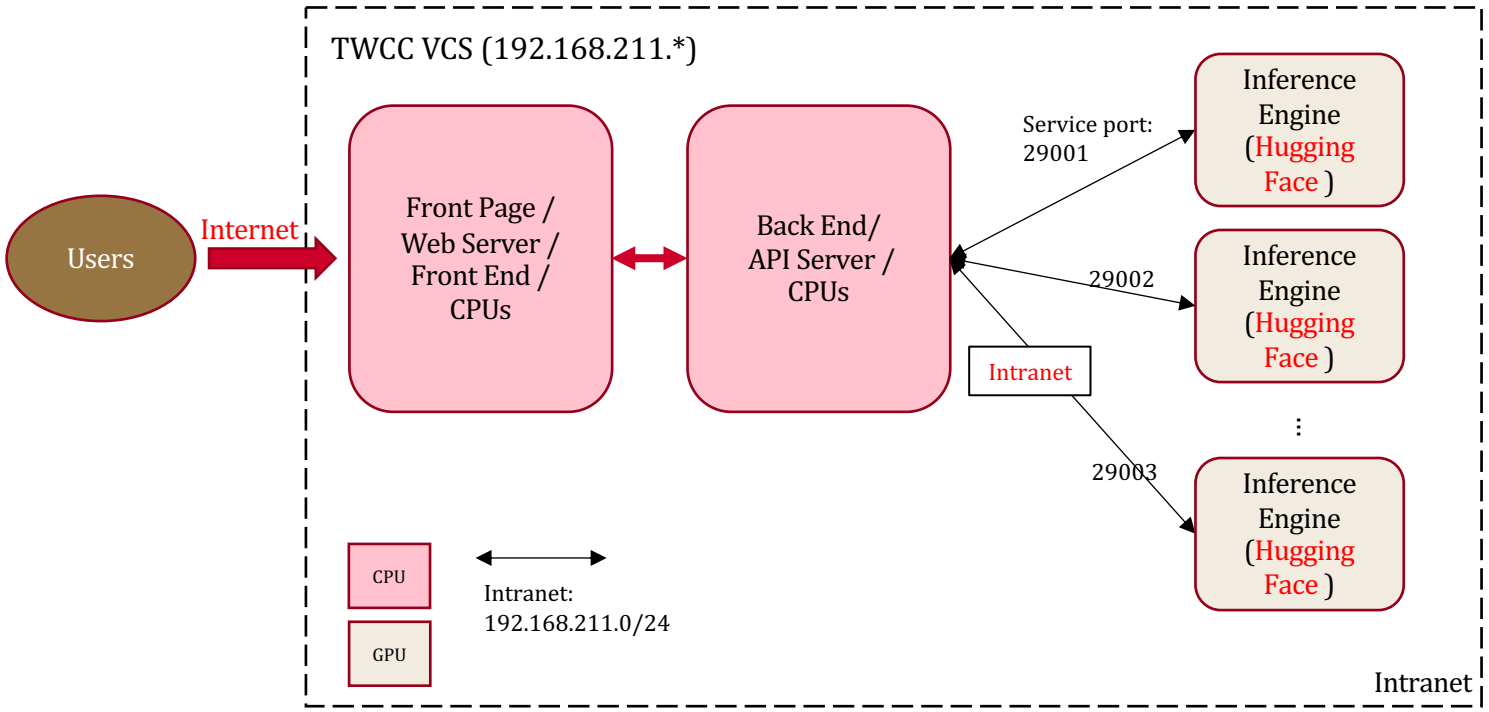

Inference Engine Optimization

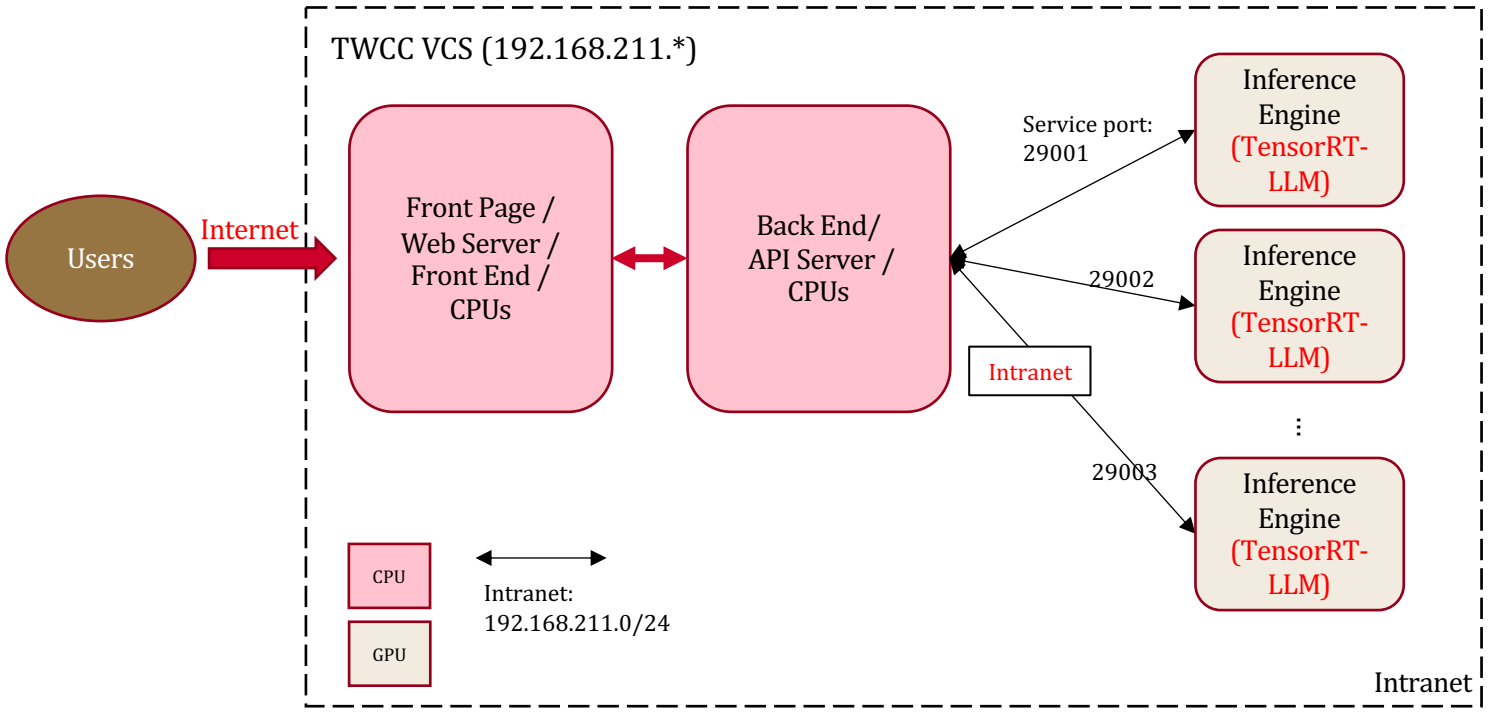

Performance Benchmark on the Demo Website

Performance benchmarking of the both inference engines

Benchmarks of inferencing Traditional Chinese LLM model on NVIDIA GPUs

- We replace the transformers inference APIs by using TensorRT-LLM APIs
- The benchmarks run on NVIDIA GPUs, including V100/A100/A6000/H100
  - LLM Model: TAIDE Model based on LLaMA2-7B with the version number, b11.
  - Batch size = 1
  - Tensor Parallel = 1 and Pipeline Parallel = 1

Inference APIs	Tokens/sec.			Computing time			Words/sec.		
TAIDE 7B (b11) with FP16	V100	A100	A6000	V100	A100	A6000	V100	A100	A6000
Hugging Face	10.63	5.38	10.88	26.13	29.53	29.01	35.45	40.05	39.35
NVIDIA TRT-LLM	19.44	17.21	17.51	52.30	103.35	51.10	63.59	125.65	62.13
Speedup	1.83	3.20	1.61	2.001	3.500	1.761	1.794	3.138	1.579
Speedup, compared to V100	1.829	3.613	1.787	2.001	3.954	1.955	1.794	3.545	1.753
INT8 model									
Speedup									

Benchmarks of inferencing Traditional Chinese LLM model on NVIDIA GPUs

- The benchmarks run on NVIDIA GPUs, including V100/A100/A6000/H100
  - LLM Model: TAIDE Model based on LLaMA2-7B with the version number, b11.
  - Batch size = 1
  - Tensor Parallel = 1 and Pipeline Parallel = 1

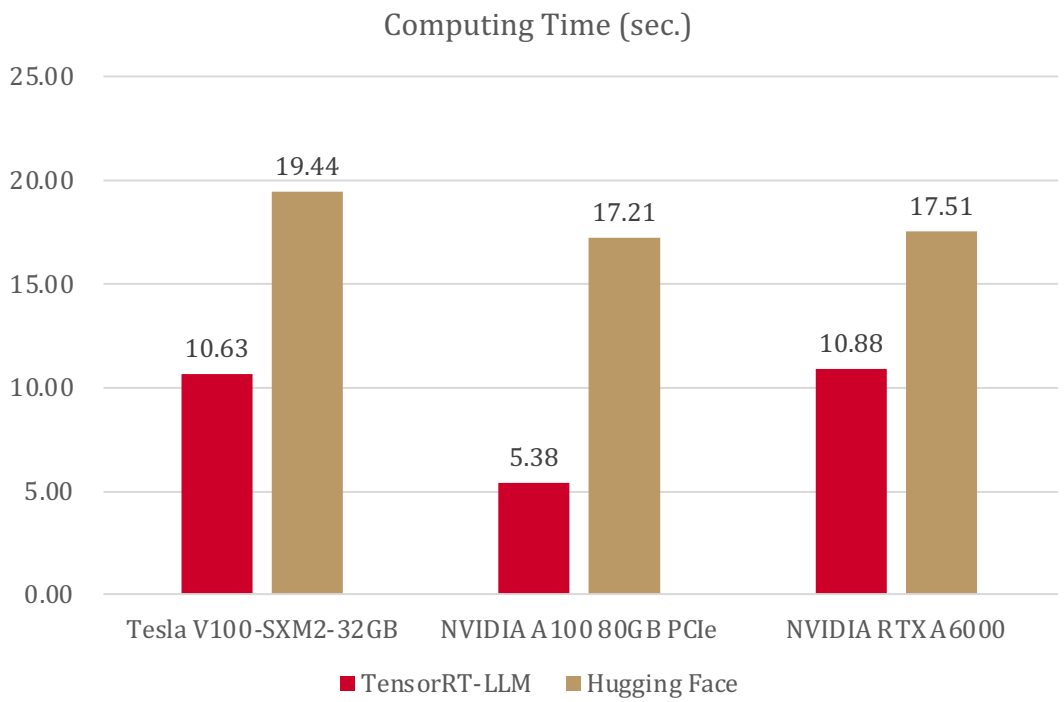

Benchmarks of inferencing Traditional Chinese LLM model on NVIDIA GPUs

- The benchmarks run on NVIDIA GPUs, including V100/A100/A6000/H100
  - LLM Model: TAIDE Model based on LLaMA2-7B with the version number, b11.
  - Batch size = 1
  - Tensor Parallel = 1 and Pipeline Parallel = 1
  - The speedup of generating tokens by using the NVIDIA TensorRT-LLM engines is about 3.5x, which is based on NVIDIA A100

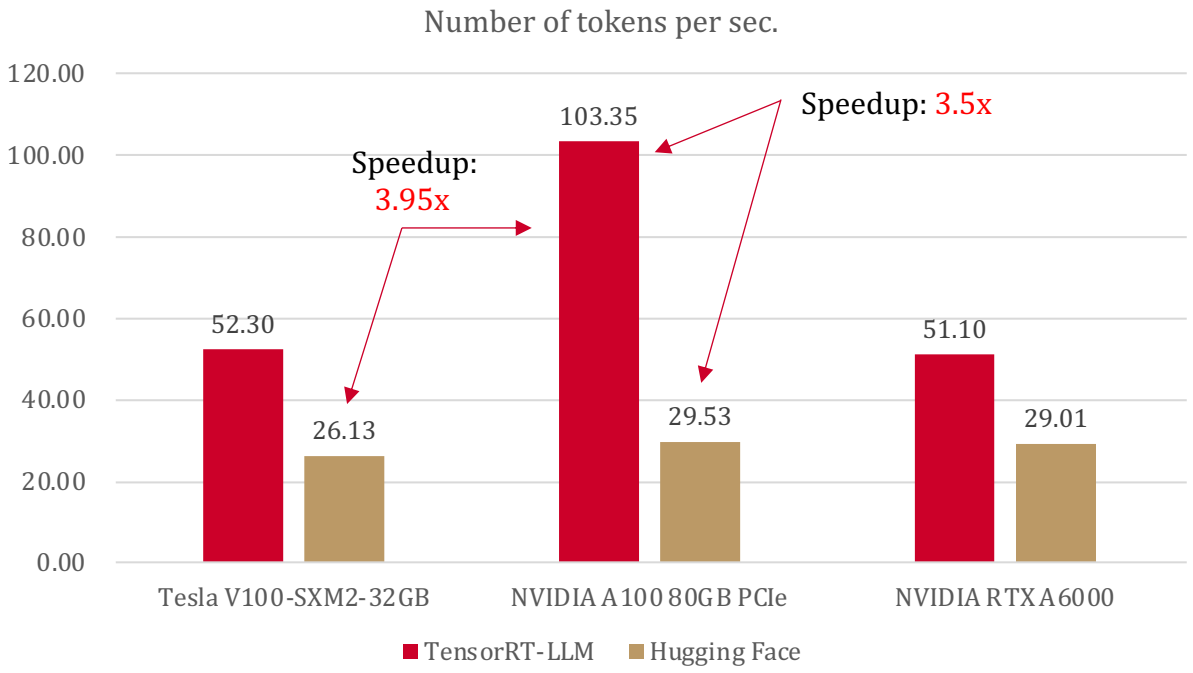

Benchmarks of inferencing Traditional Chinese LLM model on NVIDIA GPUs

- The benchmarks run on NVIDIA GPUs, including V100/A100/A6000/H100
  - LLM Model: TAIDE Model based on LLaMA2-7B with the version number, b11.
  - Batch size = 1
  - Tensor Parallel = 1 and Pipeline Parallel = 1
  - The speedup of generating words by using the NVIDIA TensorRT-LLM engines is about 3.14x, which is based on NVIDIA A100

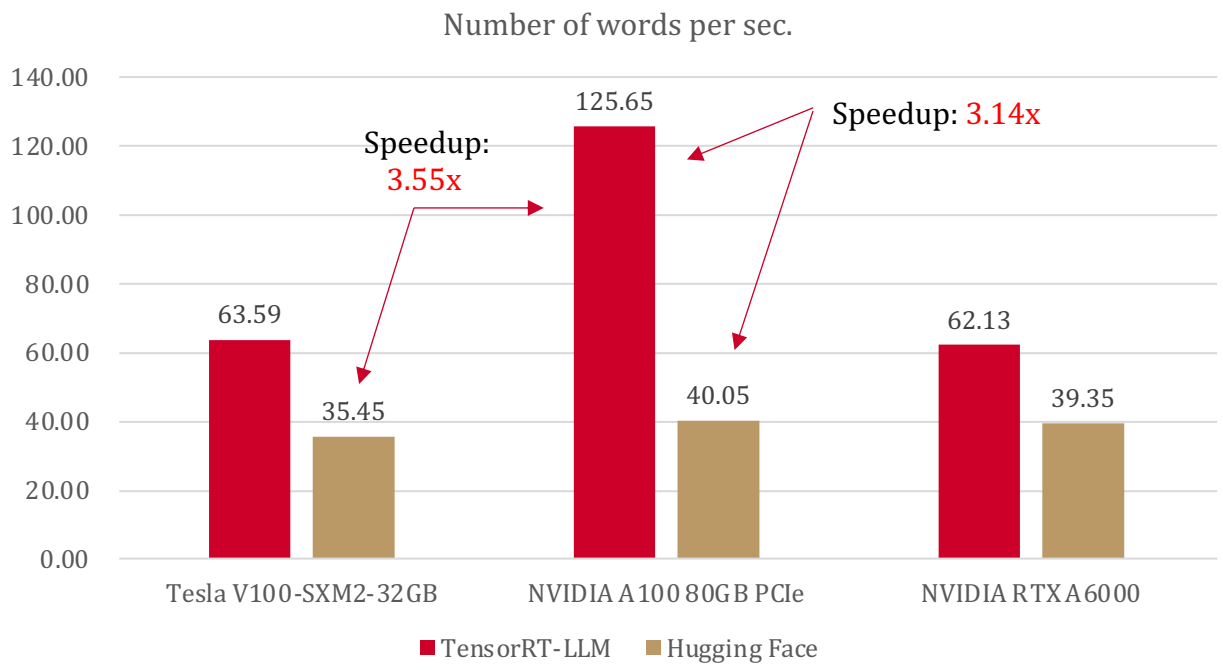

INT8 model

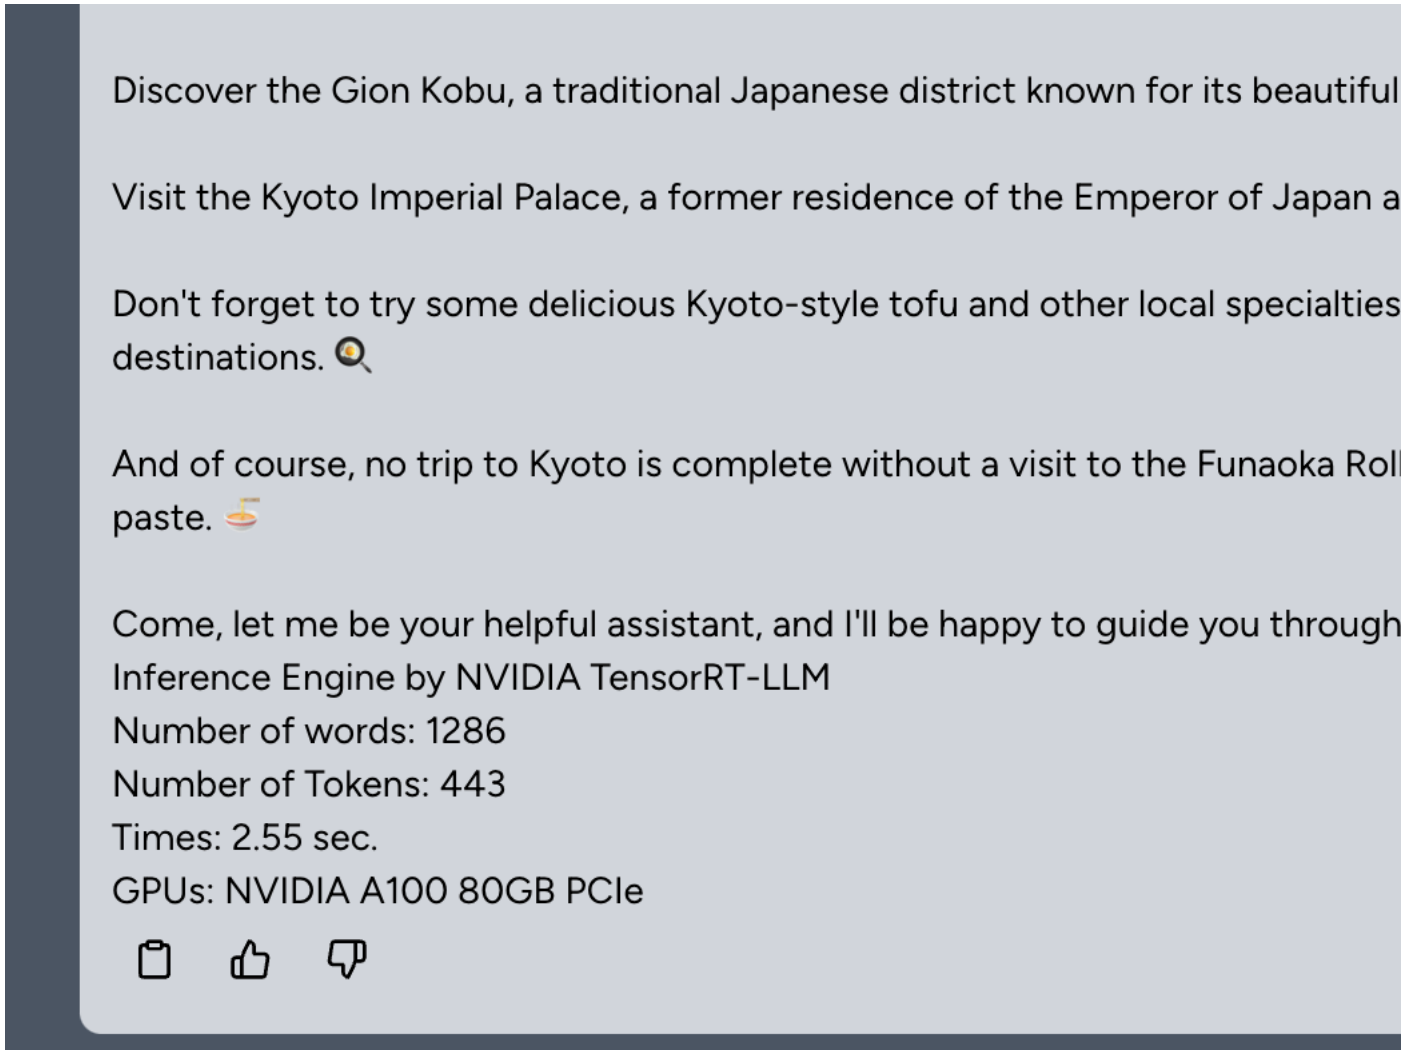
Discover the Gion Kobu, a traditional Japanese district known for its beautiful
Visit the Kyoto Imperial Palace, a former residence of the Emperor of Japan a
Don't forget to try some delicious Kyoto-style tofu and other local specialties
destinations.
And of course, no trip to Kyoto is complete without a visit to the Funaoka Roll
paste.
Come, let me be your helpful assistant, and I'll be happy to guide you through
Inference Engine by NVIDIA TensorRT-LLM
Number of words: 1286
Number of Tokens: 443
Times: 2.55 sec.
GPUs: NVIDIA A100 80GB PCIe

Problems and Solutions

- Modify {TensorRT-LLM}/docker/common/install_base.sh
  - OS: ubuntu 22.04

```
init_ubuntu() {
    apt-get update
    apt-get install -y --no-install-recommends wget gdb git-lfs python3-pip python3-dev python-is-python3 libffi-dev
    apt-get install -y --no-install-recommends screen gpustat nvtop curl iftop
    if ! command -v mpirun &> /dev/null; then
      DEBIAN_FRONTEND=noninteractive apt-get install -y --no-install-recommends openmpi-bin libopenmpi-dev
    fi
    apt-get clean
    rm -rf /var/lib/apt/lists/*
    # Remove previous TRT installation
    if [[ $(apt list --installed | grep libnvinfer) ]]; then
        apt-get remove --purge -y libnvinfer*
    fi
    if [[ $(apt list --installed | grep tensorrt) ]]; then
        apt-get remove --purge -y tensorrt*
    fi
    pip uninstall -y tensorrt
    pip install flask flask_sse datasets nltk rouge_score
}
```

Problems and Solutions

- Generation cannot properly stops at EOS (</s>) and only stops at max output length
  - Solution: output_text.replace("</s>","")
  - Another solution and reference:
    - https://github.com/NVIDIA/TensorRT-LLM/blob/a21e2f85178111fed9812bb88c2cc7411b25f0ba/examples/gpt/run.py#L299

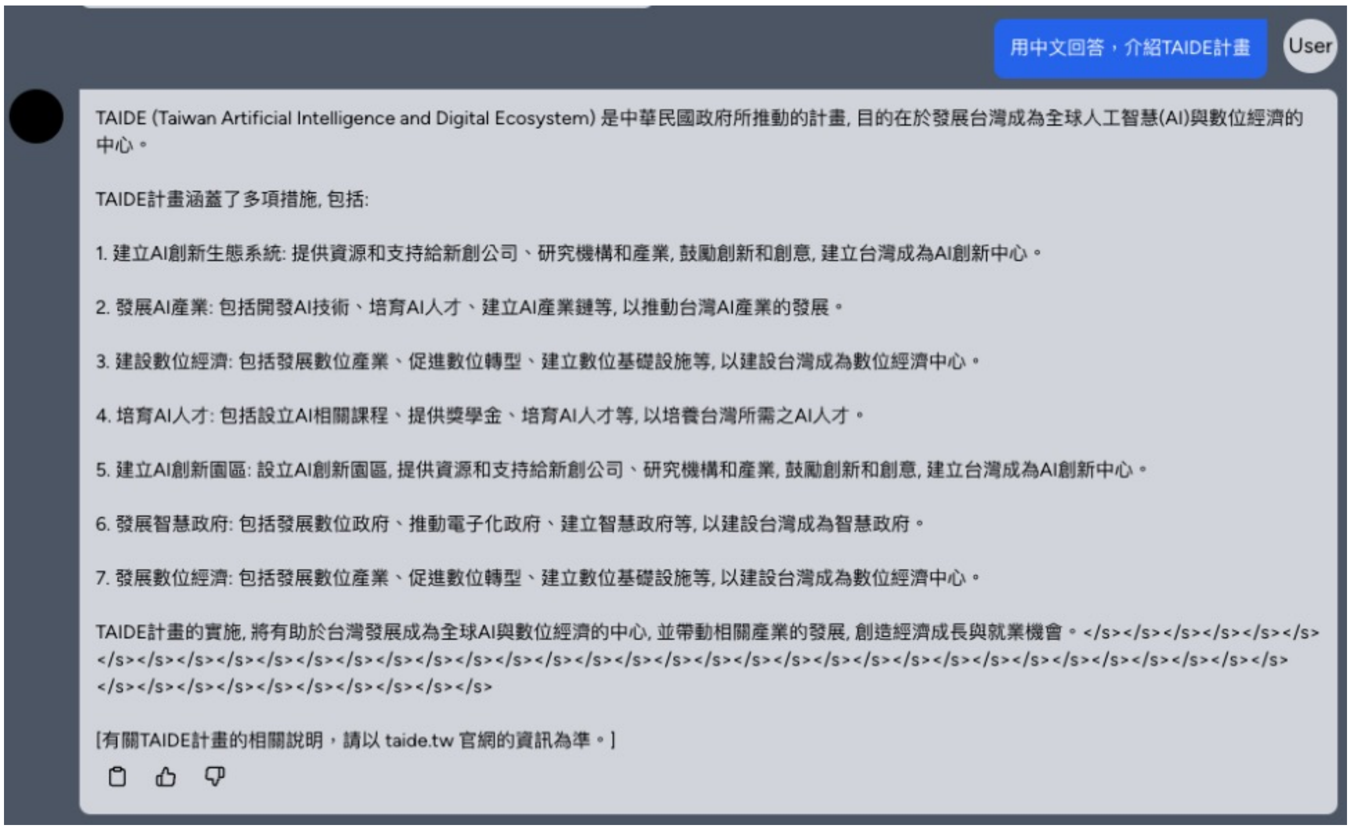

NARLabs National Applied Research Laboratories


thank you
danke
謝謝
ngiyabonga
teşekkür ederim
gracias
merci
obrigado
bedankt
dziękuję
спасибо
sukriya
kop khun krap
go raibh maith agat
mochchakkeram
terima kasih
grazie
arigatô
dakujem
takk
taiku
tapadh leat
dank je
мерси
ありがとう
감사합니다
ευχαριστώ
hvala
mauruuru
sagolun
рахмат
Баярлалаа
mersi
barka
welalin
tack
spas
faafetai lava
vinaka
blagodaram
misaotra
matondo
paldies
grazzi
mahalo
kiitos
dankie
dhanyavad
nanni
nandri
enkosi
bayarlalaa
gràçie
köszönöm
chnorakaloutioun
gratias ago
gràcies
sulpáy
djiere dieuf
tau
дякую
mamnun
chokrane
murakoze
tenki
obrigada
asante
manana
хвала
sobodi
děkuji
mèsi
didi madloba
kam sah hamnida
najis tuke
rahmat
tanemirt
rahmet
xièxie
diolch
dhanyavadagalu
shukriya
mercé
trugarez
akun
dankon
aciú
spasibi
tenki