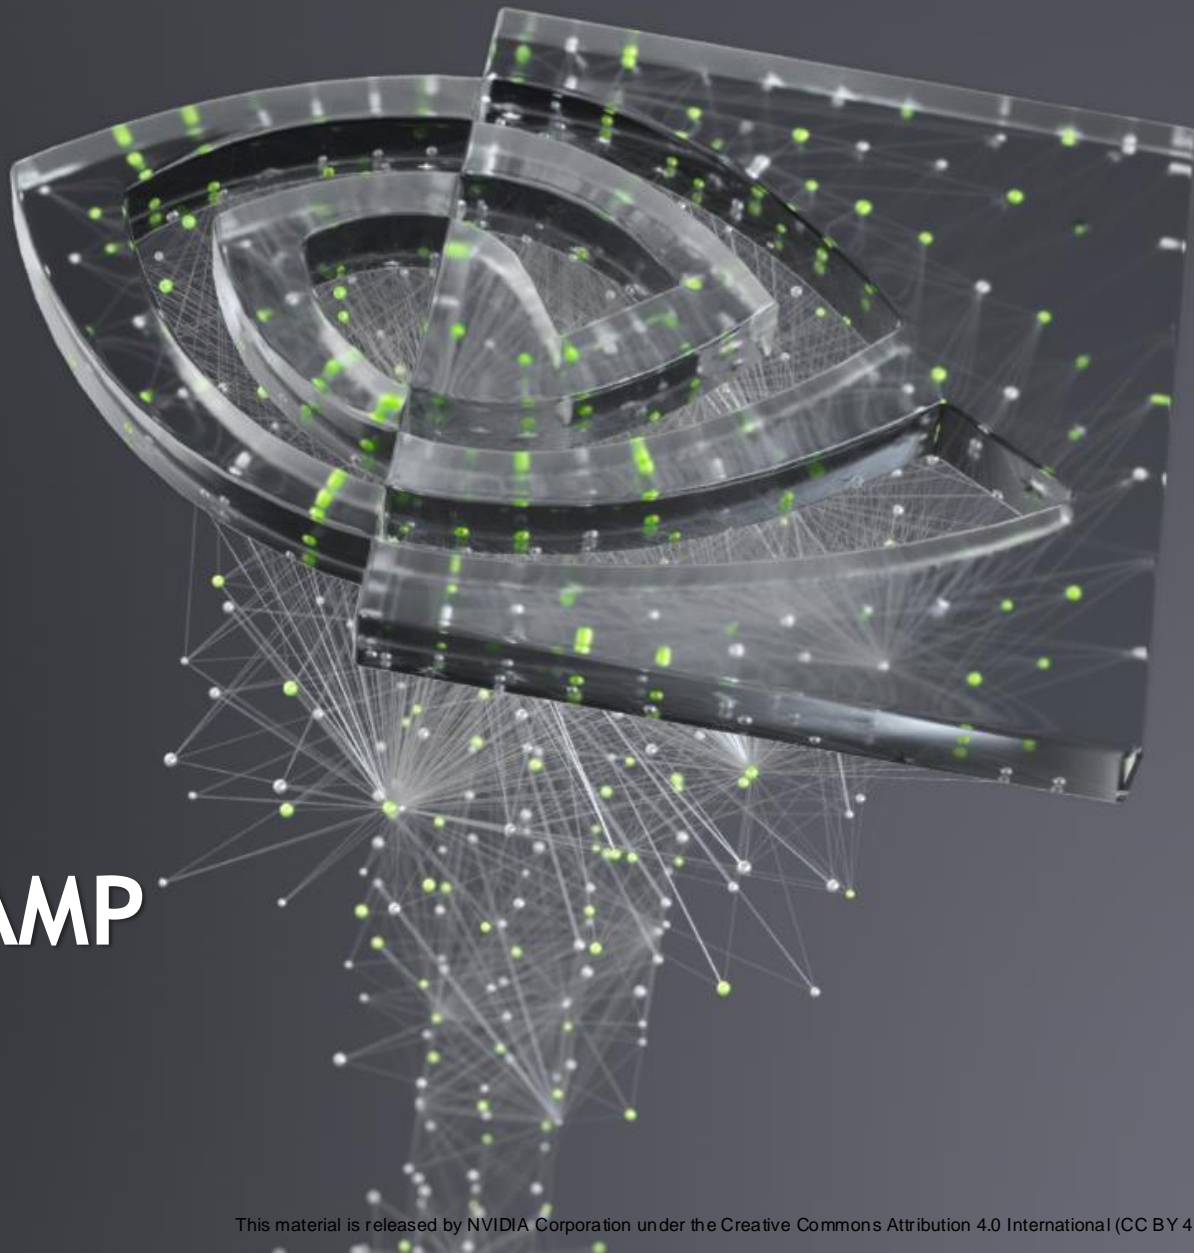




# N-WAYS GPU BOOTCAMP

## A QUICK GUIDE TO CUPY

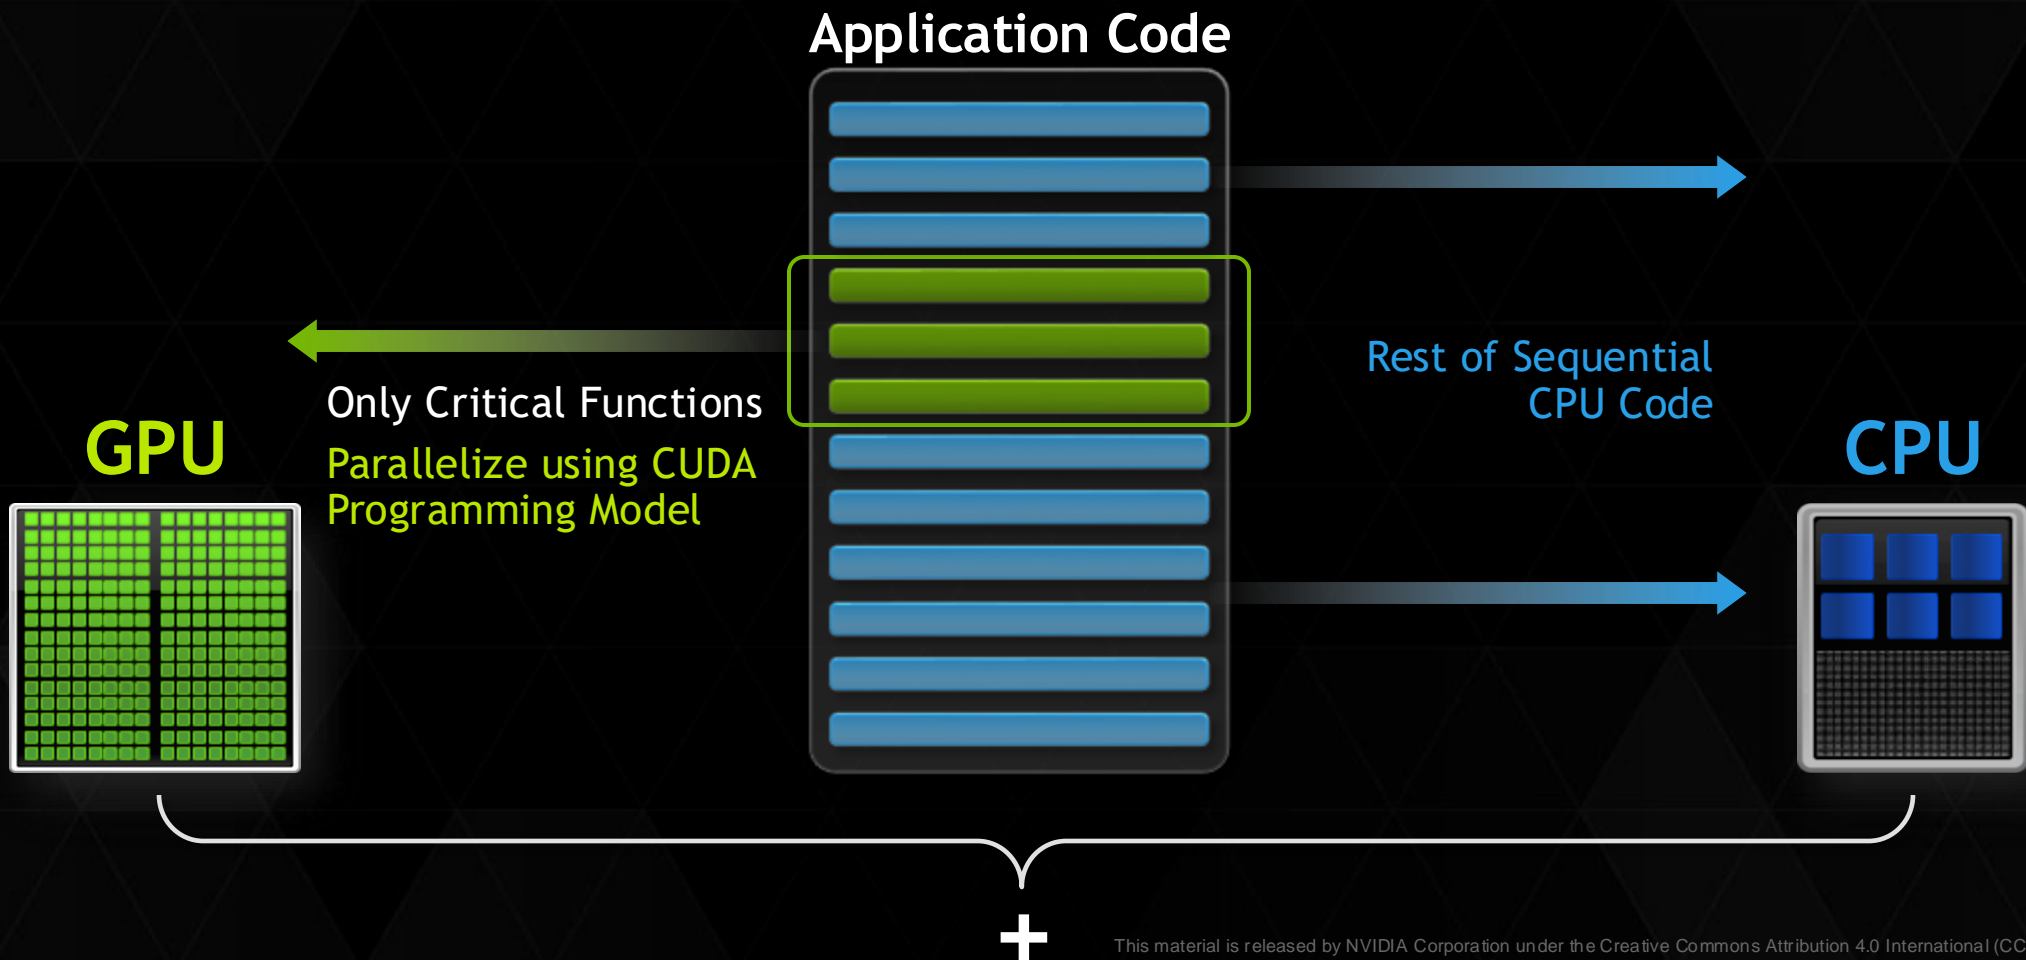


# A QUICK GUIDE TO CUPY

## What to expect?

- What is CuPy?
- Features of CuPy
- Installation Guide
- CuPy Fundamentals
- CUDA Kernels
- Summary

# GPU COMPUTING





# CUDA ARCHITECTURE PROGRAMMING MODEL

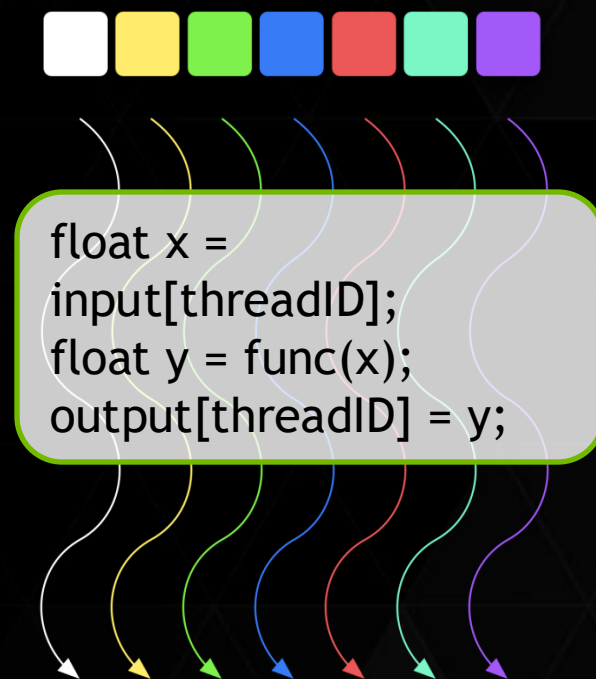
# CUDA KERNELS

- Parallel portion of application: execute as a kernel
  - Entire GPU executes kernel, many threads
- CUDA threads:
  - Lightweight
  - Fast switching
  - Tens of thousands execute simultaneously

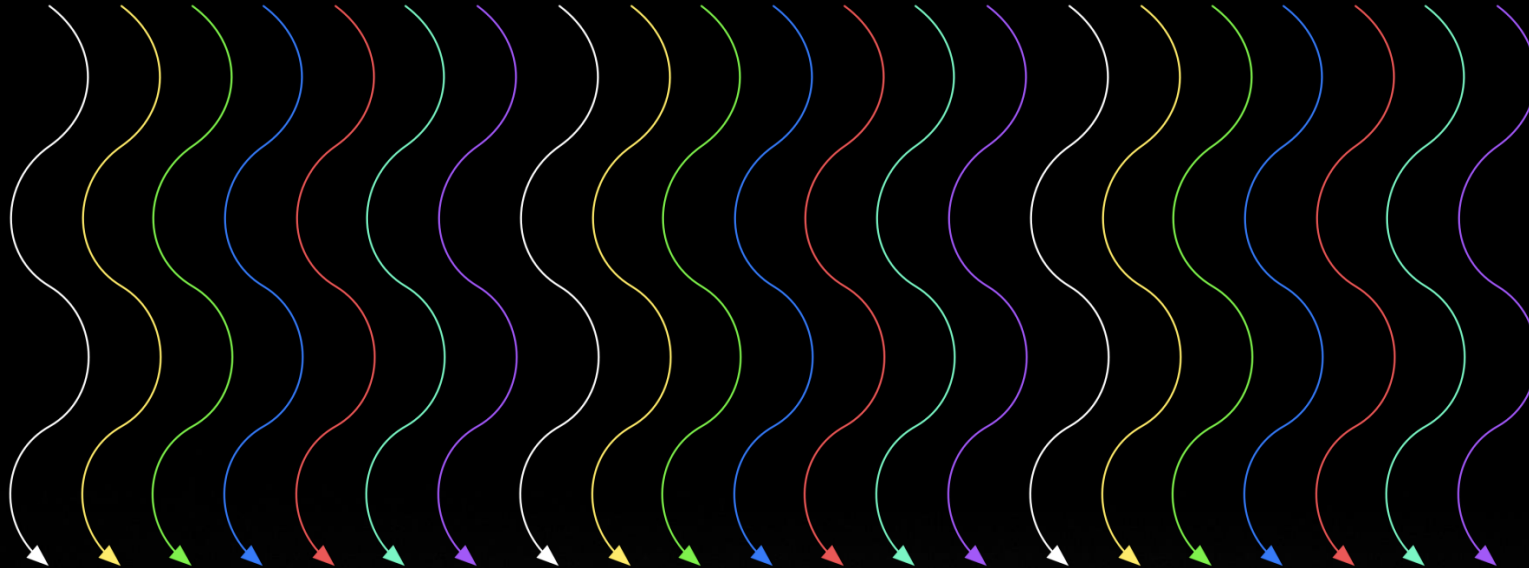
CPU	Host	Executes functions
GPU	Device	Executes kernels

# CUDA KERNELS: PARALLEL THREADS

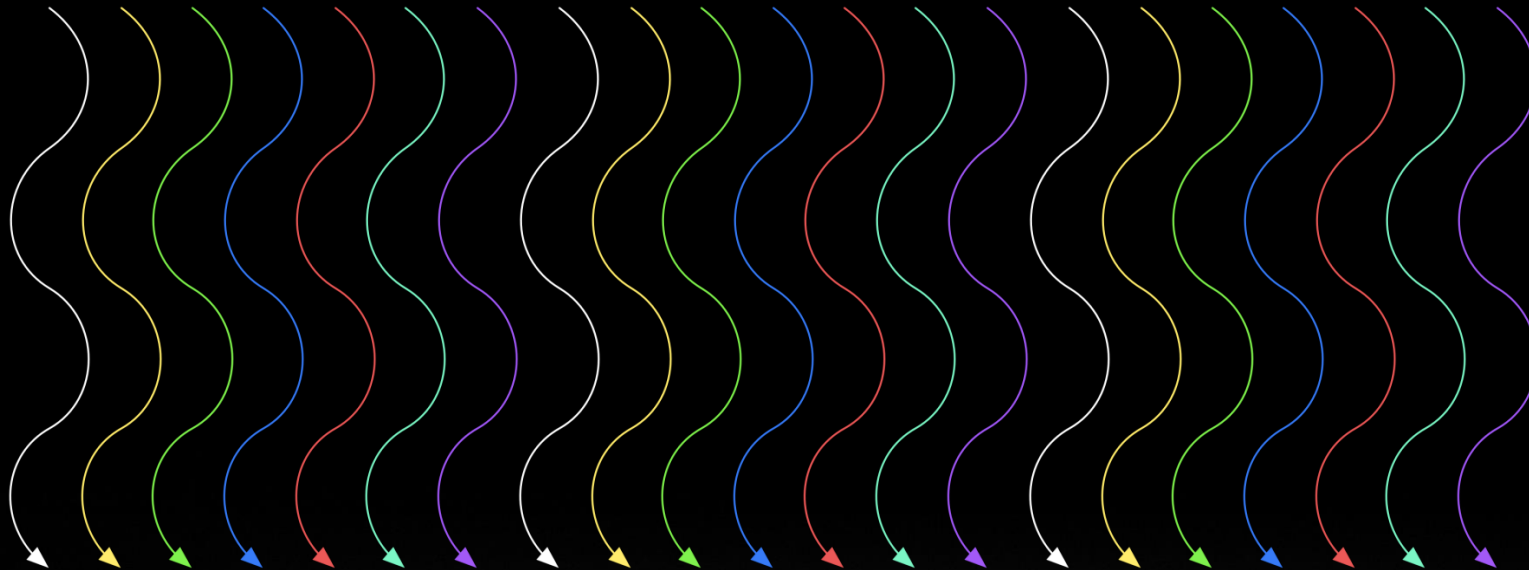
- A **kernel** is a function executed on the GPU
  - Array of threads, in parallel
- All threads execute the same code, can take different paths
  - Each thread has an ID
  - Select input/output data
  - Control decisions



# CUDA KERNELS: SUBDIVIDE INTO BLOCKS



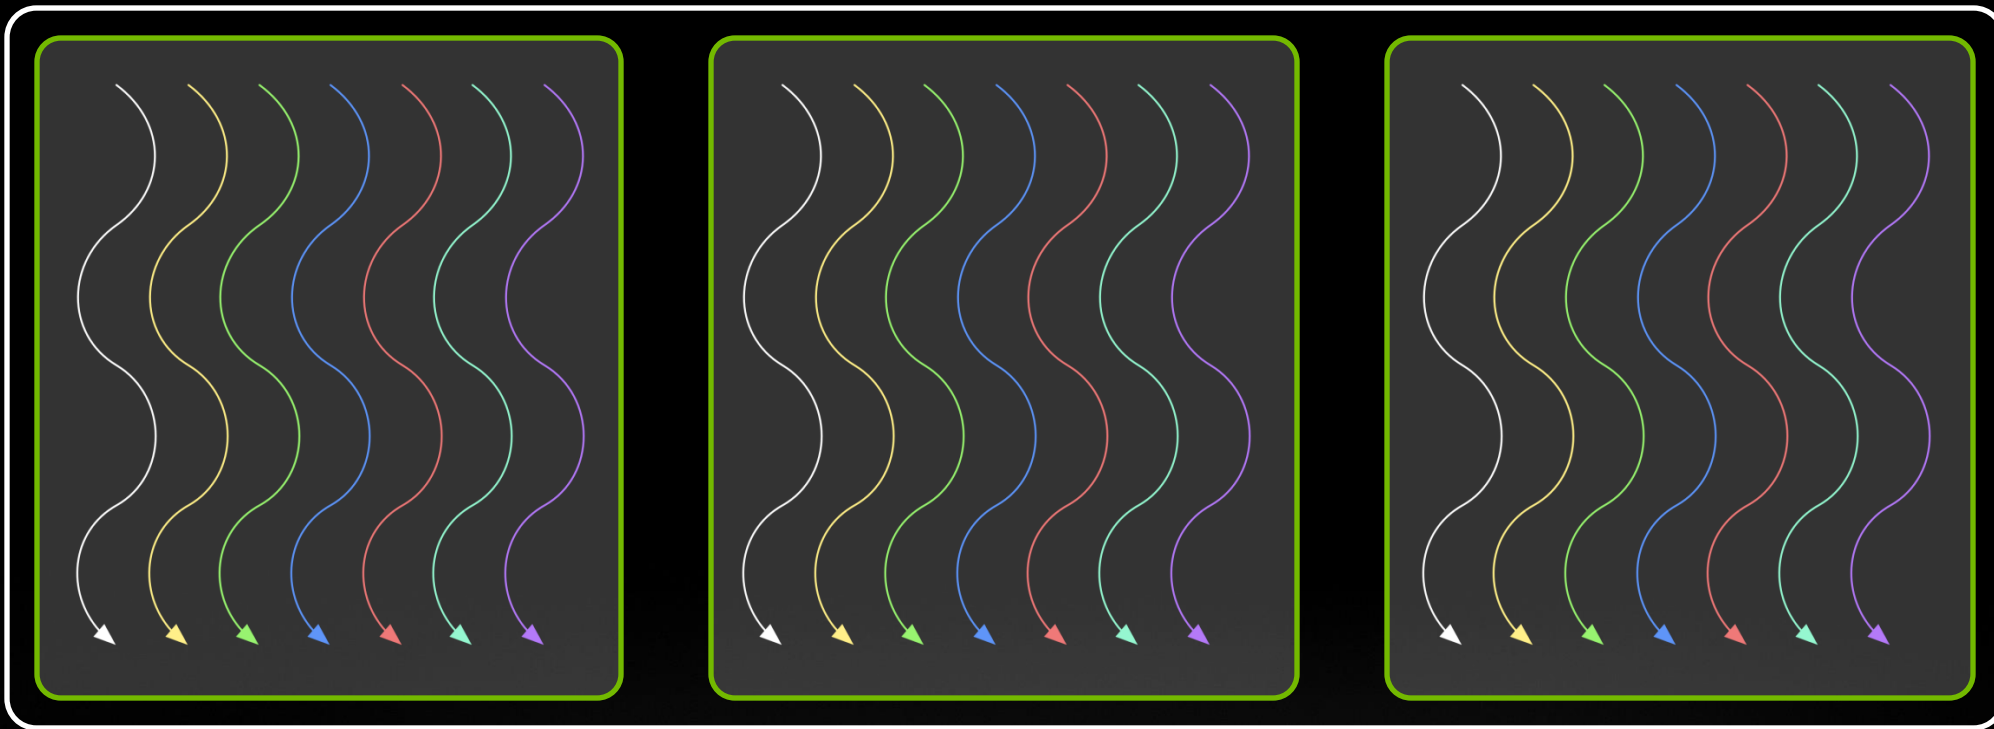
# CUDA KERNELS: SUBDIVIDE INTO BLOCKS



- Threads are grouped into **blocks**




# CUDA KERNELS: SUBDIVIDE INTO BLOCKS



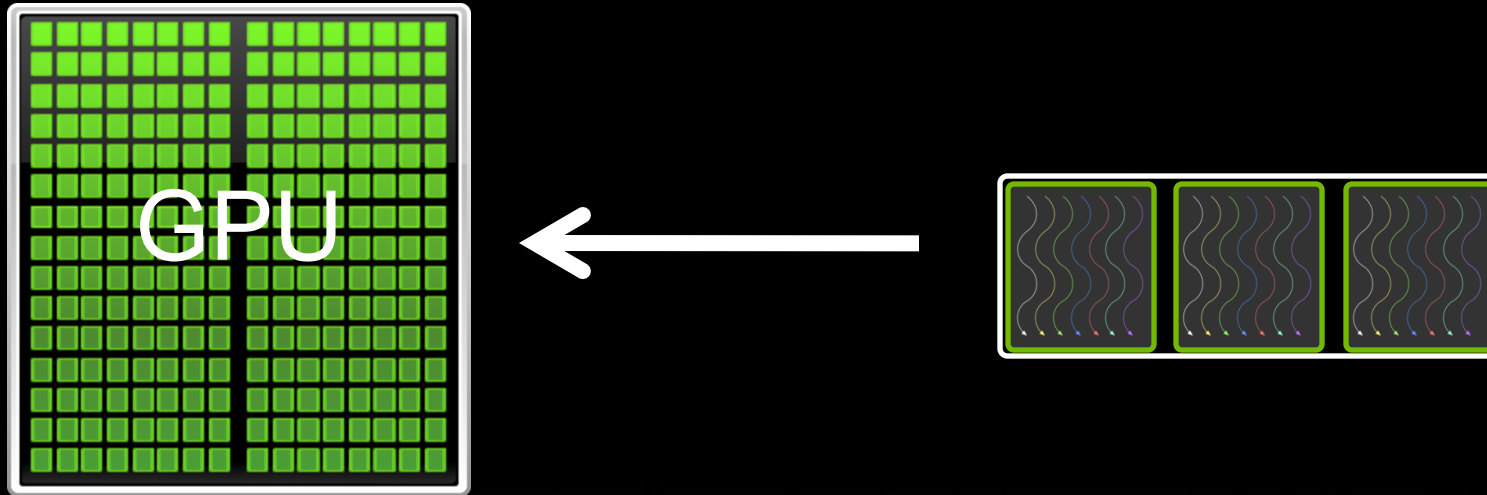
- Threads are grouped into **blocks**
- **Blocks** are grouped into **a grid**

# CUDA KERNELS: SUBDIVIDE INTO BLOCKS



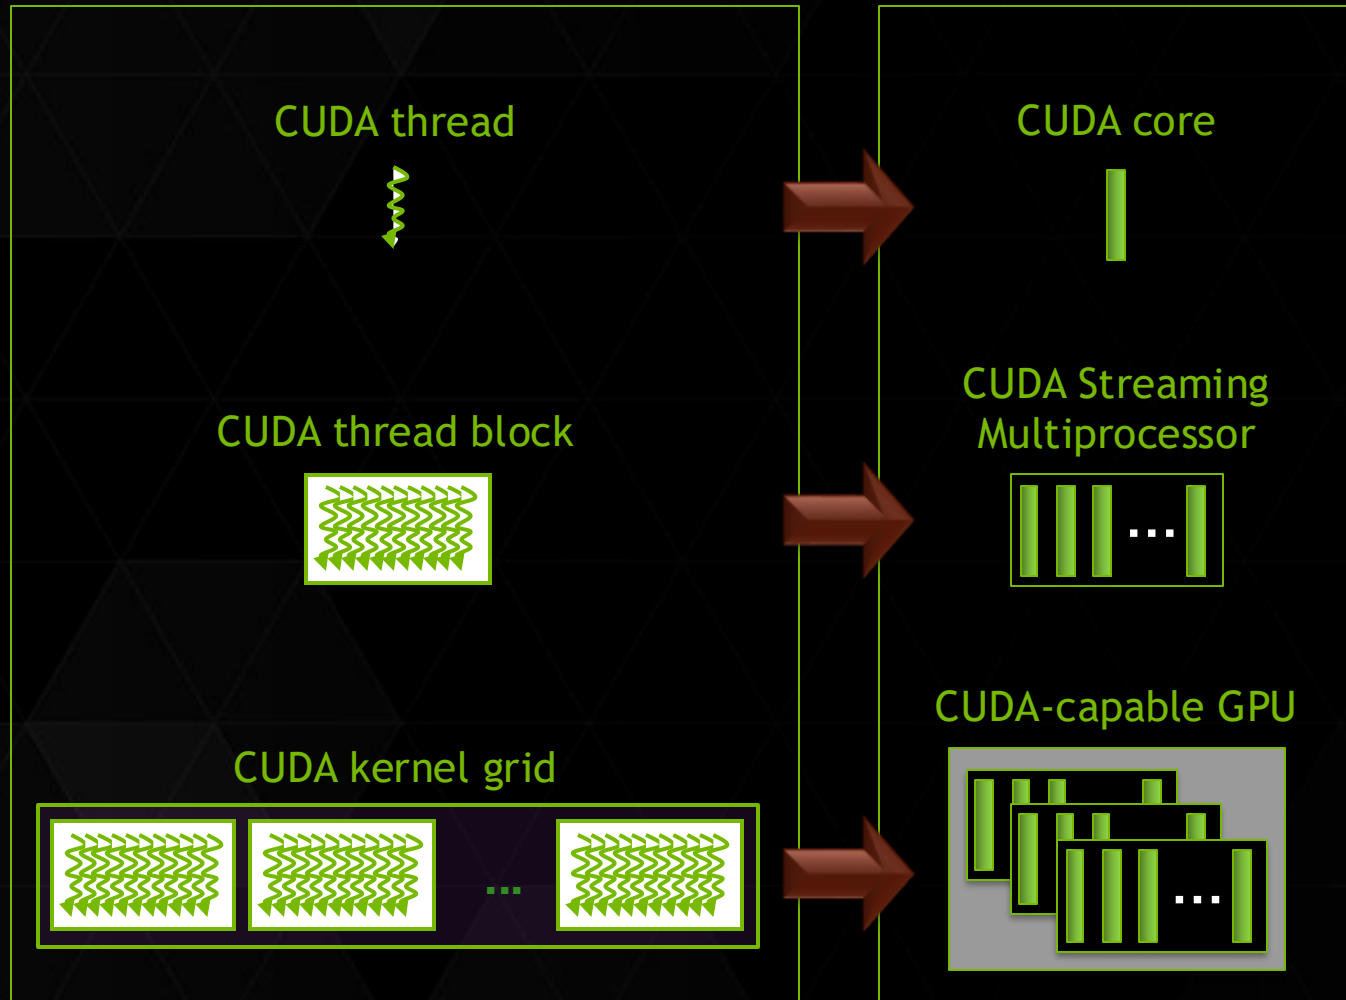
- Threads are grouped into **blocks**
- **Blocks** are grouped into **a grid**
- A **kernel** is executed as a **grid** of **blocks** of **threads**

# CUDA KERNELS: SUBDIVIDE INTO BLOCKS



- Threads are grouped into **blocks**
- **Blocks** are grouped into **a grid**
- A **kernel** is executed as a **grid** of **blocks** of **threads**

# KERNEL EXECUTION

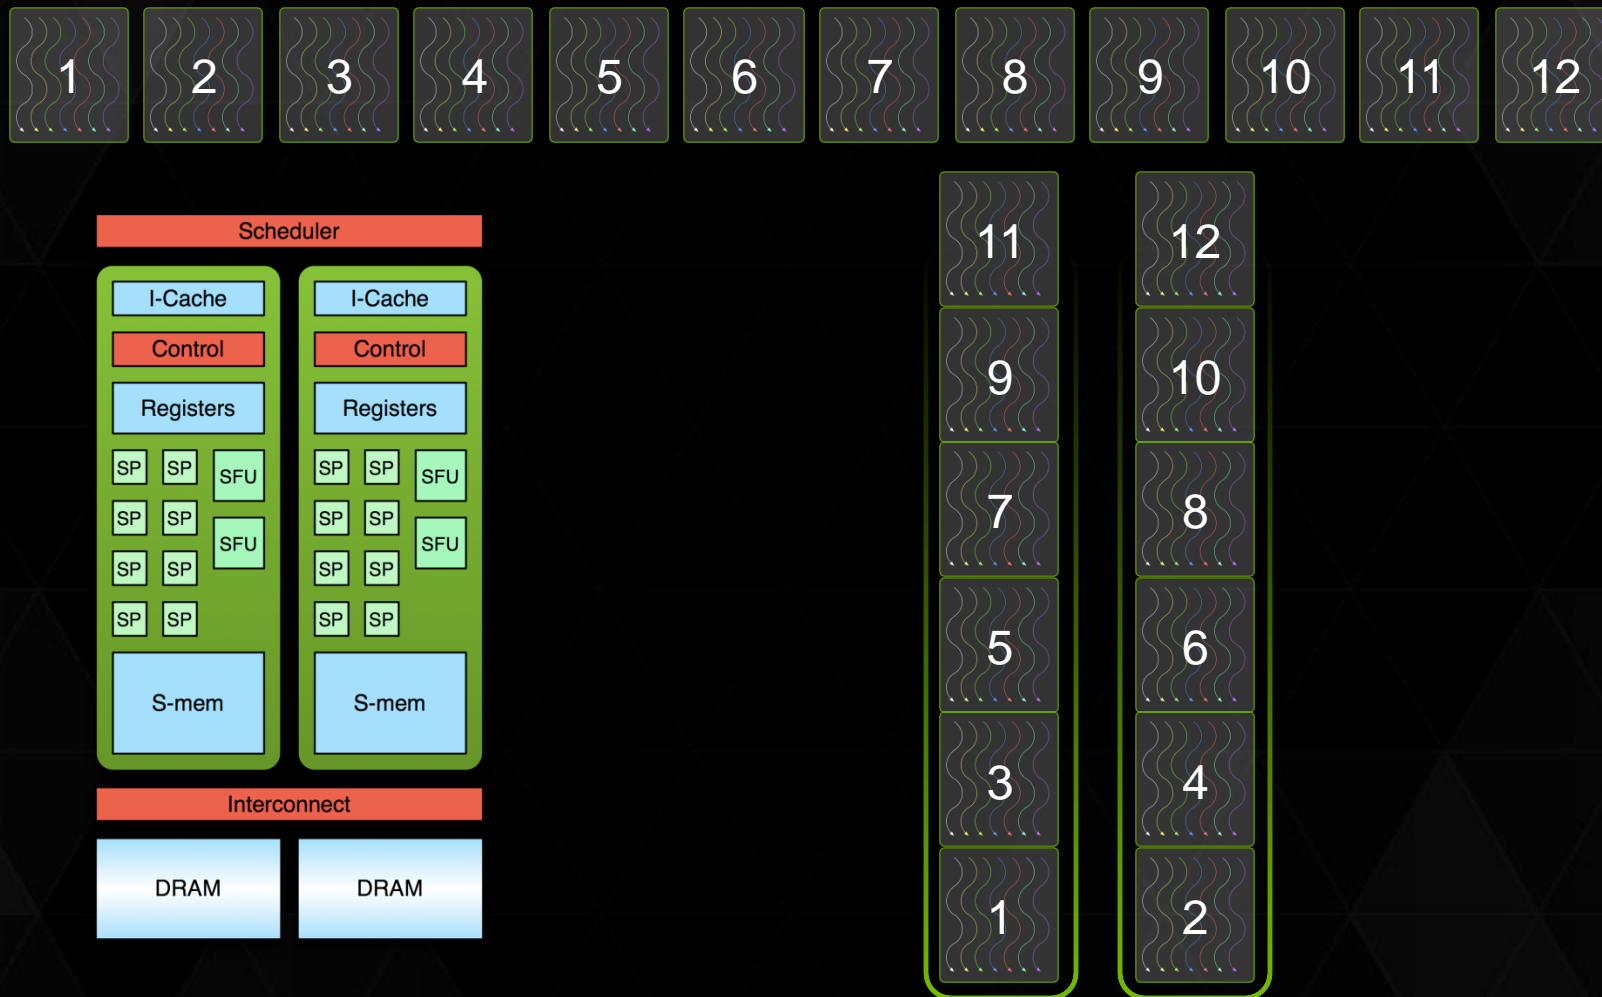


- Each thread is executed by a core
- Each block is executed by one SM and does not migrate
- Several concurrent blocks can reside on one SM depending on the blocks' memory requirements and the SM's memory resources
- Each kernel is executed on one device
- Multiple kernels can execute on a device at one time

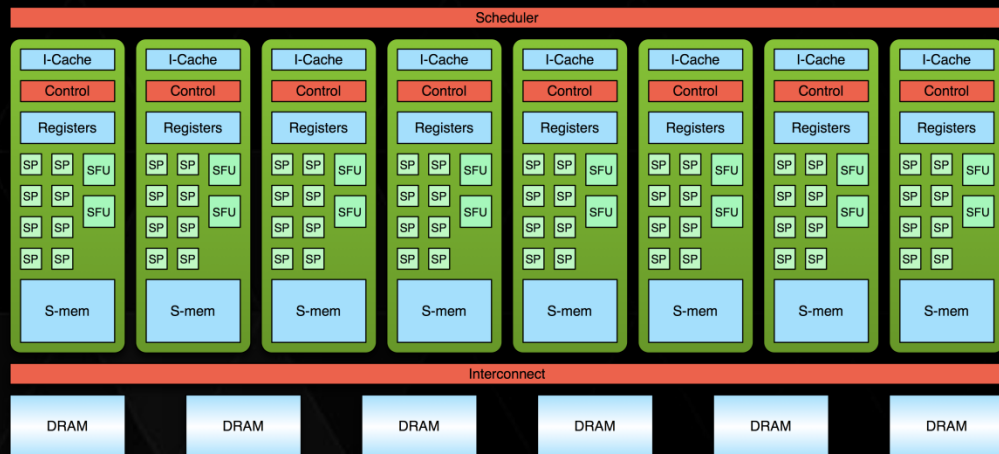
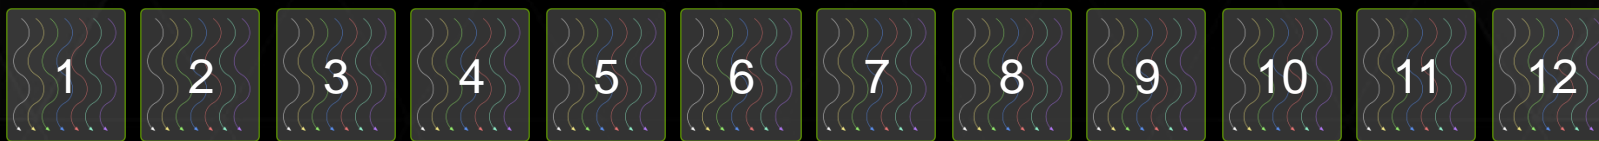
# COMMUNICATION WITHIN A BLOCK

- Threads may need to cooperate
  - Memory accesses
  - Share results
- Cooperate using **shared memory**
  - Accessible by all threads within a block
- Restriction to “within a block” permits scalability
  - Fast communication between  $N$  threads is not feasible when  $N$  large

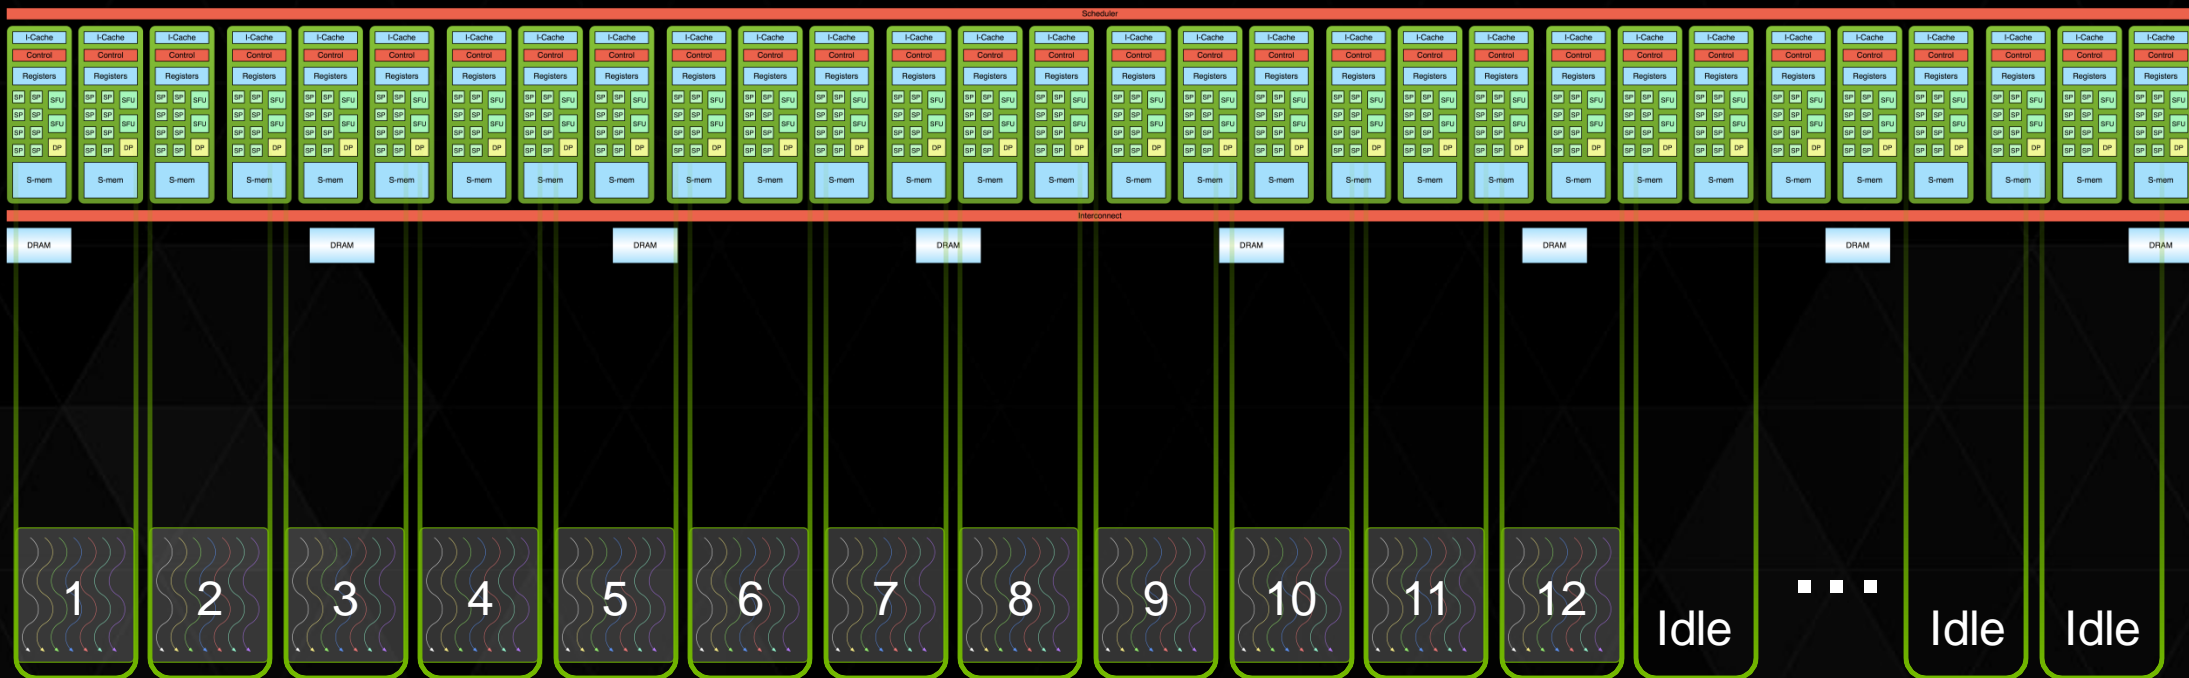

# TRANSPARENT SCALABILITY



# TRANSPARENT SCALABILITY




## TRANSPARENT SCALABILITY -





# CUDA PROGRAMMING MODEL - SUMMARY

- A kernel executes as a grid of thread blocks
- A block is a batch of threads
  - Communicate through shared memory
- Each block has a block ID
- Each thread has a thread ID





# CUDA ARCHITECTURE MEMORY MODEL

# GPU ARCHITECTURE

Two Main components

## Global memory

Analogous to RAM in a CPU server

Accessible by both GPU and CPU


48GB with bandwidth currently up to 1 TB/s

## Streaming Multiprocessors (SMs)

SMs perform the actual computations

Each SM has its own:

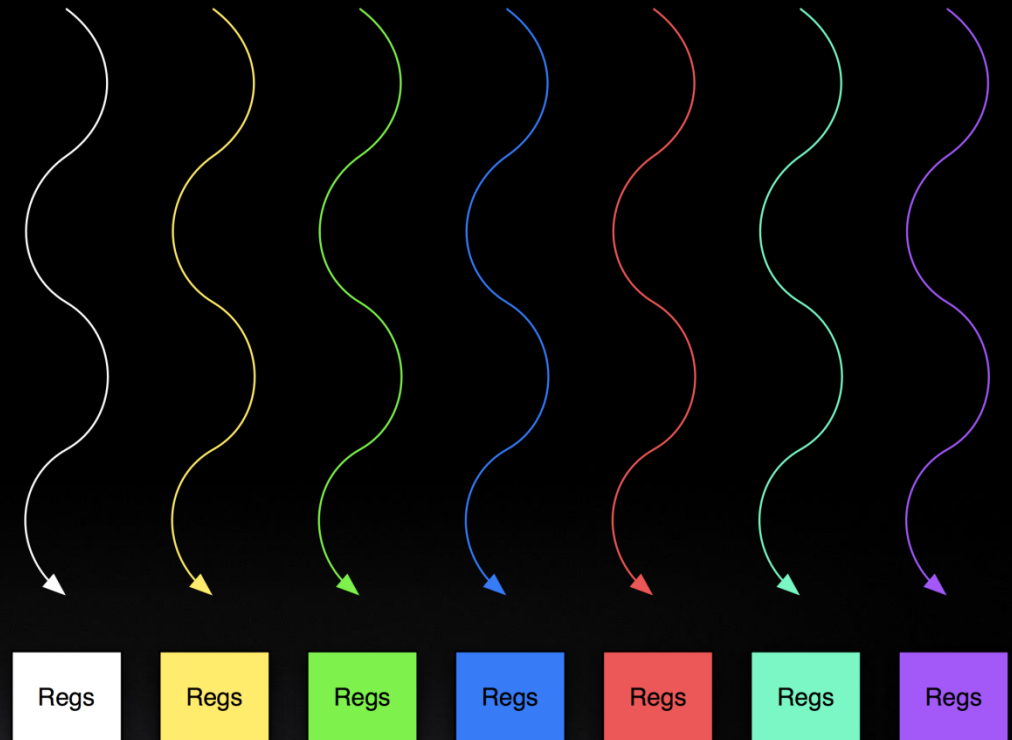
Control units, registers, execution pipelines, caches



# MEMORY HIERARCHY

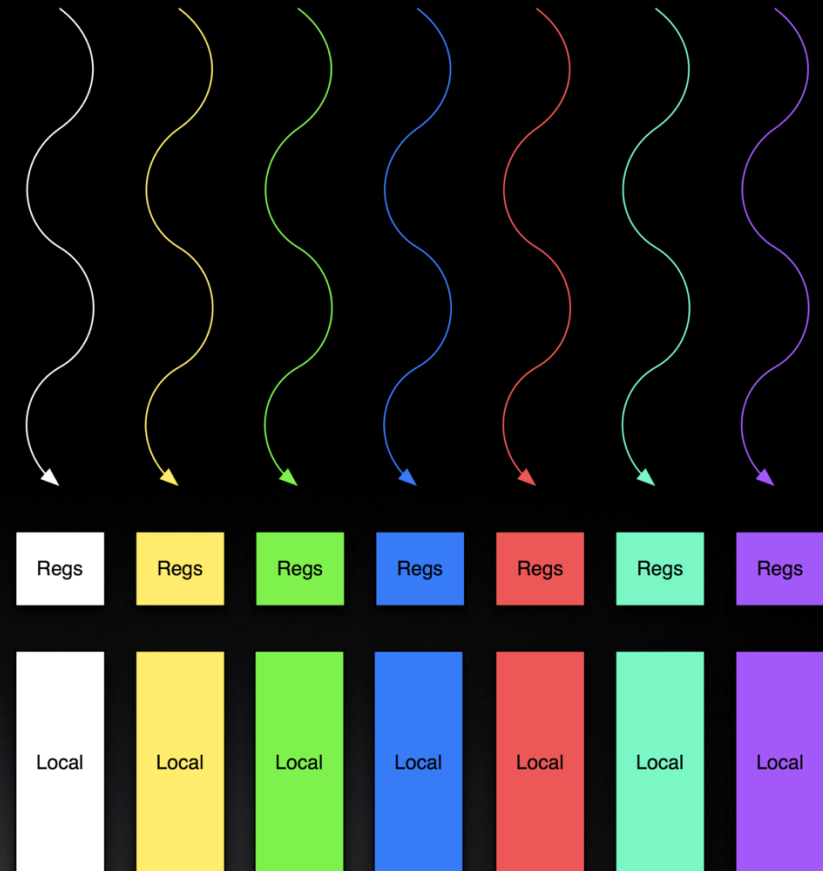
▸ Thread

▸ Registers



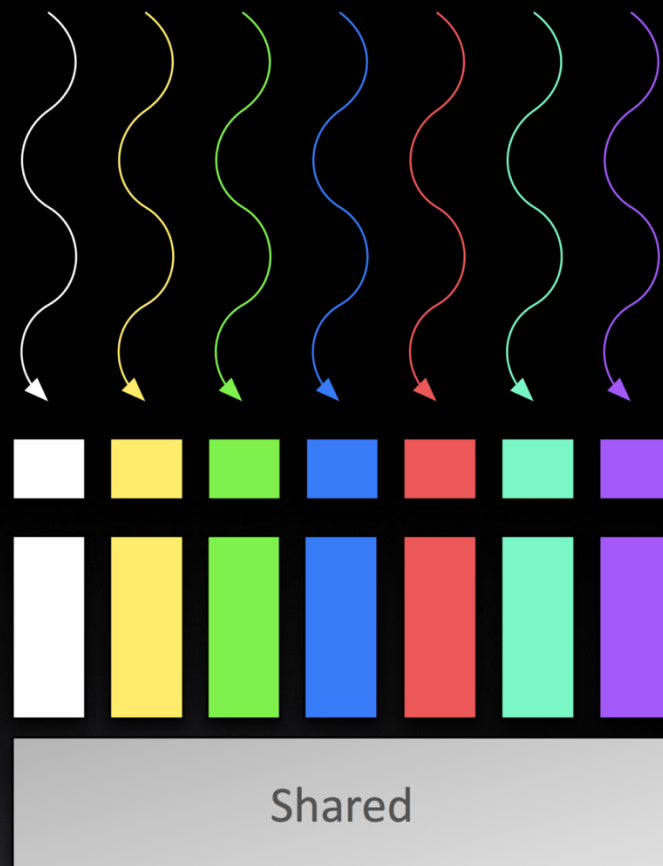
# MEMORY HIERARCHY

- Thread
  - Registers
- Thread
  - Local memory




# MEMORY HIERARCHY

- Thread
  - Registers
- Thread
  - Local memory
- Block of threads
  - Shared memory



# MEMORY HIERARCHY

- Thread
  - Registers
- Thread
  - Local memory
- Block of threads
  - Shared memory
- All blocks
  - Global memory





# WHAT IS CUPY?



# OVERVIEW OF CUPY

- CuPy is an implementation of NumPy-compatible multi-dimensional array on CUDA
- CuPy consists of :
  - ✓ `cupy.ndarray`
  - ✓ the core multi-dimensional array class
  - ✓ many functions

# OVERVIEW OF CUPY

● CuPy supports a subset of numpy.ndarray interface which include:

- ✓ Basic & advance indexing, and Broadcasting
- ✓ Data types (int32, float32, uint64, complex64,... )
- ✓ Array manipulation routine (reshape)
- ✓ Linear Algebra functions (dot, matmul, etc)
- ✓ Reduction along axis (max, sum, argmax, etc)

For more details on broadcasting visit

(<https://numpy.org/doc/stable/user/basics.broadcasting.html>)

```
>>> import numpy as np
>>> X = np.array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
#Basic indexing and slicing
>>> X[5:]
array([5, 6, 7, 8, 9])
>>> X[1:7:2]
array([1, 3, 5])

#Advance indexing
>>> X = np.array([[1, 2],[3, 4],[5, 6]])
>>> X[[0, 1, 2], [0, 1, 0]]
array([1, 4, 5])

#reduction and Linear Algebra function
>>> max(X)
9.0
>>> B = np.array([1,2,3,4], dtype=np.float32)
>>> C = np.array([5,6,7,8], dtype=np.float32)
>>> np.matmul(B, C)
70.0

#data type and array manipulation routine
>>> A = 1j*np.arange(9, dtype=np.complex64).reshape(3,3)
[[0.+0.j 0.+1.j 0.+2.j]
 [0.+3.j 0.+4.j 0.+5.j]
 [0.+6.j 0.+7.j 0.+8.j]]
```



# FEATURES OF CUPY

# FEATURES OF CUPY

- Features of CuPy includes:
  - ✓ User-define elementwise CUDA kernels
  - ✓ User-define reduction CUDA kernels
  - ✓ Fusing CUDA kernels to optimize user-define calculation
  - ✓ Customizable memory allocator and memory pool
  - ✓ cuDNN utilities
- These features are developed to support performance.
- **CuPy uses on-the-fly kernel synthesis:** when a kernel call is required, it compiles a kernel code optimized for the shapes and dtypes of given arguments, sends it to the GPU device, and executes the kernel.
- **CuPy** also caches the kernel code sent to GPU device within the process, which reduces the kernel transfer time on further calls.



# INSTALLATION GUIDE

# REQUIREMENTS

- Recommended Linux distributions are Centos and Ubuntu
- ✓ NVIDIA CUDA GPU with the Compute Capability 3.0 or larger
- ✓ CUDA Toolkit: v9.0 - v11.2
- ✓ Python: v3.5.1+ - v3.9.0+
- **Python Dependencies**
- ✓ NumPy/SciPy-compatible API in CuPy v8 is based on NumPy 1.19 and SciPy 1.5.

## 1. Wheels (precompiled binary package)

CUDA	Command
v9.0	<code>\$ pip install cupy-cuda90</code>
v9.2	<code>\$ pip install cupy-cuda92</code>
v10.0	<code>\$ pip install cupy-cuda100</code>
v10.1	<code>\$ pip install cupy-cuda101</code>
v10.2	<code>\$ pip install cupy-cuda102</code>
v11.0	<code>\$ pip install cupy-cuda110</code>
v11.1	<code>\$ pip install cupy-cuda111</code>

## 2. Conda-Forge

- ✓ `$ conda install -c conda-forge cupy`
- To install CuPy with the cuTENSOR support enabled, you can do:
- ✓ `$ conda install -c conda-forge cupy cutensor cudatoolkit=10.2`

## 3. CuPy inside Docker

You can pull CuPy Docker images from:

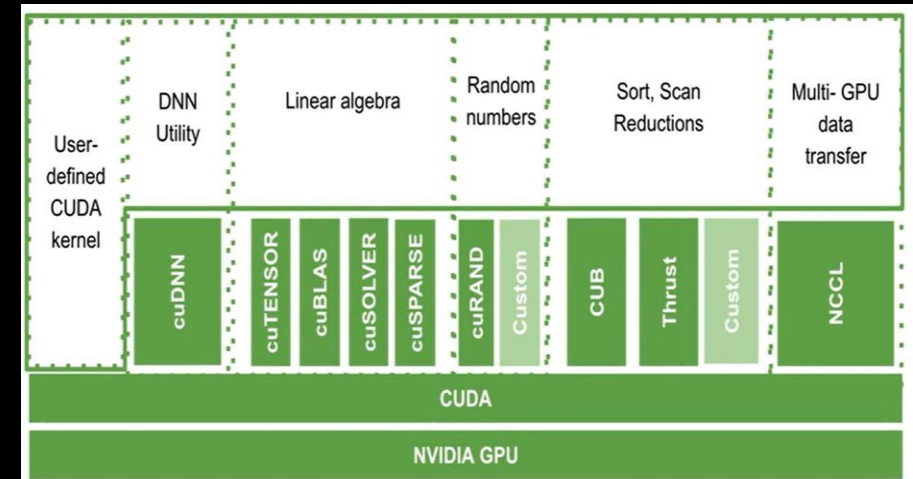
<https://hub.docker.com/r/cupy/cupy/>.

# REQUIREMENTS & ARCHITECTURE

## 4. Full RAPIDS Package

	Preferred	Advanced
METHOD	Conda	Docker + Examples Docker + Dev Env Source
RELEASE	Stable (0.18)	Nightly (0.19a)
TYPE	RAPIDS and BlazingSQL	RAPIDS Core (w/o BlazingSQL)
PACKAGES	All Packages cuDF cuML cuGraph cuSignal cuSpatial cuxfilter	
LINUX	Ubuntu 16.04 Ubuntu 18.04 Ubuntu 20.04	CentOS 7 CentOS 8 RHEL 7&8
PYTHON	Python 3.7	Python 3.8
CUDA	CUDA 10.1.2 CUDA 10.2	CUDA 11.0
COMMAND	<p><b>NOTE:</b> Ubuntu 16.04/18.04/20.04 &amp; CentOS 7/8 use the same <code>conda create</code> commands.</p> <pre>conda create -n rapids-0.18 -c rapidsai -c nvidia -c conda-forge \ -c defaults rapids-blazing=0.18 python=3.7 cudatoolkit=10.1</pre>	

## CuPy Architecture








# CUPY FUNDAMENTALS



# TERMINOLOGY

- *Host* - The CPU and its memory (host memory)
- *Device* - The GPU and its memory (device memory)
- *Kernel* - A GPU function launched by the host and executed on the device.

Host



Device



# CUPY.NDARRAY

- CuPy is a GPU array backend that implements a subset of NumPy interface

## CuPy

```
import cupy as cp

x_gpu = cp.array([1, 2, 3, 4, 5])
```

## NumPy

```
import numpy as np

x = np.array([1, 2, 3, 4, 5])
```

## Current device (GPU ID: 0)

```
import cupy as cp

gpu_0 = cp.array([1, 2, 3, 4, 5])

# Switch device
cp.cuda.Device(1).use()
gpu_1 = cp.array([1, 2, 3, 4])
```

## Switch GPU temporarily

```
import numpy as np
import cupy as cp

with cp.cuda.Device(1):
    gpu_1 = cp.array([1, 2, 3, 4])

# back to device id 0
gpu0 = cp.array([1, 2, 3, 4, 5])
```

# DATA TRANSFER

## Host → Device using `cupy.asarray`.

```
import cupy as cp
import numpy as np

X = np.array([1, 2, 3, 4, 5])
x_gpu = cp.asarray(x)
print(x_gpu)
Output:[1 2 3 4 5]
```

## Device → Host using `cupy.asnumpy` or `cupy.ndarray.get()`

```
import cupy as cp
import numpy as np

X_gpu = cp.array([1, 2, 3, 4, 5])
# copy to Host
x_cpu = cp.asnumpy(x_gpu)
print(x_cpu)
[1 2 3 4 5]

#alternative option
x_cpu_alt = x_gpu.get()
x_cpu_alt
Output:[1 2 3 4 5]
```

## Devices(GPU to GPU)

```
import cupy as cp
with cp.cuda.Device(0):
    x_gpu_0 = cp.ndarray([ 2, 3, 3])
x_gpu_0
[[[0. 0. 0.]
  [0. 0. 0.]
  [0. 0. 0.]]

  [[0. 0. 0.]
  [0. 0. 0.]
  [0. 0. 0.]]]

with cp.cuda.Device(1):
    x_gpu_1 = cp.asarray(x_gpu_0)
x_gpu_1
[[[0. 0. 0.]
  [0. 0. 0.]
  [0. 0. 0.]]

  [[0. 0. 0.]
  [0. 0. 0.]
  [0. 0. 0.]]]
```

# GPU & CPU AGNOSTIC CODE

- Using `cupy.get_array_module()`

```
>>> import cupy as cp
>>> import numpy as np
>>>
>>> #example: log(1 + exp(x))
>>> x_cpu = np.array([1, 2, 3, 4, 5])
>>> x_gpu = cp.get_array_module(x_cpu)
>>> result = x_gpu.maximum(0, x_cpu) + x_gpu.log1p(x_gpu.exp(-abs(x_cpu)))
>>> result
Output: [1.31326169 2.12692801 3.04858735 4.01814993 5.00671535]
>>>
>>>
>>> #An explicit conversion to host
>>> x_gpu = cp.array([6, 7, 8, 9, 10])
>>> result = cp.asnumpy(x_gpu) + x_cpu
>>> result
Output: [ 7  9 11 13 15]
>>> #An explicit conversion to device
>>> result = x_gpu + cp.asarray(x_cpu)
>>> result
Output: [ 7  9 11 13 15]
>>>
```

# HOW MUCH FASTER IS CUPY THAN NUMPY ?

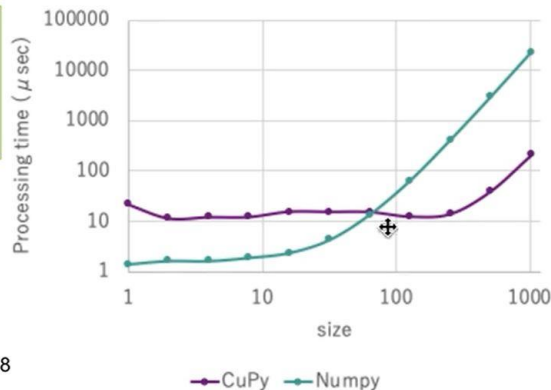
## Dot product

```
a = xp.ones((size, size), 'f')
b = xp.ones((size, size), 'f')

def f():
    xp.dot(a, b)
```

For a rough estimation, if the array size is larger than L1 cache of your CPU, CuPy gets faster than NumPy.

Try on Google Colab! <http://bit.ly/cupywest2018>



## Add Function

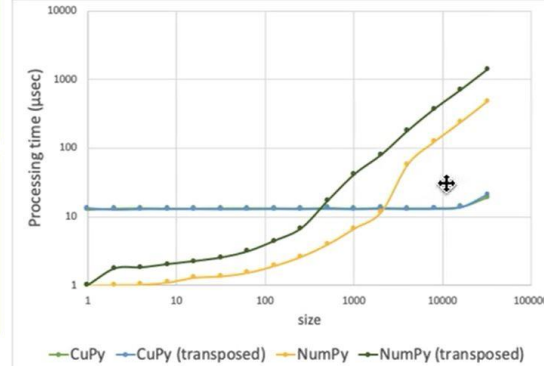
```
a = xp.ones((size, 32), 'f')
b = xp.ones((size, 32), 'f')

def f():
    a + b
```

```
# Transpose
a = xp.ones((32, size), 'f').T
b = xp.ones((size, 32), 'f')
```

```
def f():
    a + b
```

Xeon Gold 6154 CPU @ 3.00GHz  
Tesla V100-PCIE-16GB





# CUDA KERNELS

# CUPY CUDA KERNELS

- CUDA Kernels can be defined in Cupy as follows:
  - ✓ Elementwise Kernels
  - ✓ Reduction Kernels
  - ✓ Raw Kernels
  - ✓ Kernel Fusion
- These kernels are user-defined based.

# ELEMENTWISE KERNEL

- The **ElementwiseKernel** class is used to define this type of kernel.
- This kernel consists of four parts which includes:
  - ✓ a list of input argument
  - ✓ a list of output argument
  - ✓ a loop body code
  - ✓ a kernel name
- Variable name starting with underscore “\_”, “n”, and “i” are regarded as reserved keywords.



# ELEMENTWISE KERNEL

## ● Example : $z = x * w + b$

```
import cupy as cp
input_list = 'float32 x , float32 w, float32 b'
output_list = 'float32 z'
code_body = 'z = (x * w) + b'
# elementwisekernel class defined

dnnLayerNode = cp.ElementwiseKernel(input_list, output_list, code_body, 'dnnLayerNode')

x = cp.arange(9, dtype=cp.float32).reshape(3,3)
w = cp.arange(9, dtype=cp.float32).reshape(3,3)
b = cp.array([-0.5], dtype=cp.float32)
z = cp.empty((3,3), dtype=cp.float32)

# kernel call with argument passing

dnnLayerNode(x,w,b,z)

print(z)

#output
[[-0.5  0.5  3.5]
 [ 8.5 15.5 24.5]
 [35.5 48.5 63.5]]
```

Diagram annotations:

- Data type (points to float32 in input\_list)
- Input argument list (points to x, w, b in input\_list)
- Output argument list (points to z in output\_list)
- A loop body code (points to the code\_body string)

# ELEMENTWISE KERNEL: GENERIC-TYPE KERNELS

## ● Example : $z = x * w + b$

```
import cupy as cp
input_list = 'T x , T w, T b'
output_list = 'T z'

code_body = 'z = (x * w) + b'

# elementwisekernel class defined

dnnLayerNode = cp.ElementwiseKernel(input_list, output_list, code_body, 'dnnLayerNode')

x = cp.arange(9, dtype=cp.float32).reshape(3,3)
w = cp.arange(9, dtype=cp.float32).reshape(3,3)
b = cp.array([-0.5], dtype=cp.float32)
z = cp.empty((3,3), dtype=cp.float32)

# kernel call with argument passing

dnnLayerNode(x,w,b,z)

print(z)

#output
[[-0.5  0.5  3.5]
 [ 8.5 15.5 24.5]
 [35.5 48.5 63.5]]
```

Generic type placeholder

### Multiple generic placeholder

```
import cupy as cp
input_list = 'T x , W w, B b'
output_list = 'T z'
.....
.....
.....
#output
.....
.....
```

Different types of placeholder

# REDUCTION KERNEL

- Reduction kernel is implemented through the **ReductionKernel** class.
- In order to implement this kernel class, the following parts must be defined:
  - ✓ **Identity value**: to initialize reduction value.
  - ✓ **Mapping expression**: Used for the pre-processing of each element to be reduced.
  - ✓ **Reduction expression**: It is an operator to reduce the multiple mapped values. The special variables **a** and **b** are used for its operands.
  - ✓ **Post mapping expression**: It is used to transform the resulting reduced values. The special variable **a** is used as its input. Output should be written to the output parameter.

# REDUCTION KERNEL

Example:  $z = \sum_{i=1} x_i w_i + b$

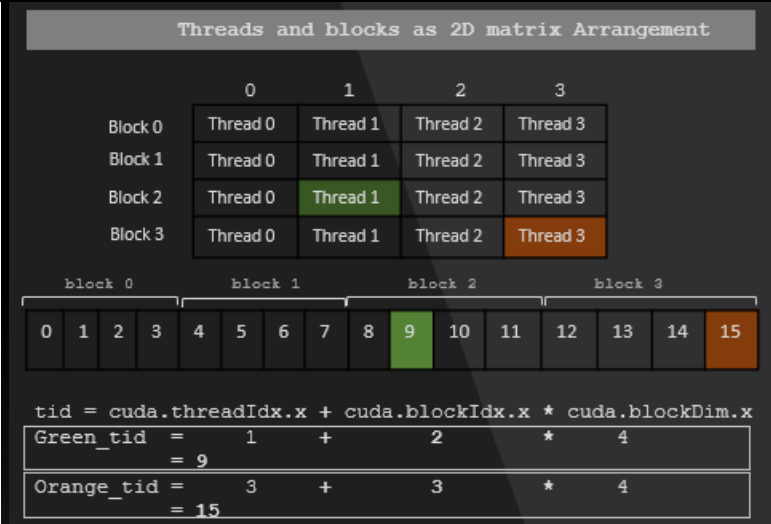
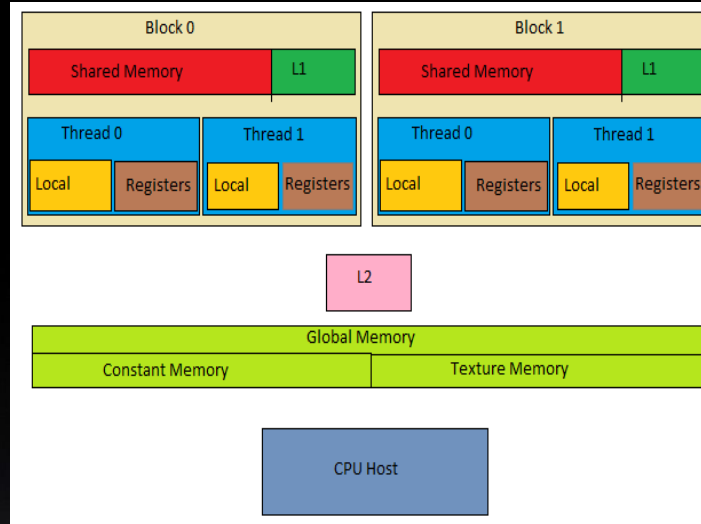
```
import cupy as cp
dnnLayer = cp.ReductionKernel(
    'T x, T w, T bias', ← input params. The bias represents b from the above equation.
    'T z', ← output params
    'x * w', ← map
    'a + b', ← reduce
    'z = a + bias', ← post-reduction map
    '0', ← identity value
    'dnnLayer' ← kernel name
)
x = cp.arange(10, dtype=cp.float32).reshape(2,5)
w = cp.arange(10, dtype=cp.float32).reshape(2,5)
bias = -0.1
z = dnnLayer(x,w,bias) ← kernel call
print(z)
#output
284.9
```

inputs

# RAW KERNEL

- Raw kernels enable the direct use of kernels from CUDA source, and it is defined through the *RawKernel* class.
- The *RawKernel* object allows you to call the kernel with **CUDA's cuLaunchKernel interface**. In other words, you

have control over:



- ✓ grid size
- ✓ block size
- ✓ shared memory size
- ✓ and stream.

# RAW KERNEL EXAMPLE

```
import cupy as cp

add_kernel = cp.RawKernel(r'''
extern "C" __global__
void add_func(const float* x1, const float* x2, float* y) {
    int tid = blockDim.x * blockIdx.x + threadIdx.x;
    y[tid] = x1[tid] + x2[tid];
}
''', 'add_func')
```

```
N = 100
shape = (10, 10)
x1 = cp.arange(N, dtype=cp.float32).reshape(shape)
x2 = cp.arange(N, dtype=cp.float32).reshape(shape)
y = cp.zeros((shape), dtype=cp.float32)
```

```
add_kernel((10,), (10,), (x1, x2, y))
```

grid size    block size    arguments

#output

```
[[ 0.  2.  4.  6.  8. 10. 12. 14. 16. 18.]
 [20. 22. 24. 26. 28. 30. 32. 34. 36. 38.]
 [40. 42. 44. 46. 48. 50. 52. 54. 56. 58.]
 [60. 62. 64. 66. 68. 70. 72. 74. 76. 78.]
 [80. 82. 84. 86. 88. 90. 92. 94. 96. 98.]
[100. 102. 104. 106. 108. 110. 112. 114. 116. 118.]
[120. 122. 124. 126. 128. 130. 132. 134. 136. 138.]
[140. 142. 144. 146. 148. 150. 152. 154. 156. 158.]
[160. 162. 164. 166. 168. 170. 172. 174. 176. 178.]
[180. 182. 184. 186. 188. 190. 192. 194. 196. 198.]]
```

This also yield the same output:

```
add_kernel((1,), (100,), (x1, x2, y))
```

# RAW MODULES

- The **RawModule** class is used to defining a large raw CUDA C source or loading an existing CUDA binary.
- It is initialized by a CUDA C source code having several kernels (functions) such that needed kernels are retrieved by calling **the `get_function()` method**.

```
import cupy as cp
loaded_from_source = r'''
extern "C" {
__global__ void sum(const float* A, const float* B, float* C, int N)
{
    int tid = blockDim.x * blockIdx.x + threadIdx.x;
    if(tid < N)
    {
        C[tid] = A[tid] + B[tid];
    }
}

__global__ void multiply(const float* A, const float* B, float* C, int N)
{
    int tid = blockDim.x * blockIdx.x + threadIdx.x;
    if(tid < N)
    {
        C[tid] = A[tid] * B[tid];
    }
}
}'''
```

# EXAMPLE OF RAW MODULE

```
import cupy as cp
loaded_from_source = r'''
extern "C" {
__global__ void sum_ker(const float* a, const float* b, float* c, int N)
{
    int tid = blockDim.x * blockIdx.x + threadIdx.x;
    if(tid < N)
    {
        c[tid] = a[tid] + b[tid];
    }
}
__global__ void multiply_ker(const float* a, const float* b, float* c, int N)
{
    int tid = blockDim.x * blockIdx.x + threadIdx.x;
    if(tid < N)
    {
        c[tid] = c[tid] * b[tid];
    }
}
}'''
```

```
ker_sum((1,), (25,), (a, b, c))
print(y)
##output1
[[ 1.  2.  3.  4.  5.]
 [ 6.  7.  8.  9. 10.]
 [11. 12. 13. 14. 15.]
 [16. 17. 18. 19. 20.]
 [21. 22. 23. 24. 25.]]
```

```
ker_times((5,), (5,), (a, b, c))
print(y)
##output2
[[ 0.  1.  2.  3.  4.]
 [ 5.  6.  7.  8.  9.]
 [10. 11. 12. 13. 14.]
 [15. 16. 17. 18. 19.]
 [20. 21. 22. 23. 24.]]
```

```
Module      = cp.RawModule(code = load_raw_module) ← Loading raw module
ker_sum     = module.get_function('sum_ker') ← Accessing the sum_ker kernel
ker_times   = module.get_function('multiply_ker') ← Accessing the multiple_ker
```

```
a = cp.arange(25, dtype=cp.float32).reshape(5,5)
b = cp.ones((5,5), dtype=cp.float32)
c = cp.zeros((5,5), dtype=cp.float32)
```

← Creating arguments (input and output parameters)



# KERNEL FUSION

- **Kernel fusion** is a decorator that fuses functions. It can be used to define an **elementwise** or **reduction** kernels easily.

```
import cupy as cp
```

```
@cp.fuse(kernel_name='dnnlayerNode') ← decorator
```

```
def dnnlayerNode(x, w, bias):
```

```
    return (x * w) + bias
```

Function scope

```
x = cp.arange(9, dtype=cp.float32).reshape(3,3)
```

```
w = cp.arange(9, dtype=cp.float32).reshape(3,3)
```

```
bias = cp.array([-0.5], dtype=cp.float32)
```

```
z = dnnlayerNode(x, w, bias)
```

```
print(z)
```

```
#output
```

```
[[ -0.5   0.5   3.5]
```

```
 [ 8.5  15.5  24.5]
```

```
 [35.5  48.5  63.5]]
```

```
import cupy as cp
```

```
@cp.fuse
```

```
def sumlayer(x, w):
```

```
    return cp.sum(x * w, axis = -1)
```

```
x = cp.arange(10, dtype=cp.float32)
```

```
w = cp.arange(10, dtype=cp.float32)
```

```
z = sumlayer(x,w)
```

```
print(z)
```

```
#output
```

```
285.0
```

# KERNEL FUSION: MERITS & DEMERITS


## ● Merits

- ✓ Relax the bandwidth bottleneck
- ✓ Reduce memory consumption
- ✓ Speedup function calls

## ● Demerits

- ✓ No support for `cupy.matmul()` and `cupy.reshape()` operations
- ✓ Support only `reduction` and `element-wise` operations

# CUPY SPEEDUP





# SUMMARY

# SUMMARY

- CuPy is an implementation of NumPy-compatible multi-dimensional array on CUDA

## Installation

- ✓ Wheels (precompiled binary package)
- ✓ Conda-Forge
- ✓ CuPy inside Docker
- ✓ Conda (full RAPIDS package)

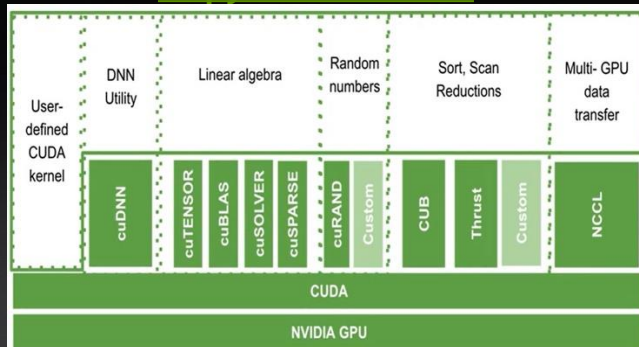
## Data Movement

- ✓ Host to Device (CPU → GPU) using `cupy.asarray`.
- ✓ Device to Host (GPU → CPU) using `cupy.asnumpy` or `cupy.ndarray.get()`
- ✓ between devices (GPU to GPU), `cupy.ndarray` is used.

## Cupy User-defined Kernels

- ✓ Elementwise Kernels
- ✓ Reduction Kernels
- ✓ Raw Kernels
- ✓ Kernel Fusion

## Cupy Architecture



## You want to save GPU memory?

```
import cupy as cp
size = 32768
a = cp.ones((size, size)) # 8GB
b = cp.ones((size, size)) # 8GB
cp.dot(a, b)               # 8GB
```



Traceback (most recent call last):

```
...
cupy.cuda.memory.OutOfMemoryError: out of memory to
allocate 8589934592 bytes (total 17179869184 bytes)
```

## Try Unified Memory! (Supported only on V100)

- Just edit 2 lines to enable unified memory

```
import cupy as cp
```

```
pool = cp.cuda.MemoryPool(cp.cuda.malloc_managed)
cp.cuda.set_allocator(pool.malloc)
```

```
size = 32768
a = cp.ones((size, size)) # 8GB
b = cp.ones((size, size)) # 8GB
cp.dot(a, b)               # 8GB
```

# REFERENCES

- <https://developer.nvidia.com/hpc-sdk>
- <https://docs.cupy.dev/en/stable/>
- <https://cupy.dev/>
- CuPy Documentation Release 8.5.0, Preferred Networks, inc. and Preferred Infrastructure inc., Feb 26, 2021.
- Bhaumik Vaidya, Hands-On GPU-Accelerated Computer Vision with OpenCV and CUDA, Packt Publishing, 2018.
- Crissman Loomis and Emilio Castillo, CuPy Overview: NumPy Syntax Computation with Advanced CUDA Features, GTC Digital March, March 2020.
- <https://www.gpuhackathons.org/technical-resources>
- <https://rapids.ai/start.html>



# THANK YOU

