

NCHC Open Hackathons 2024

Day Final

Team5-parallel-minds, NTHU

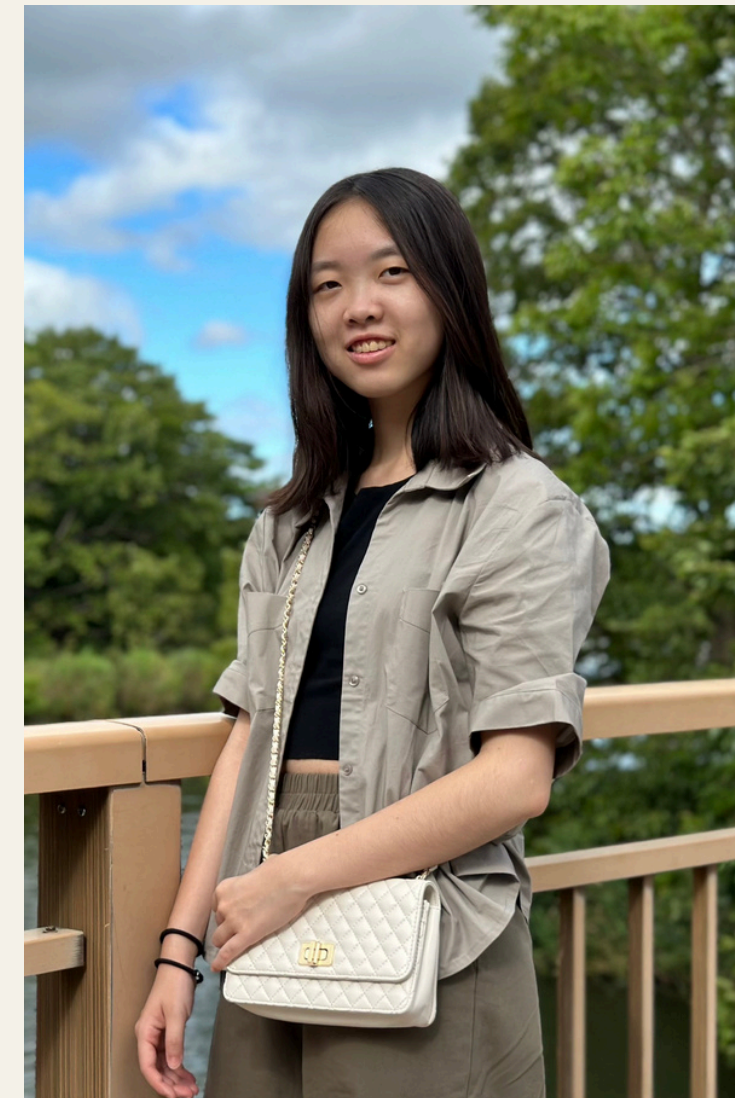
team-11-parallel-minds



程詩柔



謝之豫



熊恩伶

Firefly algorithm (visualization)

1. 設定螢火蟲總數(population),空間維度(dimen),訓練次數(max_iter)
 2. 將所有螢火蟲隨機放在空間中並隨機給定亮度(fitness)
 3. 每個iteration，每隻螢火蟲都要朝附近最亮的螢火蟲前進(位置更新)
 4. 更新所有螢火蟲的亮度(fitness更新)
- 迴圈時間複雜度 $\text{max_iter} * \text{population} * \text{dimen} * \text{population}$

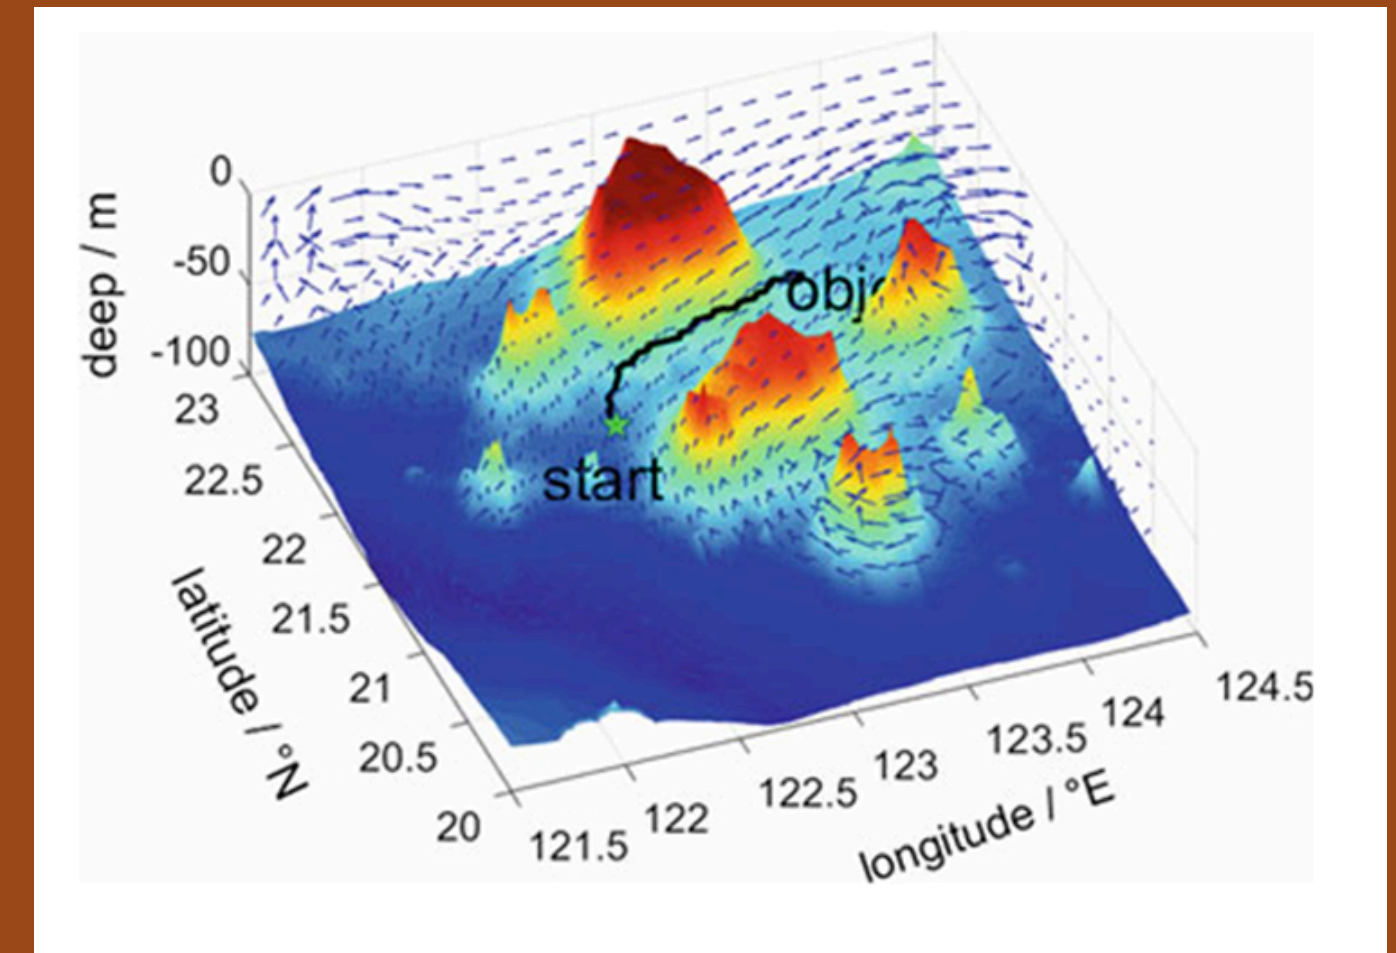
why is it important? Robot path planning

Firefly algorithm

- textbook: Nature-Inspired Computation in Navigation and Routing Problems Algorithms, Methods and Applications Editors: Xin-She Yang, Yu-Xin Zhao
- Vehicle path planning : Particle swarm optimization algorithm vs Firefly algorithm
- 準確率 : Firefly algorithm > Particle swarm optimization algorithm
- 執行時間: Firefly algorithm > Particle swarm optimization algorithm

Goals :

- accelerate firefly algorithm with cuda



Progress -Week1

Stats System View										
MPI Event Trace	Time	Total Time	Instances	Avg	Med	Min	Max	StdDev	Range	
NVTX GPU Projection Summary	98.9%	329.187 s	50	6.584 s	6.678 s	3.189 s	9.520 s	1.901 s	:fun() calculate fitness	
NVTX GPU Projection Trace										
NVTX Push/Pop Range Summary	1.1%	3.807 s	6125000	621 ns	320 ns	79 ns	28.061 ms	20.046 µs	:update firefly position	
NVTX Push/Pop Range Trace	0.0%	403.836 µs	1	403.836 µs	403.836 µs	403.836 µs	403.836 µs	0 ns	:write result file	
NVTX Range Kernel Summary	0.0%	231.226 µs	49	4.718 µs	4.020 µs	1.490 µs	21.599 µs	3.719 µs	:update best fitness	
NVTX Range Summary	0.0%	134.489 µs	1	134.489 µs	134.489 µs	134.489 µs	134.489 µs	0 ns	:pop initialize	
NVTX Start/End Range Summary	0.0%	4.390 µs	1	4.390 µs	4.390 µs	4.390 µs	4.390 µs	0 ns	:FA() initialize parameter	
Network Devices Congestion										
NvVideo API Summary										
OS Runtime Summary										
OpenACC Summary										

```
vector<double> fun(const vector<vector<double>>& pop) {  
    //nvtxRangePushA("fun() calculate fitness");  
    vector<double> result;  
    for (int i = 0; i < pop.size(); i++) {  
        double funsum = 0;  
        for (int j = 0; j < D; j++) {  
            double x = pop[i][j];  
            funsum += x * x - 10 * cos(2 * M_PI * x);  
        }  
        funsum += 10 * D;  
        result.push_back(funsum);  
    }  
    return result;  
    //nvtxRangePop();  
}
```

bottleneck:

fitness亮度更新計算更新耗時最久

strategy:

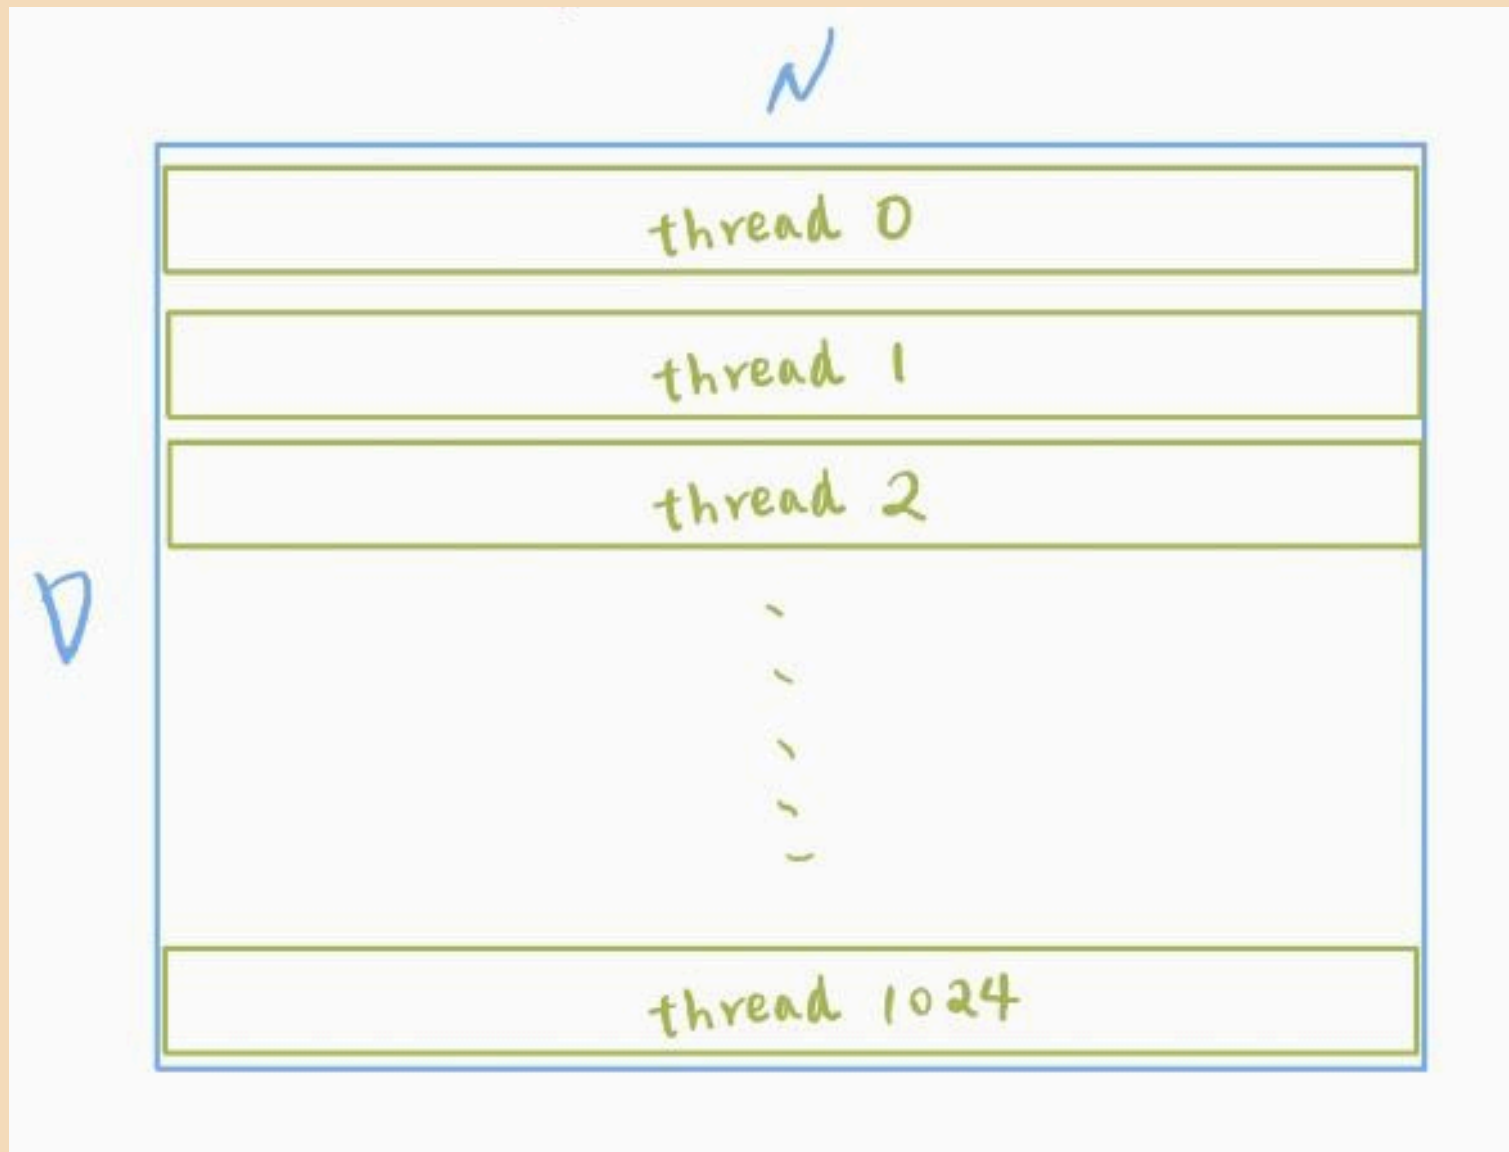
將fun()寫成kernel放上gpu做運算

Progress -Week1

將fun_kernel用1024 thread 平行運算

Problem : 反而比cpu code 還慢!!

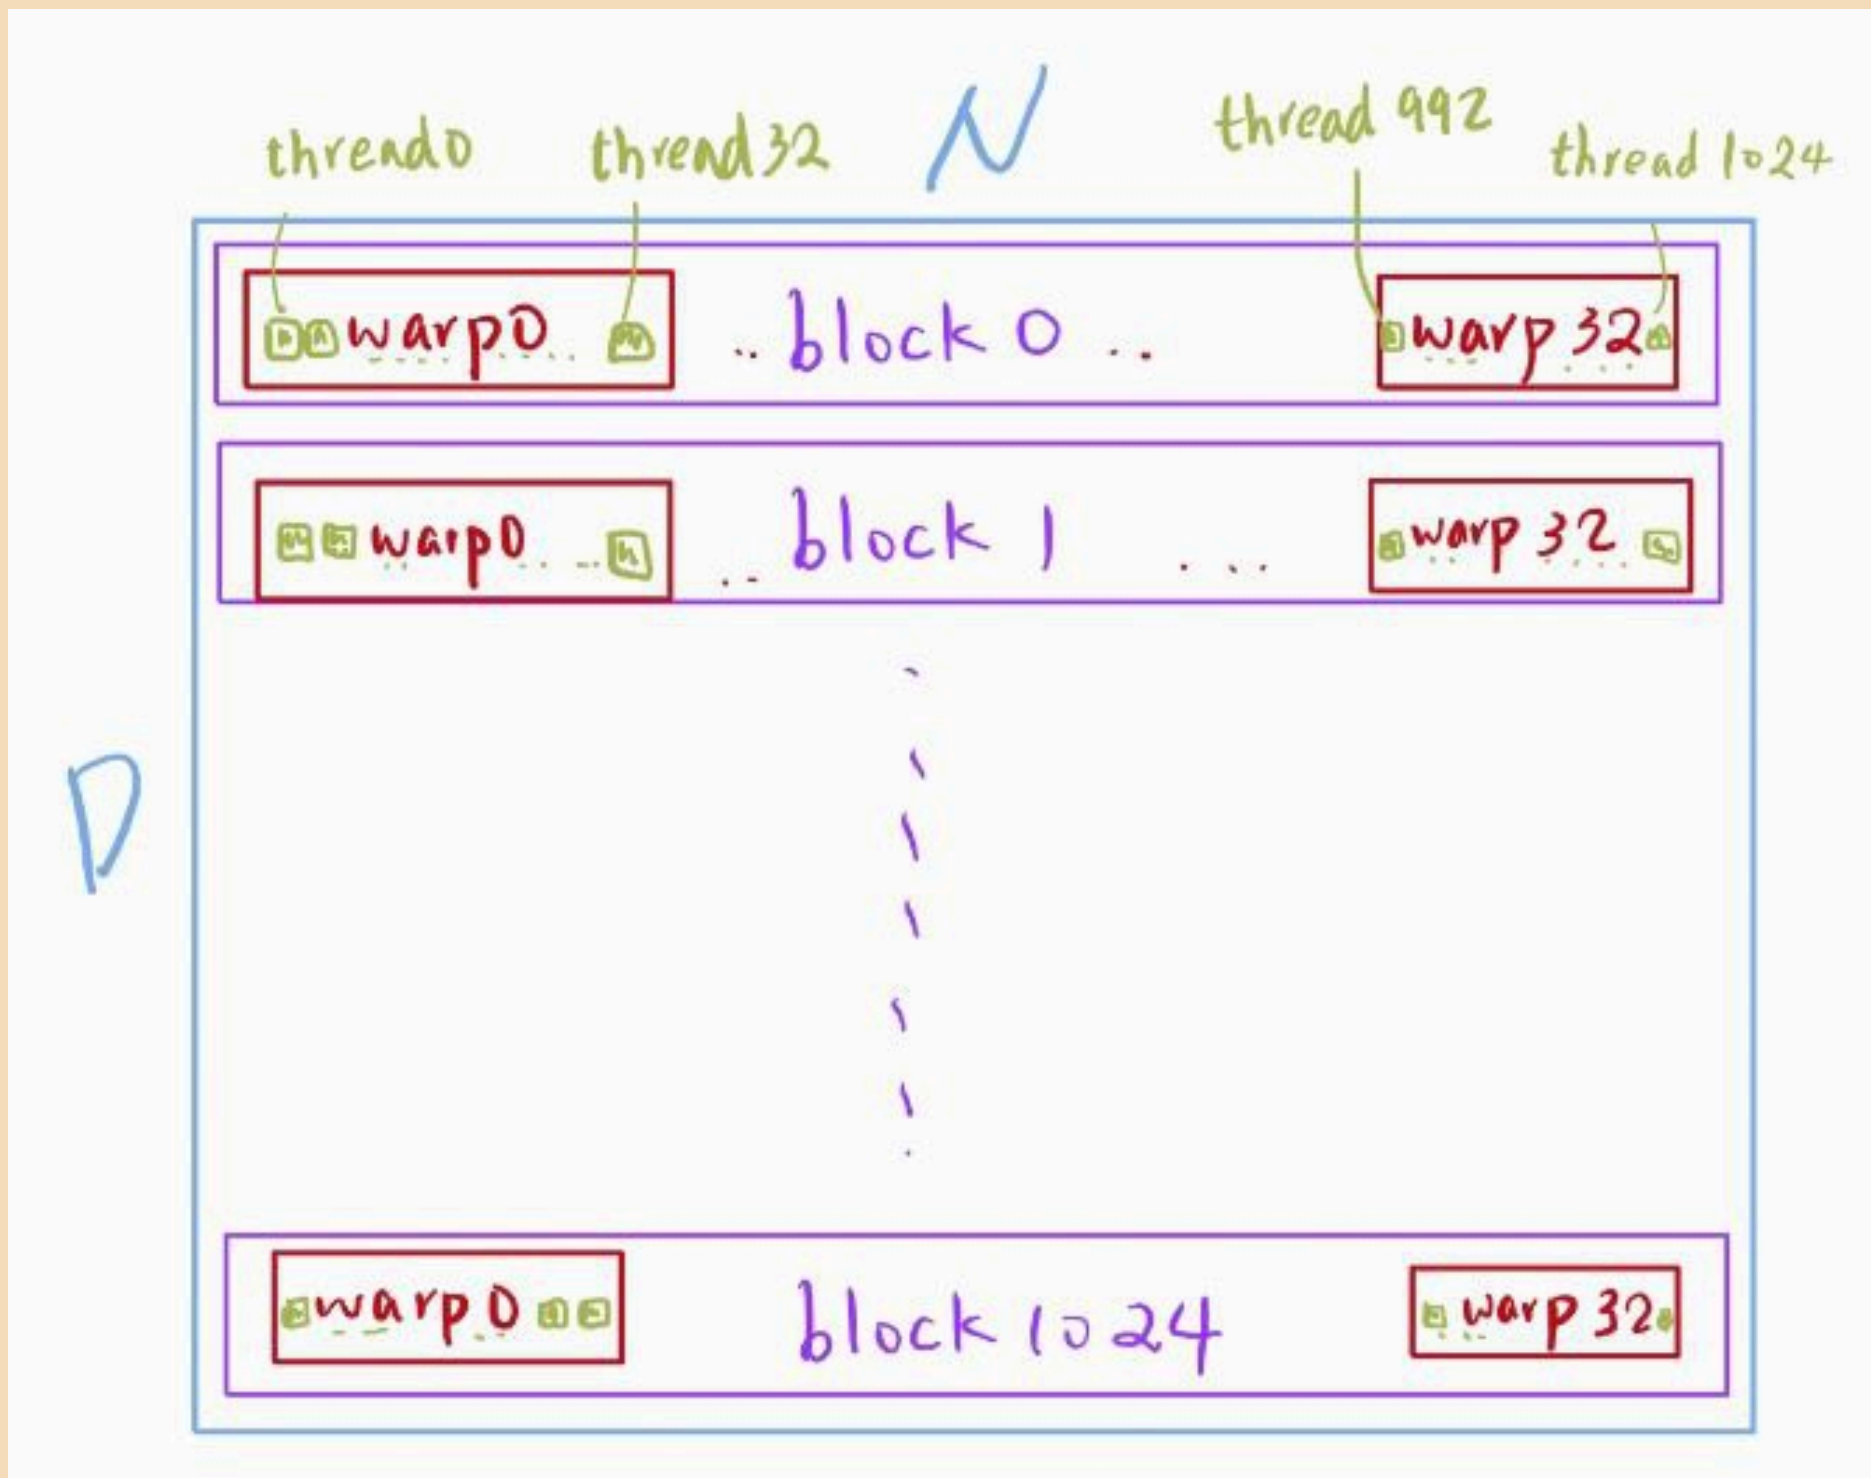
Reason: 平行度太低+頻繁的Memcpy很耗時



low parallelization

Progress -Week2

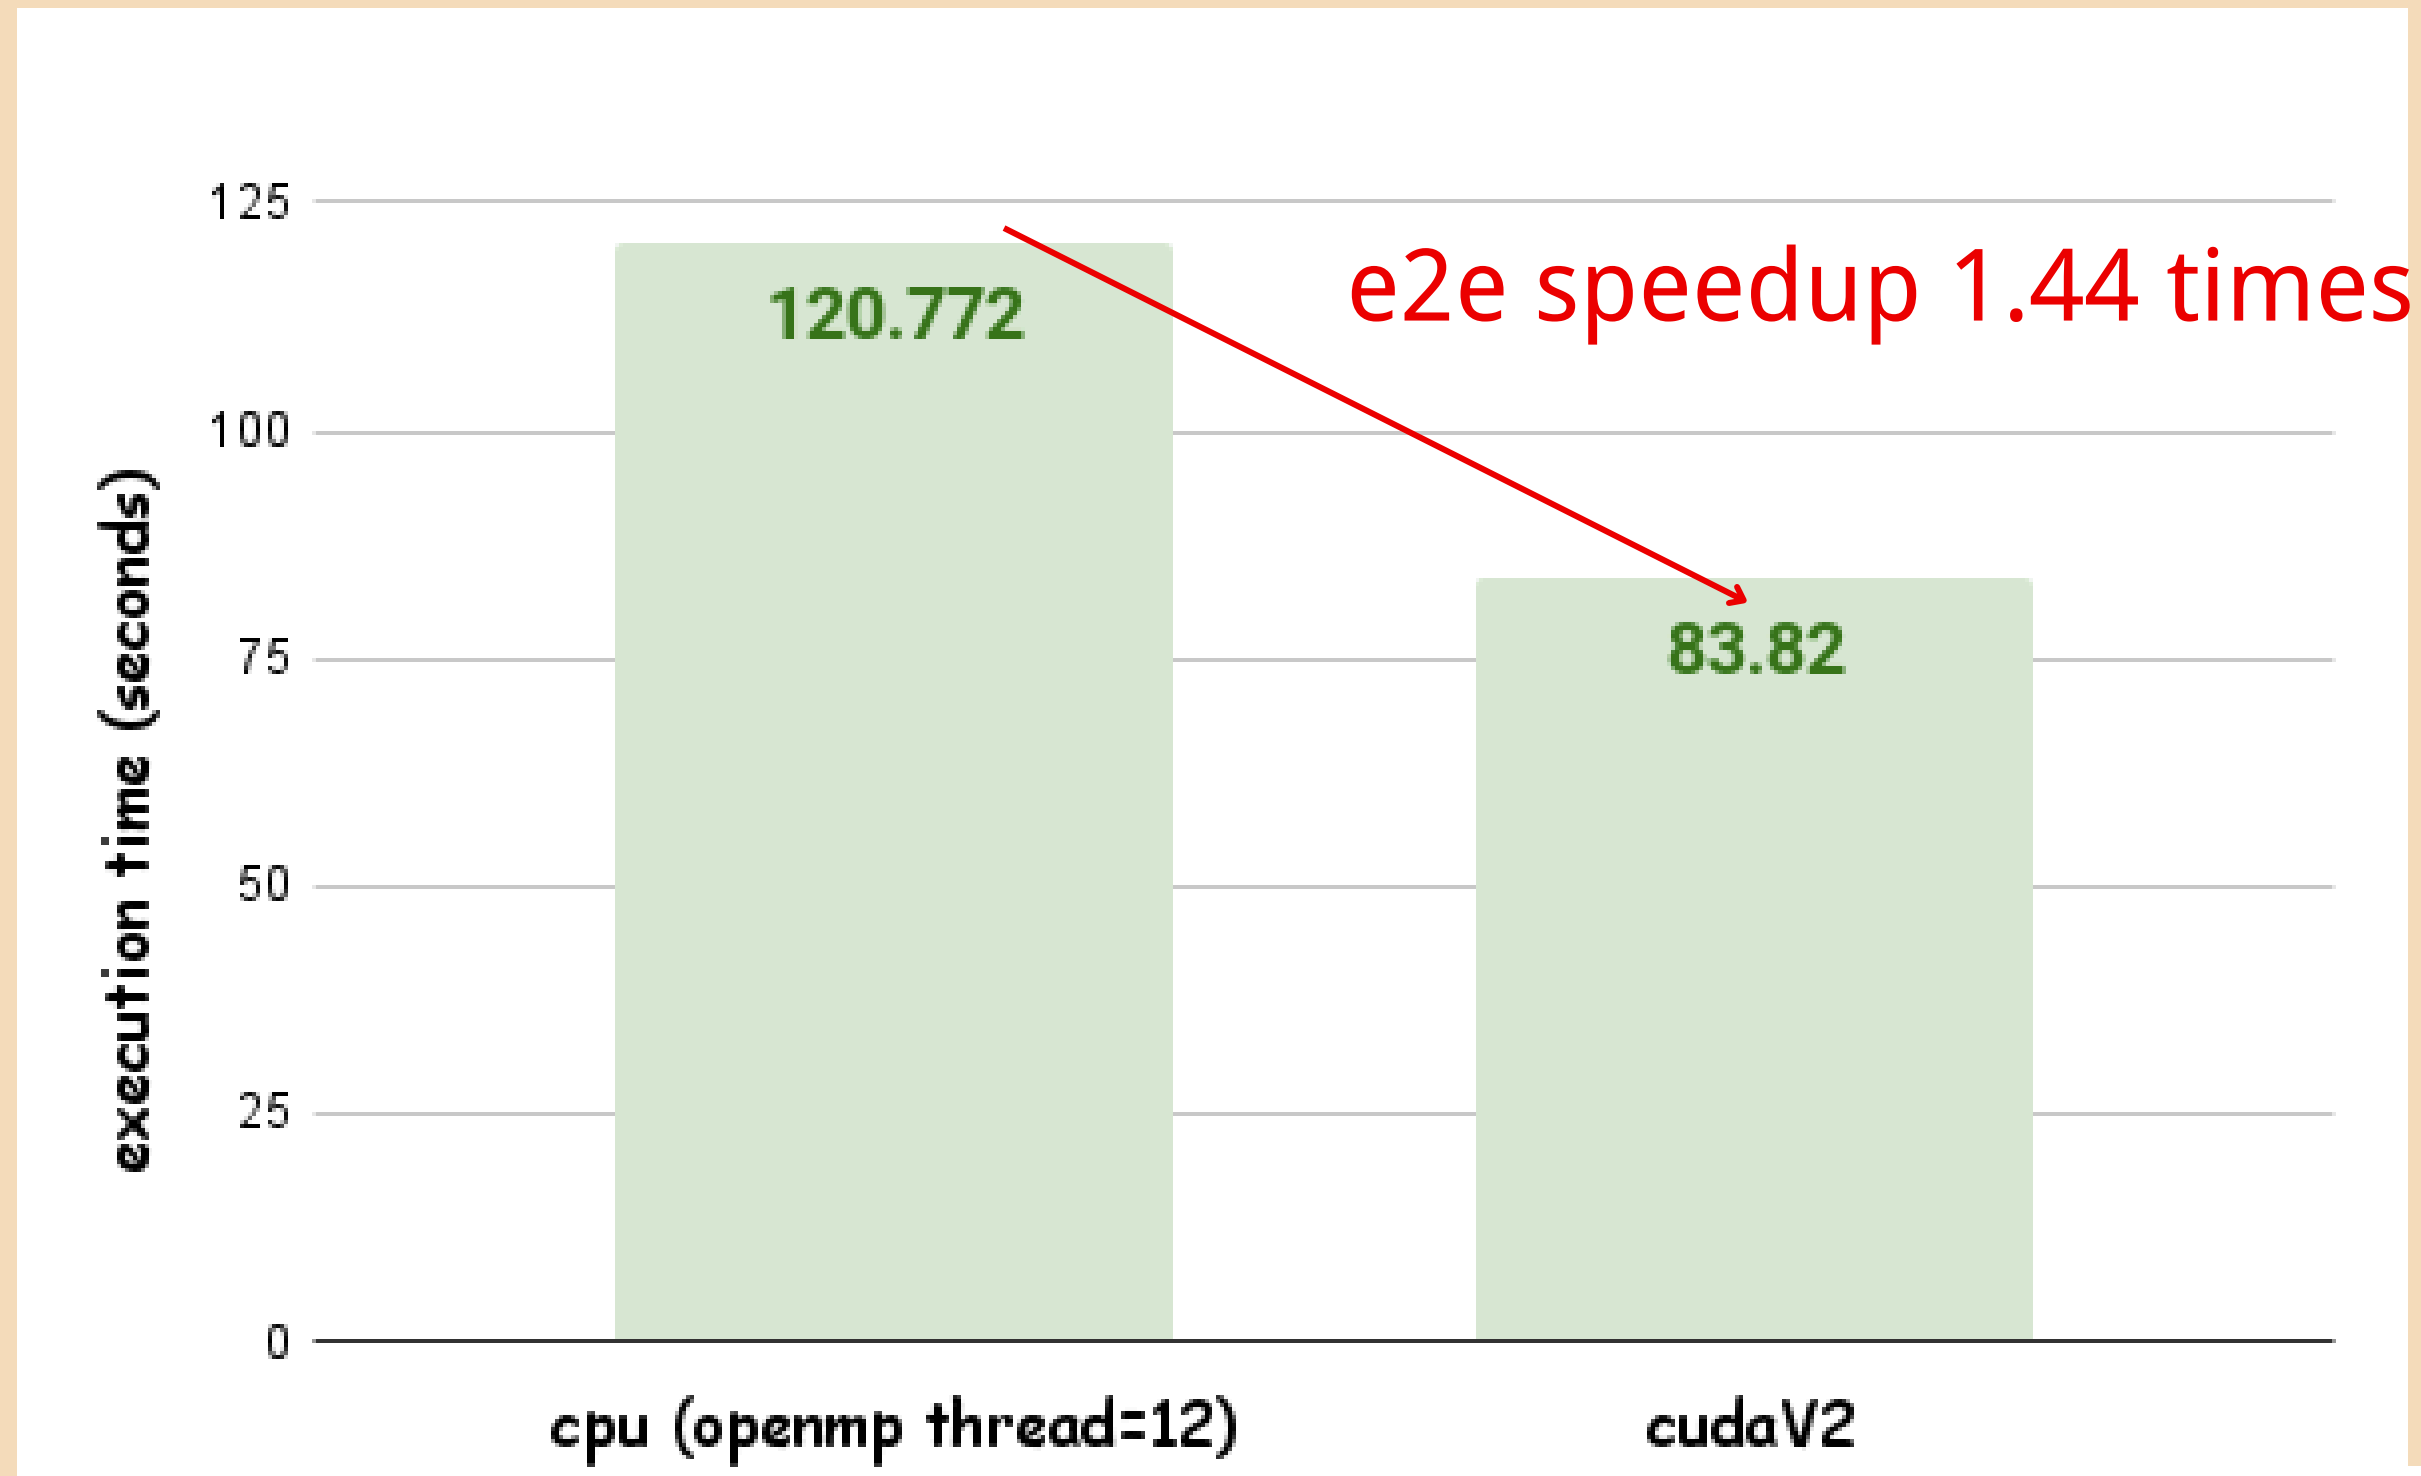
high parallelization



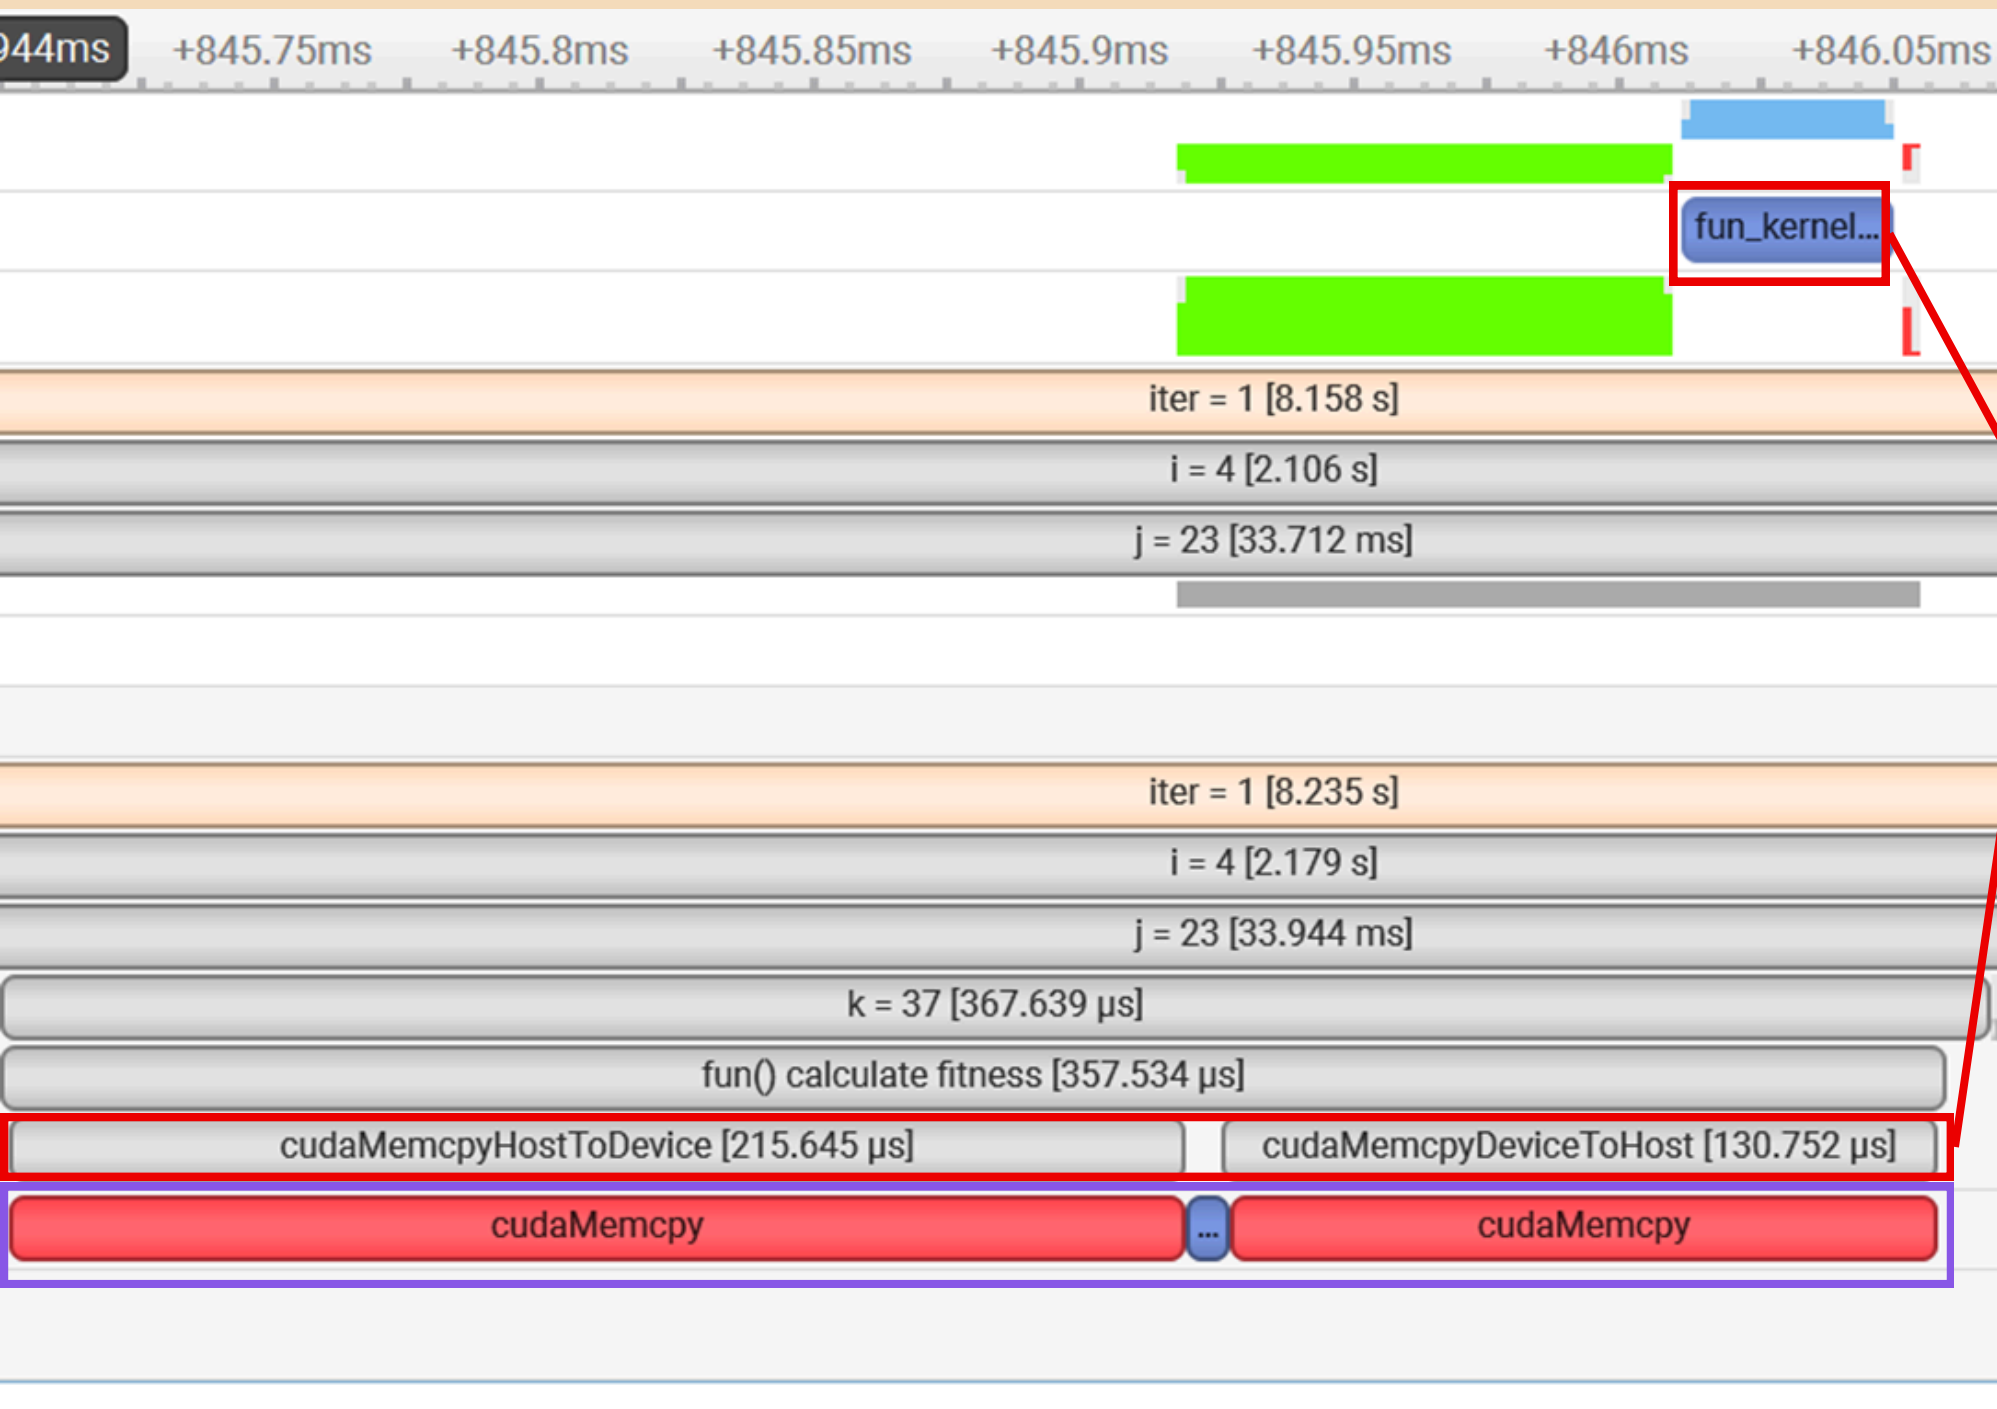
- 使用1024blocks中的32個wraps
- 每個wrap中的32thread
- WarpReduce

Progress -Week2

problem size **population:1024 dimen:512 max_iter 3**



Progress -Week2



bottleneck:
kernel計算時間佔10%
剩餘90%均為Memcpy的時間

strategy:
透過將fun_kernel計算前後也放上gpu做運算節省Memcpy時間

Progress -Week3

```
if (fitness[i] > fitness[k]) {  
    r_distance += pow(pop[i * fa.D + j] - pop[k * fa.D + j], 2);  
    double Beta = fa.B * exp(-fa.G * r_distance);  
    double xnew = pop[i * fa.D + j] + Beta * (pop[k * fa.D + j] - pop[i * fa.D + j]) + steps;  
  
    xnew = min(max(xnew, fa.Lb[0]), fa.Ub[0]);  
    pop[i * fa.D + j] = xnew;  
}
```

update_pop_kernel

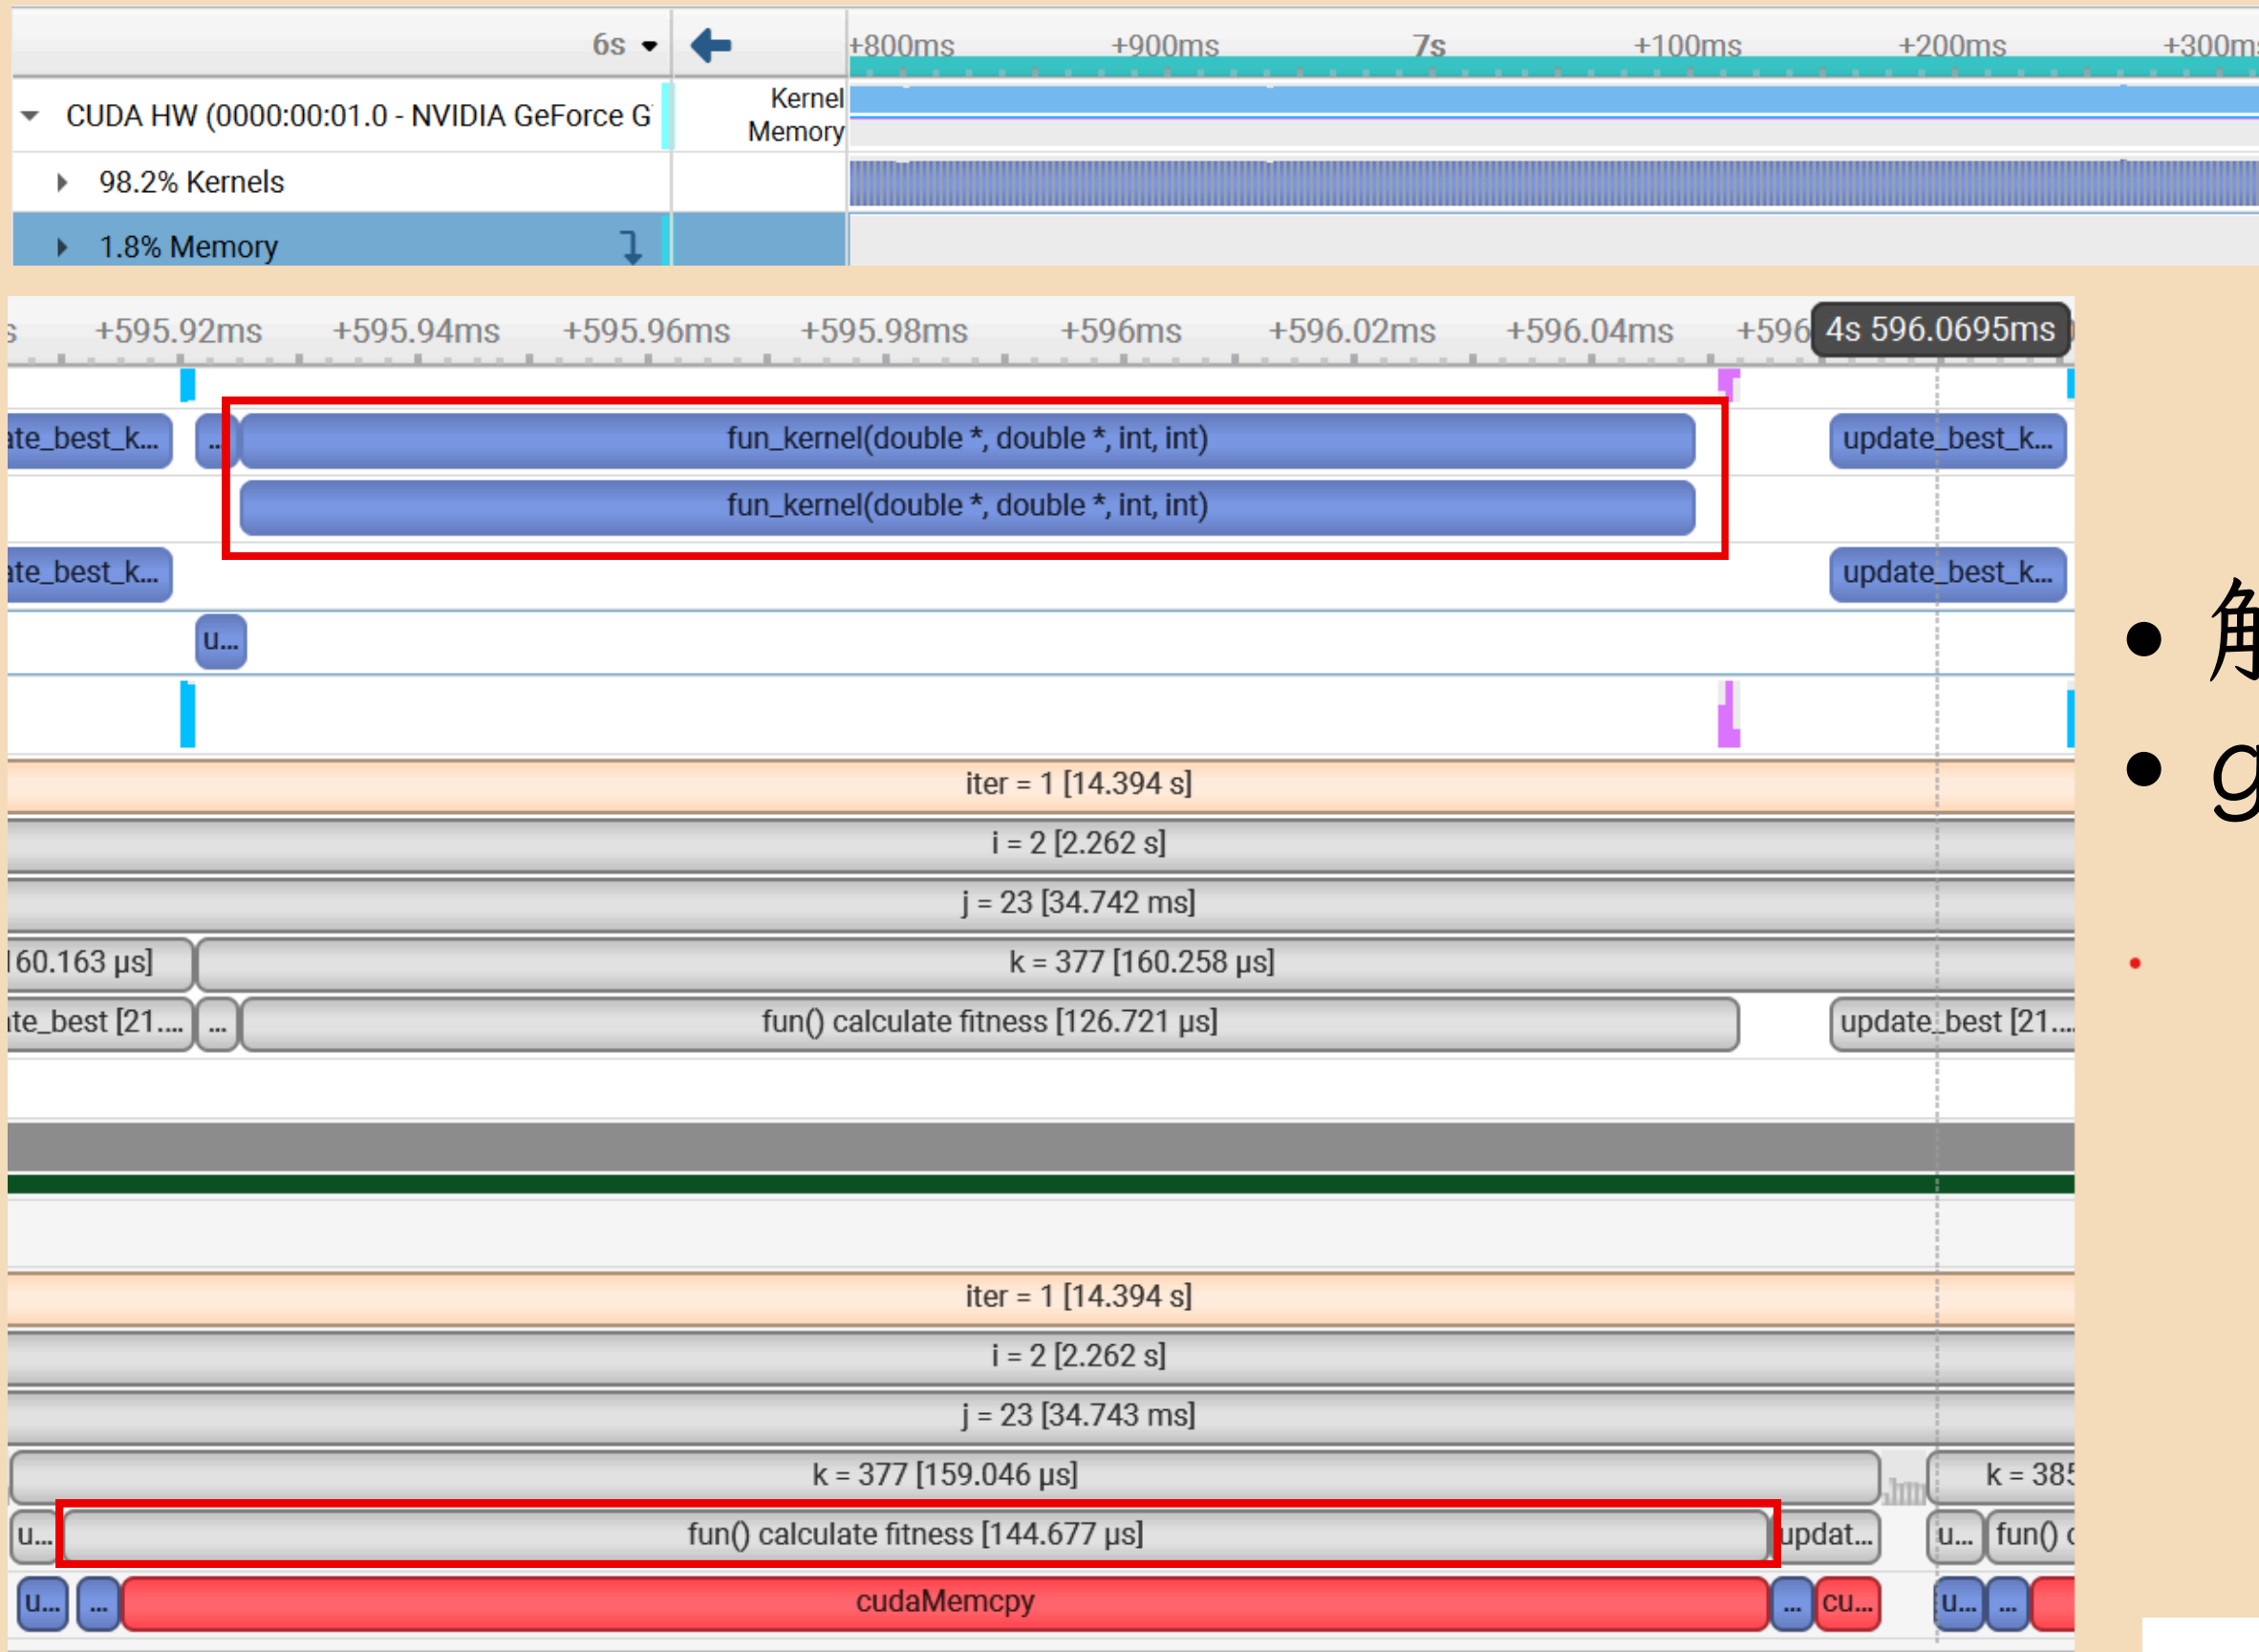
```
// Update fitness after position update  
fitness = fa.fun(pop);
```

fun_kernel

```
auto best_iter = min_element(fitness.begin(), fitness.end()); //取得min fitness  
best_ = *best_iter; //取得min fitness  
int arr_ = distance(fitness.begin(), best_iter); //取得min fitness index位置  
  
for (int j = 0; j < fa.D; j++) {  
    best_para_[j] = pop[arr_ * fa.D + j]; //將min_fitness整個dimention位置存入best_para_  
}
```

update_best_kernel

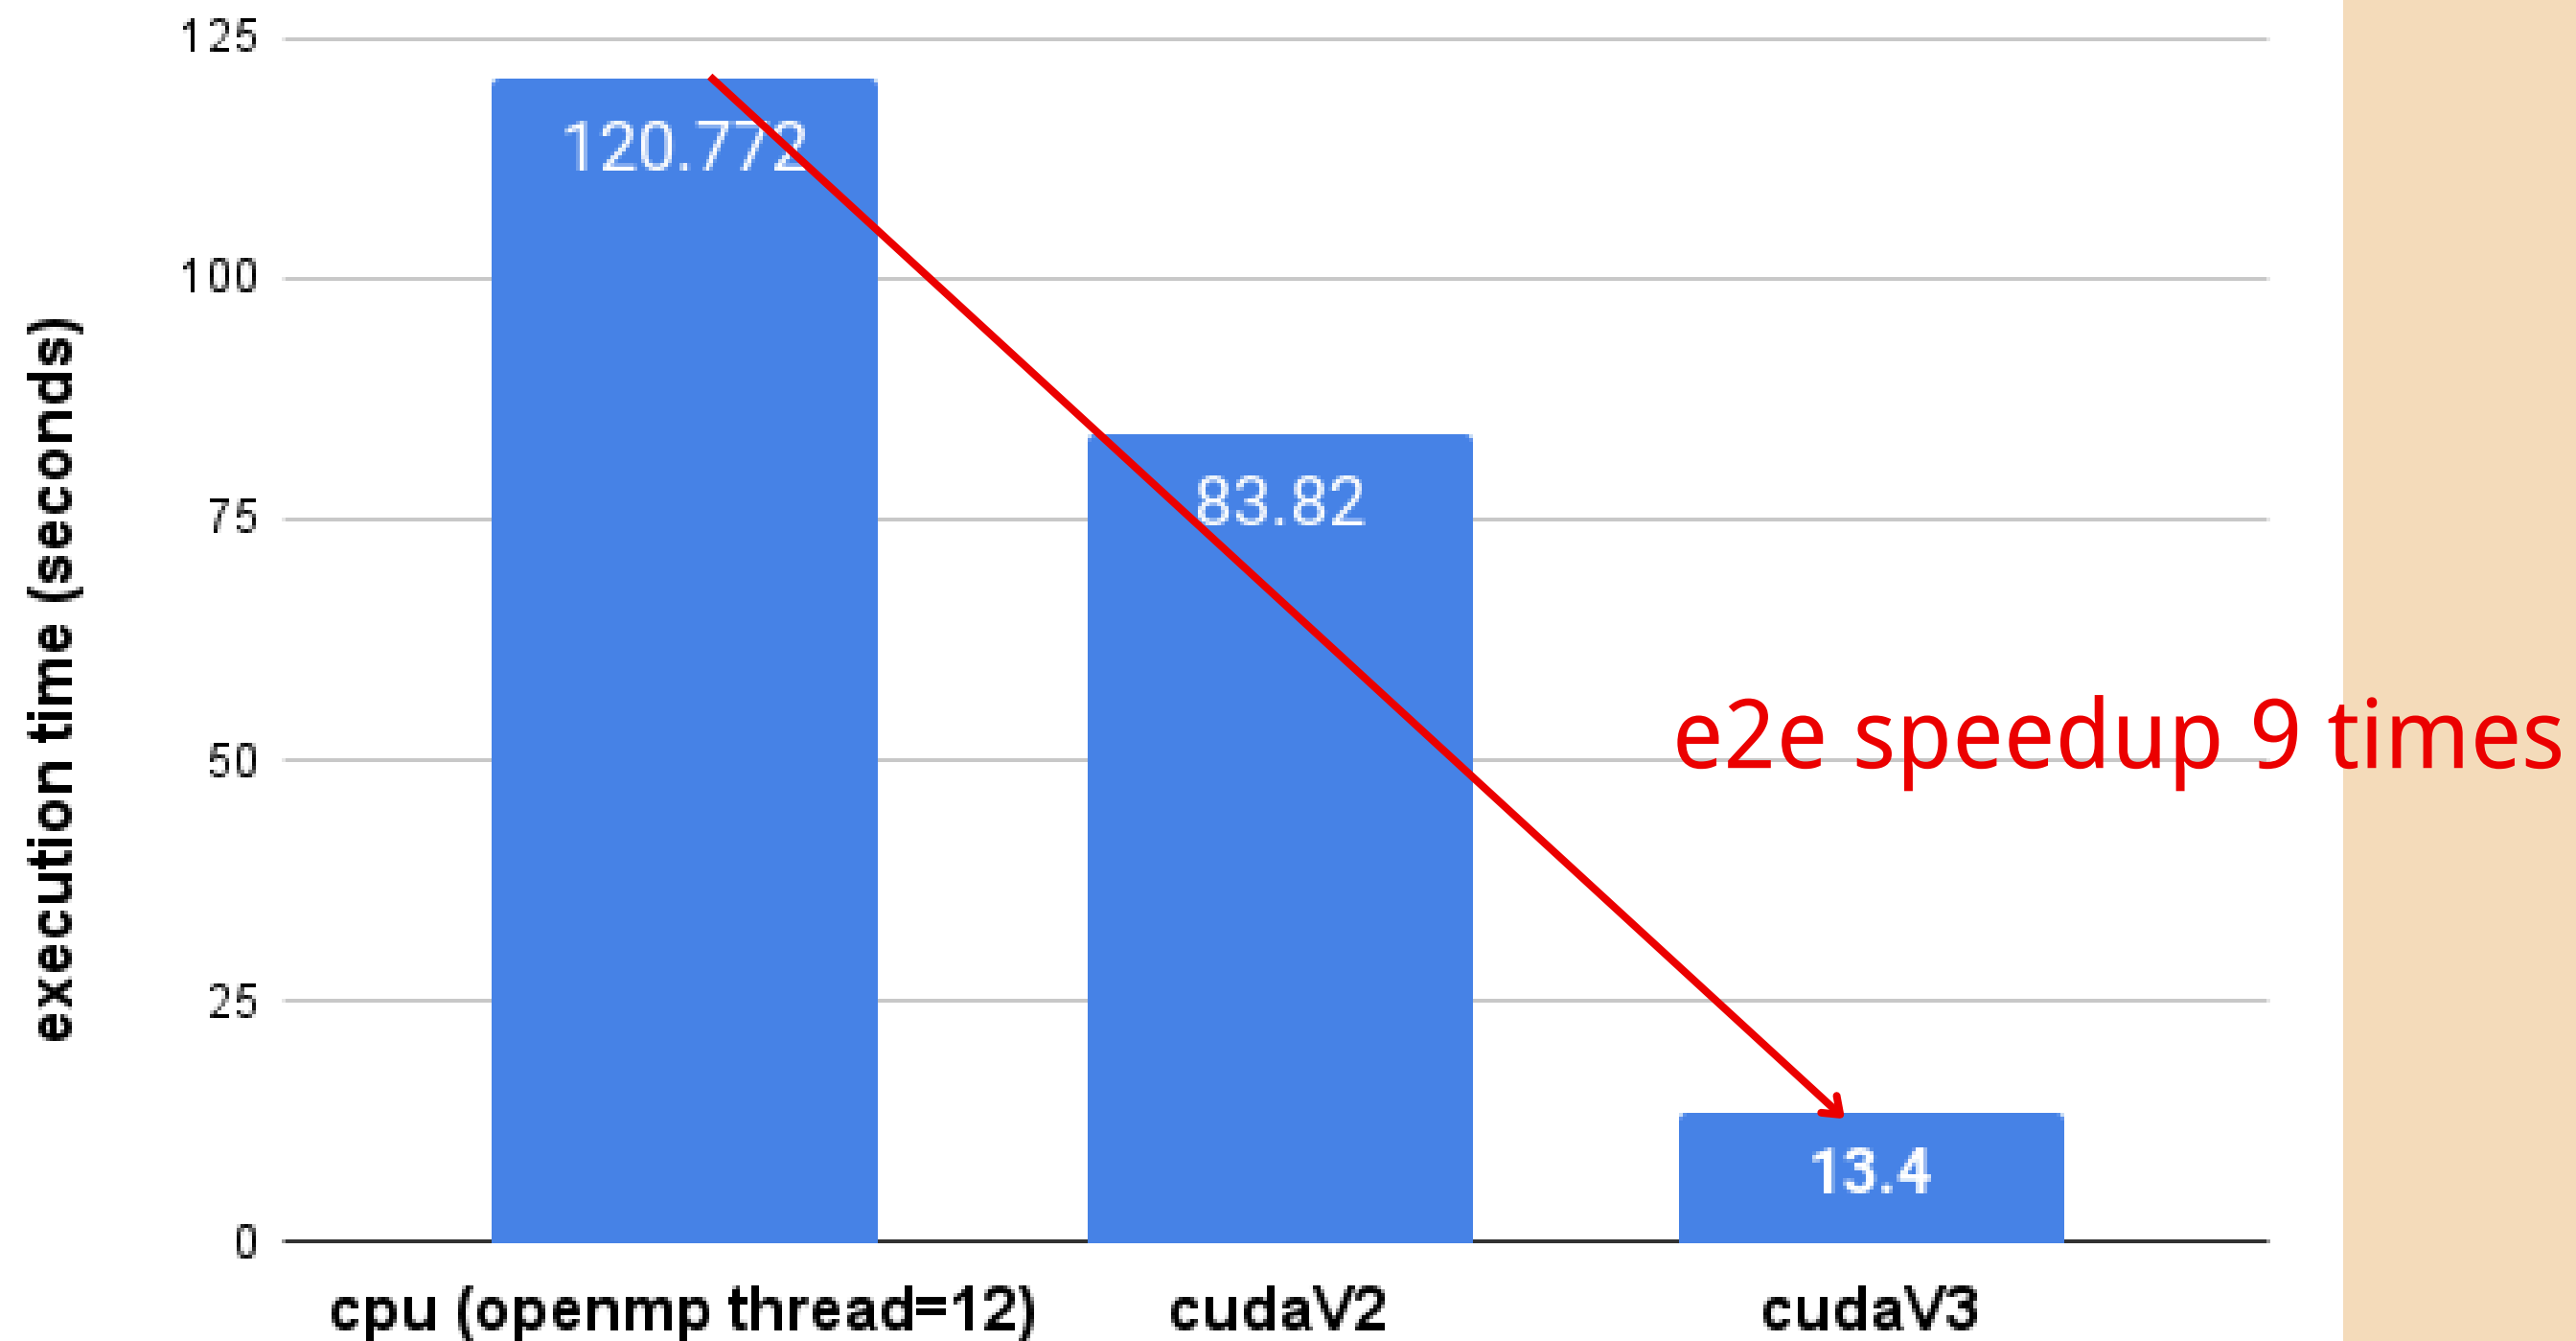
Progress -Week3



- 解決Week2遇到的memcpy
- gpu utilization > 90%


Progress -Week3

problem size : population:1024 dimen:512 max_iter 3



Energy Efficiency improvement

Energy Efficiency



Energy efficiency improve 10 times

What problems have you encountered?

- 平行運算時遇到race condition：謹慎處理data dependency避免同步讀寫資料
- 盡可能減少Memcpy時間
- 能夠運用多個blocks中多個wrap運算提高平行度



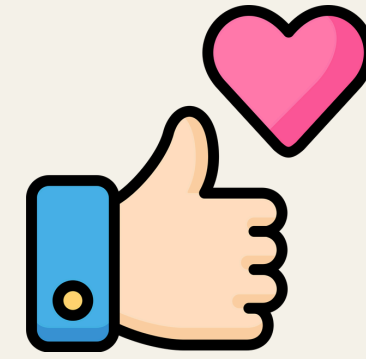
Wishlist

- 可以贊助H100給清大的超算團隊嗎？
- 免費cuda教學課程



Final Thoughts

超級值得!!!!



- 1.讓我更熟悉cuda語法，並成功用cuda加速firefly algorithm
- 2.知道如何解決Memcpy耗時情況
- 3.學會cuda加速小技巧：WarpReduce。
- 4.更熟悉如何使用nsys profile並解讀report

What's next?

- 1.優化fun_kernel : use coalesced memory access
- 2.優化update_best_kernel : 使用多個blocks平行計算



謝謝大家!!!

