


Mentors

NCHC Open Hackathon 2024

Final presentation

GPU-Based Acceleration VVM Team

Leo Chen
NVIDIA



Min Yu
NVIDIA



Miles Hsieh
NTU/LCDM



Jerry Chen
NTU/AEL



Stacy Guo
NTU/AEL



Shao Tseng
NTU/AEL




Aries Fan
NTU/LCDM

OpenACC
More Science, Less Programming



special thanks to 尚恩



Building Digital Twin of Taiwan



Building Digital Twin of Taiwan




Clouds, Rain, and Pollution




***14000 CPU hour for a 24-hr simulation**

Atmospheric Physical and Chemical Processes Simulation using Vector Vorticity equation cloud-resolving Model (VVM)

- Fortran code
- 4-core CPU with MPI
- Domain: 128x128x50
- integrates 720 steps
(2 hours)



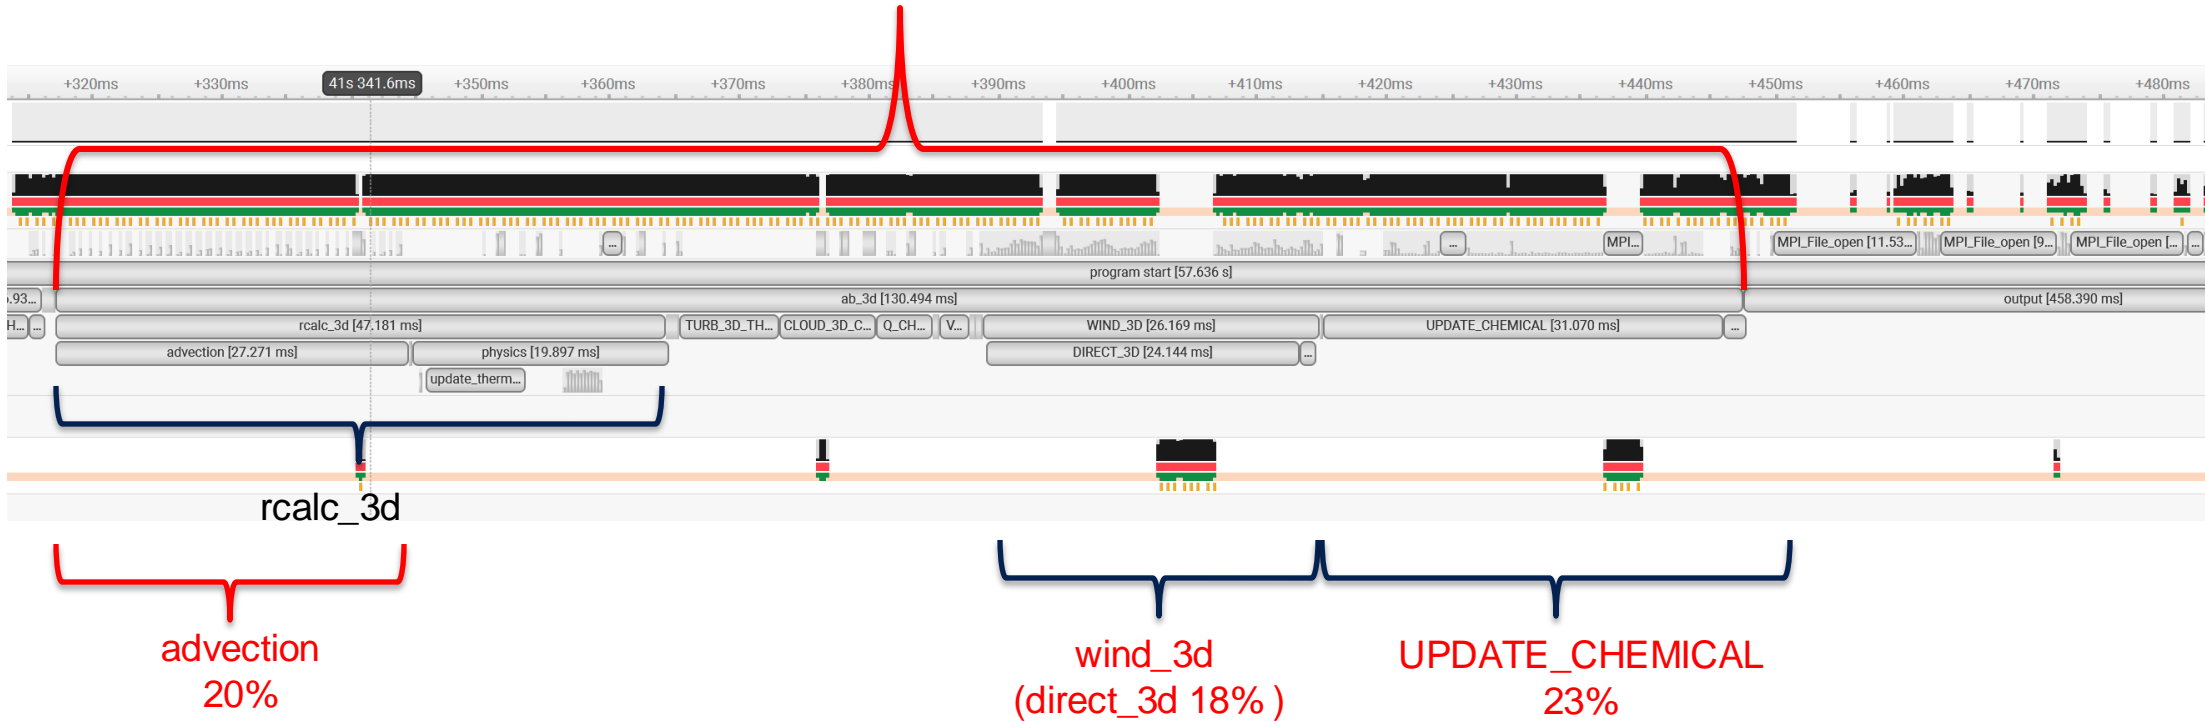
HACKATHON EVOLUTION



Chandrasekaran et al., (2018)
doi: 10.1109/MCSE.2018.042781332.

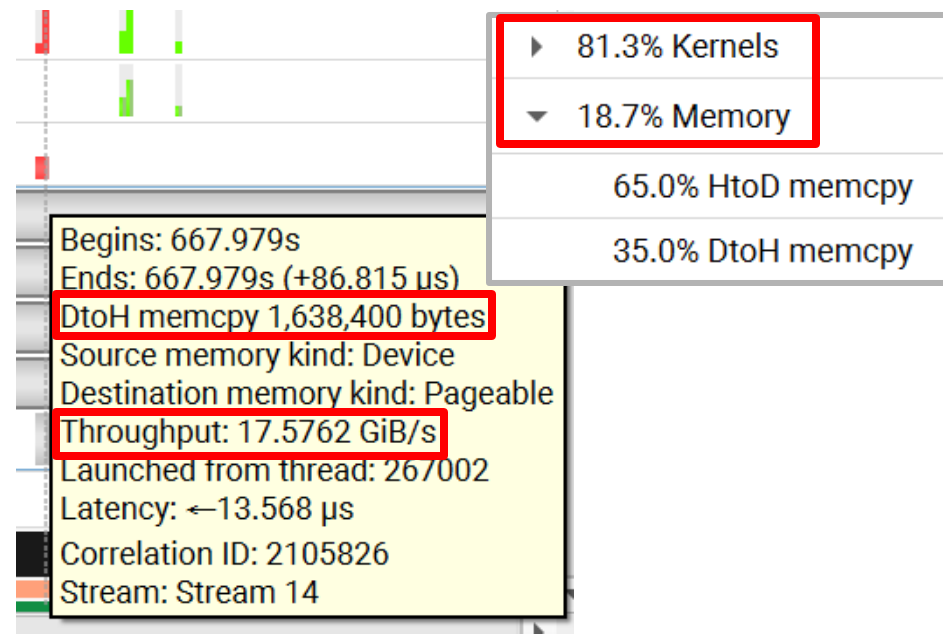
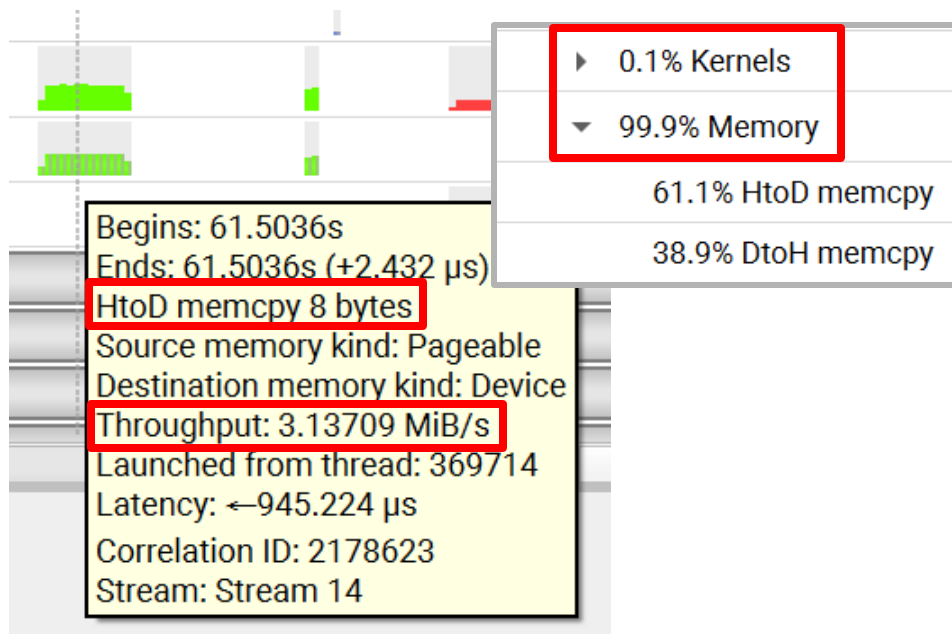
initial strategy

ab_3d (one step), 130.5ms



- porting path: openACC multi-core → openACC GPU

problems and solutions




- **Improve data locality:**



- Pre-load reusable data before calling the advection module and present them in the module
- Create local variables in the GPU memory
- pruning legacy codes: reduce loops from 19 to 8
- asynchronize multiple loops

Results and Final Profile

GPU



GPU Porting Comparison



4-core CPU

- 4 core CPU baseline
- OpenACC data region
- Collapse construct
- Data transfer order
- Data transfer by point
- Fuse loops
- Asynchronize
- Isolated compute construct


Energy Efficiency

INPUTS			
# CPU Cores	4		
# GPUs (A100)	1		
Application Speedup	18.7x		
Node Replacement	4.7x		
GPU NODE POWER SAVINGS			
	AMD Dual Rome 7742	8x A100 80GB SXM4	Power Savings
Compute Power (W)	5,143	6,500	-1,358
Networking Power (W)	217	93	124
Total Power (W)	5,360	6,593	-1,233
Node Power efficiency	0.8x		
ANNUAL ENERGY SAVINGS PER GPU NODE			
	AMD Dual Rome 7742	8x A100 80GB SXM4	Power Savings
Compute Power (kWh/year)	45,048	56,940	(11,892)
Networking Power (kWh/year)	1,902	814	1,088
Total Power (kWh/year)	46,950	57,754	(10,804)
\$/kWh	0.20		
Annual Cost Savings	(2,160.71)		
3-year Cost Savings	(6,482.13)		
Metric Tons of CO ₂	(8)		
Gasoline Cars Driven for 1 year	(2)		
Seedlings Trees grown for 10 years	(127)		

- Time consumption of advection in one time step
- GPU acceleration is more apparent in larger-domain-computation
- We're able to achieve over 1860% speedup in the 384*384 domain.
- **TaiwanVVM with 1024*1024 domain can achieve over 750% speedup in a sample test.**

unsolved problems: update_chemical module

- precision issues (not sure)



```
subroutine ITER(gdt,k)
implicit none

real gdt
integer k
integer j,n,r,s !loop count
real :: YP_local(1:i1,1:j1), YL_local(1:i1,1:j1)

YP_local = 0.
YL_local = 0.
!$acc data copyin(PL_scheme, KeffT, keff, RC, ysum, gdt, k, H2O) copy(ynew, YP, YL)
!$acc parallel loop private(j, kreact, YP_local, YL_local)
do n=1, nchsp !chemical species
if (PL_scheme(n)%active .EQV. .TRUE.) then !reactive or not (True or False)
if (PL_scheme(n)%name == H2O%name) cycle !don't do calculation of H2O
do j=1, PL_scheme(n)%nr_PL !chemical reactions (depend on chemical speices)
if (RC(PL_scheme(n)%PL(j)%r_nr)%raddep == 1) then !photolysis or not (1 or 0)
select case (PL_scheme(n)%PL(j)%formula) !select case of different formulas (8 types)
case (0)
if(PL_scheme(n)%PL(j)%PorL == 1) then !production or loss (1 or 0)
YP_local = YP_local + PL_scheme(n)%PL(j)%coef * keff(:, :, RC(PL_scheme(n)%PL(j)%r_nr)%Kindex, k) !Production
else
YL_local = YL_local + PL_scheme(n)%PL(j)%coef * keff(:, :, RC(PL_scheme(n)%PL(j)%r_nr)%Kindex, k) !Loss
endif
case (1~7) !other cases
end select
endif
enddo !end of chemical reactions
do r=1, i1
do s=1, j1
!$acc atomic
YP(r,s,n) = YP(r,s,n) + YP_local(r,s) !data output
!$acc atomic
YL(r,s,n) = YL(r,s,n) + YL_local(r,s) !data output
enddo
enddo
ynew(:, :, n) = max(0.0, (ysum(:, :, n)+gdt*YP(:, :, n))/(1.0+gdt*YL(:, :, n))) !data output
else
ynew(:, :, n) = ysum(:, :, n)
endif
enddo !end of chemical species
!$acc end data
end subroutine ITER
```

ITT=12: Contouring: -0 to 0.00012 interval 1e-05

ITT=360: Contouring: -0.016 to 0.001 interval 0.001

ITT=720: Contouring: -0.03 to 0.004 interval 0.001

Wishlist and future plans

- keep working on update_chemical module
- focus on solving elliptical equations using CUDA
- cuSPARSE or cuFFT

$$\mu \frac{\partial w}{\partial t} + \left(\frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2} \right) w + \frac{\partial}{\partial z} \left[\frac{1}{\rho_0} \frac{\partial}{\partial z} (\rho_0 w) \right] = - \frac{\partial \eta}{\partial x} + \frac{\partial \xi}{\partial y} \quad (\text{A.6})$$

(Wu and Arakawa, 2011)

以GPU加速計算建構臺灣天氣的數位孿生

GBA-VVM團隊來自臺灣大學雲動力模擬暨大氣環境實驗室，將 VVM加速了18倍！！



天氣與生活息息相關，從日常是否帶傘到空汙政策規劃都需仰賴氣象資訊。臺灣地狹人稠，對劇烈天氣的時機、強度與影響範圍的掌握尤為重要。VVM是一個優異的高解析度大氣模式，能有效模擬臺灣常見的午後雷陣雨及空汙天氣，但其預報模擬需大量CPU運算，受限於硬體資源有限常導致模擬時間過長，難以提供即時預報參考，且高耗能的計算亦不符合減碳趨勢。為此，將VVM運算改由GPU加速，不僅縮短模擬時間，也符合減能需求。透過程式重整與清理降低計算複雜度，並以GPU進行VVM中氣象參數及污染物的傳送計算，實現了18倍的加速效果。此成果與開發經驗有助於提升VVM其他模組效能，並將VVM轉型為以GPU為核心的高速運算氣象模式，為建構臺灣天氣數位孿生系統奠定關鍵基礎。

Optimize data management

Optimize program construct

	4 core CPU baseline	Data transfer by point	OpenACC data region	Fuse loops	Collapse construct	Asynchronize directional advection	Change data transfer order	Isolated compute construct
128*128	130	7300(--)	790(--)	560(--)	440(--)	190(--)	120(1.08x)	110(1.18x)
384*384	3370				985(3.42x)	600(5.62x)	440(7.66x)	180(18.72x)



Thank You

OpenACC
More Science, Less Programming